

Architecture of Enterprise Applications 2

Messaging - kafka

Haopeng Chen

REliable, ***IN***telligent and ***SC***alable Systems Group (***REINS***)

Shanghai Jiao Tong University



Shanghai, China

<http://reins.se.sjtu.edu.cn/~chenhp>

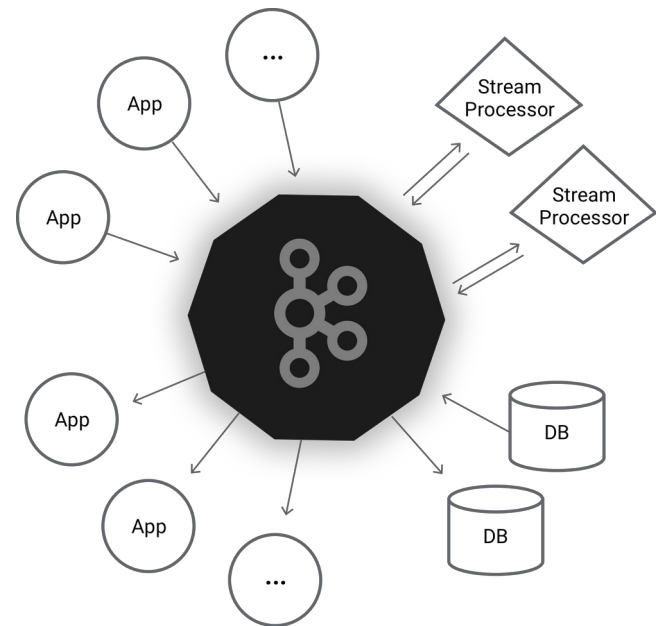
e-mail: chen-hp@sjtu.edu.cn

- Contents
 - Kafka
 - Introduction
 - Quick start
 - Communication with Kafka
- Objective
 - 能够根据系统需求，使用 Kafka 消息中间件来实现异步通信，包括跨编程语言的通信

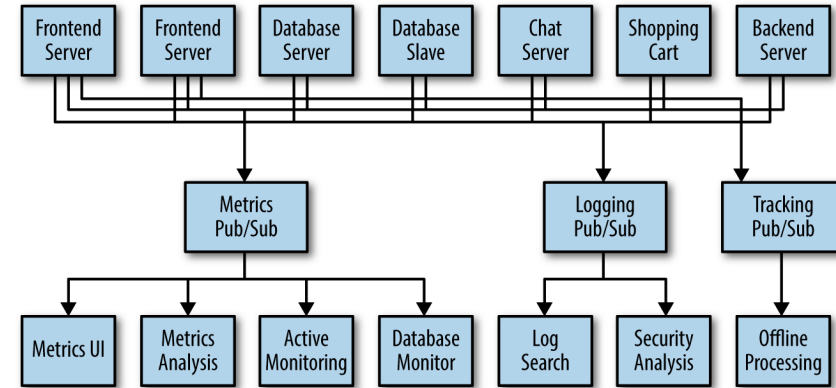
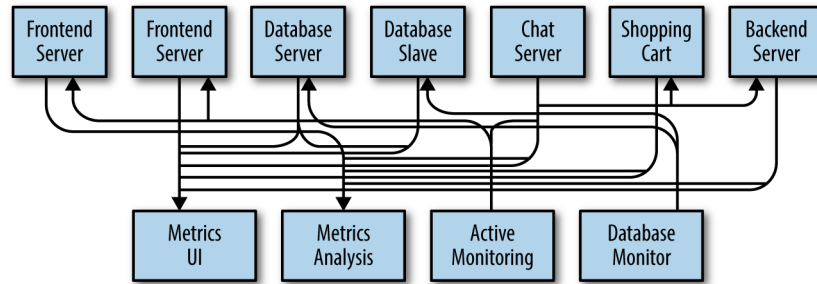
- from : <https://kafka.apache.org/books-and-papers>
- Apache Kafka Quick Start
 - <https://kafka.apache.org/quickstart>



- <http://kafka.apache.org>
- Apache Kafka
 - is an open-source **distributed event streaming**.
 - To **publish** (write) and **subscribe to** (read) streams of events, including continuous import/export of your data from other systems.
 - To **store** streams of events durably and reliably for as long as you want.
 - To **process** streams of events as they occur or retrospectively.



Publish/Subscribe Messaging



- Logs are *append-only* data structures that capture an *ordered sequence* of events.

```
# create the logfile
```

```
touch users.log
```

```
# generate four dummy records in our log
```

```
echo "timestamp=1597373669,user_id=1,purchases=1" >> users.log
```

```
echo "timestamp=1597373669,user_id=2,purchases=1" >> users.log
```

```
echo "timestamp=1597373669,user_id=3,purchases=1" >> users.log
```

```
echo "timestamp=1597373669,user_id=4,purchases=1" >> users.log
```

```
# append a new record to the log
```

```
echo "timestamp=1597374265,user_id=1,purchases=2" >> users.log
```

```
# print the contents of the log
```

```
cat users.log
```

```
# output
```

```
timestamp=1597373669,user_id=1,purchases=1 ❶
```

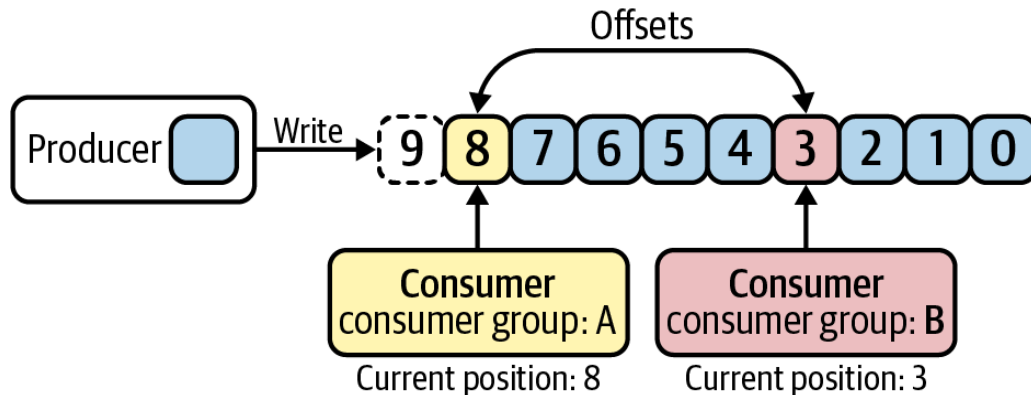
```
timestamp=1597373669,user_id=2,purchases=1
```

```
timestamp=1597373669,user_id=3,purchases=1
```


```
timestamp=1597373669,user_id=4,purchases=1
```

```
timestamp=1597374265,user_id=1,purchases=2 ❷
```


- Kafka refers to the position of each entry in its distributed log as an *offset*.
 - multiple consumer groups can each read from the same log, and maintain their own positions in the log/stream they are reading from.




- Topics
 - *homogeneous topics* that contain only one type of data, or *heterogeneous topics* that contain multiple types of data



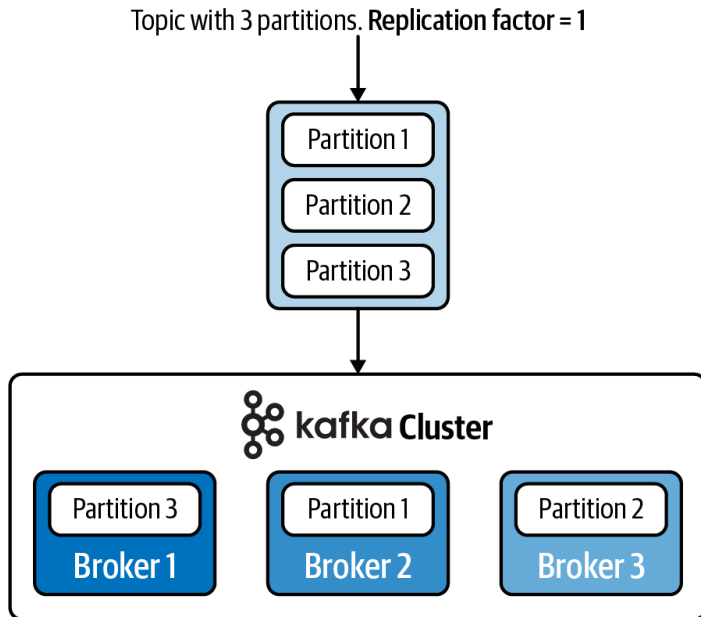
- Kafka topics are broken into smaller units called *partitions*.
 - Partitions are individual logs where data is produced and consumed from
 - The number of partitions for a given topic is configurable, and having more partitions in a topic generally translates to more parallelism and throughput



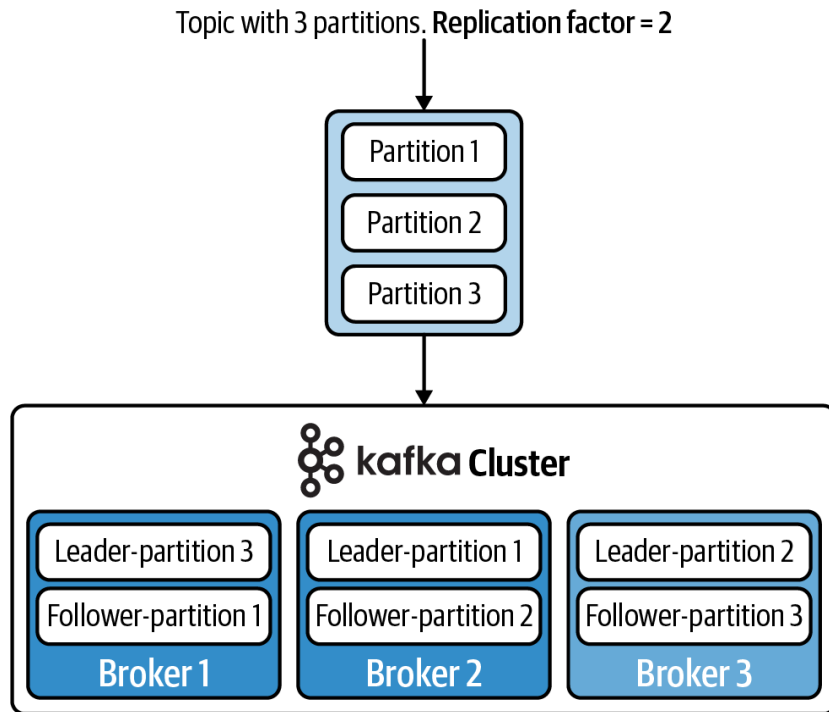
- An event is a timestamped key-value pair that records *something that happened*.
 - Application-level headers contain optional metadata about an event.
 - Keys are also optional, but play an important role in how data is distributed across partitions.
 - Each event is associated with a timestamp.
 - The value contains the actual message contents, encoded as a byte array.



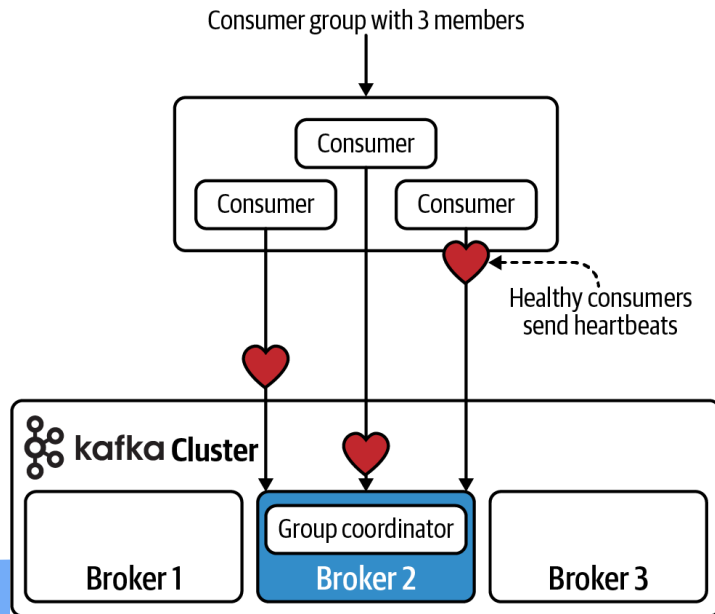
- Kafka operates as a cluster, and multiple machines, called *brokers*, are involved in the storage and retrieval of data.
 - Kafka clusters can be quite large, and can even span multiple data centers and geo- graphic regions.





- To achieve fault tolerance and high availability, you can set a replication factor when configuring the topic.



- Consumer groups are made up of multiple cooperating consumers, and the membership of these groups can change over time.
 - every consumer group is assigned to a special broker called the *group coordinator*, which is responsible for receiving heartbeats from the consumers, and triggering a *rebalance* of work whenever a consumer is marked as dead.



- Every active member of the consumer group is eligible to receive a partition assignment.



- Kafka has four core APIs:
 - The [Producer API](#) allows an application to **publish** a stream of records to one or more Kafka topics.
 - The [Consumer API](#) allows an application to **subscribe** to one or more topics and **process** the stream of records produced to them.
 - The [Streams API](#) allows an application to act as a **stream processor**, **consuming** an input stream from one or more **topics** and **producing** an output stream to one or more output **topics**, effectively transforming the input streams to output streams.
 - The [Connector API](#) allows building and running reusable producers or consumers that **connect Kafka topics to existing applications or data systems**. For example, a connector to a relational database might capture every change to a table.

- Quickstart
 - <https://kafka.apache.org/quickstart>

STEP 1: GET KAFKA

[Download](#) the latest Kafka release and extract it:

```
1 | $ tar -xzf kafka_2.13-3.2.1.tgz
2 | $ cd kafka_2.13-3.2.1
```


STEP 2: START THE KAFKA ENVIRONMENT

NOTE: Your local environment must have Java 8+ installed.

Run the following commands in order to start all services in the correct order:

```
1 | # Start the ZooKeeper service
2 | # Note: Soon, ZooKeeper will no longer be required by Apache Kafka.
3 | $ bin/zookeeper-server-start.sh config/zookeeper.properties
```

Open another terminal session and run:

```
1 | # Start the Kafka broker service
2 | $ bin/kafka-server-start.sh config/server.properties
```

Once all services have successfully launched, you will have a basic Kafka environment running and ready to use.

```
1. bin/zkServer.sh start
2. conf/zoo.cfg
   admin.serverPort = 8888
```

STEP 3: CREATE A TOPIC TO STORE YOUR EVENTS

Kafka is a distributed *event streaming platform* that lets you read, write, store, and process [events](#) (also called *records* or *messages* in the documentation) across many machines.

Example events are payment transactions, geolocation updates from mobile phones, shipping orders, sensor measurements from IoT devices or medical equipment, and much more. These events are organized and stored in [topics](#). Very simplified, a topic is similar to a folder in a filesystem, and the events are the files in that folder.

So before you can write your first events, you must create a topic. Open another terminal session and run:

```
$ bin/kafka-topics.sh --create --topic quickstart-events --bootstrap-server localhost:9092
```

All of Kafka's command line tools have additional options: run the `kafka-topics.sh` command without any arguments to display usage information. For example, it can also show you [details such as the partition count](#) of the new topic:

```
1 $ bin/kafka-topics.sh --describe --topic quickstart-events --bootstrap-server localhost:9092
2 Topic:quickstart-events PartitionCount:1 ReplicationFactor:1 Configs:
3 Topic: quickstart-events Partition: 0 Leader: 0 Replicas: 0 Isr: 0
```

STEP 4: WRITE SOME EVENTS INTO THE TOPIC

A Kafka client communicates with the Kafka brokers via the network for writing (or reading) events. Once received, the brokers will store the events in a durable and fault-tolerant manner for as long as you need—even forever.

Run the console producer client to write a few events into your topic. By default, each line you enter will result in a separate event being written to the topic.

```
1 | $ bin/kafka-console-producer.sh --topic quickstart-events --bootstrap-server localhost:9092
2 | This is my first event
3 | This is my second event
```

You can stop the producer client with `Ctrl-C` at any time.

STEP 5: READ THE EVENTS

Open another terminal session and run the console consumer client to read the events you just created:

```
1 | $ bin/kafka-console-consumer.sh --topic quickstart-events --from-beginning
2 | This is my first event                                     --bootstrap-server localhost:9092
3 | This is my second event
```

You can stop the consumer client with `Ctrl-C` at any time.

Feel free to experiment: for example, switch back to your producer terminal (previous step) to write additional events, and see how the events immediately show up in your consumer terminal.

Because events are durably stored in Kafka, they can be read as many times and by as many consumers as you want. You can easily verify this by opening yet another terminal session and re-running the previous command again.

- Access Kafka with Java
 - To use Producer API and Consumer API, you can use the following maven dependency:

```
<dependency>  
  <groupId>org.apache.kafka</groupId>  
  <artifactId>kafka-clients</artifactId>  
  <version>3.5.1</version>  
</dependency>
```

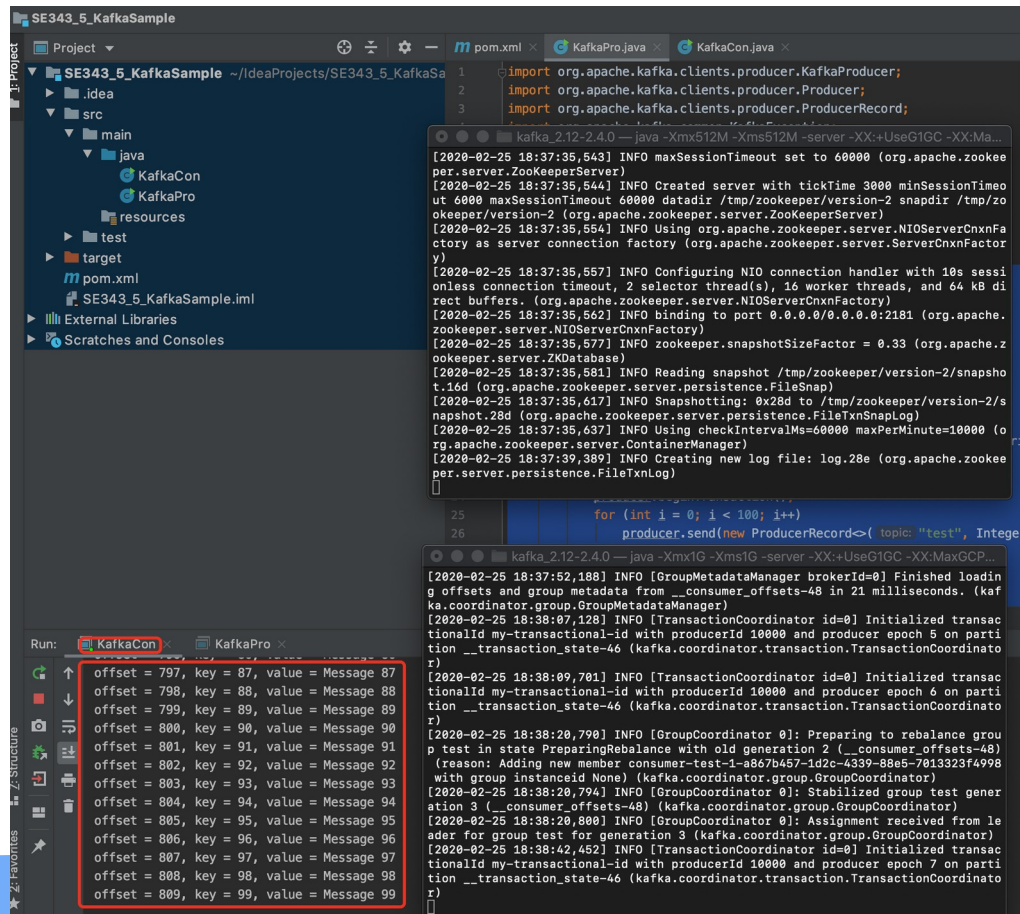
Apache Kafka

KafkaPro.java

```
public class KafkaPro {  
    public static void main(String[] args) throws Exception {  
        Properties props = new Properties();  
        props.put("bootstrap.servers", "localhost:9092");  
        props.setProperty("transactional.id", "my-transactional-id");  
  
        Producer<String, String> producer = null;  
  
        try {  
            producer = new KafkaProducer<String, String>(props, new StringSerializer(), new  
StringSerializer());  
            producer.initTransactions();  
  
            producer.beginTransaction();  
            for (int i = 0; i < 100; i++)  
                producer.send(new ProducerRecord<>("test", Integer.toString(i), "Message " +  
Integer.toString(i)));  
            producer.commitTransaction();  
        } catch (ProducerFencedException e) {  
            producer.close();  
        } catch (OutOfOrderSequenceException e) {  
            producer.close();  
        } catch (AuthorizationException e) {  
            producer.close();  
        } catch (KafkaException e) {  
            producer.abortTransaction();  
        }  
        producer.close();  
    }  
}
```

KafkaCon.java

```
public class KafkaCon {  
    public static void main(String[] args) throws Exception {  
        Properties props = new Properties();  
        props.setProperty("bootstrap.servers", "localhost:9092");  
        props.setProperty("group.id", "test");  
        props.setProperty("enable.auto.commit", "true");  
        props.setProperty("auto.commit.interval.ms", "1000");  
        props.setProperty("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");  
        props.setProperty("value.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");  
  
        KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props);  
        consumer.subscribe(Arrays.asList("test"));  
  
        while (true) {  
            ConsumerRecords<String, String> records = consumer.poll(Duration.ofMillis(100));  
            for (ConsumerRecord<String, String> record : records)  
                System.out.printf("offset = %d, key = %s, value = %s%n", record.offset(), record.key(),  
record.value());  
        }  
    }  
}
```



The screenshot displays an IDE environment with a project named `SE343_5_KafkaSample`. The project structure includes `src/main/java` with `KafkaCon` and `KafkaPro` classes, and `resources`. The `KafkaPro.java` file is open, showing imports for `org.apache.kafka.clients.producer.KafkaProducer`, `org.apache.kafka.clients.producer.Producer`, and `org.apache.kafka.clients.producer.ProducerRecord`. The `KafkaCon` class is also visible, containing a `main` method that runs a loop to send messages to a topic named "test".

The `Run` tab shows the execution of `KafkaCon`, displaying a list of messages sent to the topic "test". The messages are numbered 87 to 99, with keys ranging from 87 to 99 and values being "Message 87" to "Message 99".

The `Console` tab shows the logs of the `KafkaPro` application. The logs indicate that the application is running on `kafka_2.12-2.4.0` with `-Xmx512M` and `-Xms512M` heap sizes. The logs show the initialization of the `KafkaProducer` and the sending of messages to the topic "test".

```
import org.apache.kafka.clients.producer.KafkaProducer;
import org.apache.kafka.clients.producer.Producer;
import org.apache.kafka.clients.producer.ProducerRecord;

// ...

for (int i = 0; i < 100; i++)
    producer.send(new ProducerRecord<>("test", Integer.parseInt(i)));
```

Run: **KafkaCon** | KafkaPro

- offset = 797, key = 87, value = Message 87
- offset = 798, key = 88, value = Message 88
- offset = 799, key = 89, value = Message 89
- offset = 800, key = 90, value = Message 90
- offset = 801, key = 91, value = Message 91
- offset = 802, key = 92, value = Message 92
- offset = 803, key = 93, value = Message 93
- offset = 804, key = 94, value = Message 94
- offset = 805, key = 95, value = Message 95
- offset = 806, key = 96, value = Message 96
- offset = 807, key = 97, value = Message 97
- offset = 808, key = 98, value = Message 98
- offset = 809, key = 99, value = Message 99

Console: `kafka_2.12-2.4.0` — `java -Xmx512M -Xms512M -server -XX:+UseG1GC -XX:MaxGC...`

```
[2020-02-25 18:37:35,543] INFO maxSessionTimeout set to 60000 (org.apache.zookeeper.server.ZooKeeperServer)
[2020-02-25 18:37:35,544] INFO Created server with tickTime 3000 minSessionTimeout 60000 maxSessionTimeout 60000 dataDir /tmp/zookeeper/version-2 snapdir /tmp/zookeeper/version-2 (org.apache.zookeeper.server.ZooKeeperServer)
[2020-02-25 18:37:35,554] INFO Using org.apache.zookeeper.server.NIOServerCnxnFactory as server connection factory (org.apache.zookeeper.server.ServerCnxnFactory)
[2020-02-25 18:37:35,557] INFO Configuring NIO connection handler with 10s sessionless connection timeout, 2 selector thread(s), 16 worker threads, and 64 kB direct buffers. (org.apache.zookeeper.server.NIOServerCnxnFactory)
[2020-02-25 18:37:35,562] INFO binding to port 0.0.0.0/0.0.0.0:2181 (org.apache.zookeeper.server.NIOServerCnxnFactory)
[2020-02-25 18:37:35,577] INFO zookeeper.snapshotSizeFactor = 0.33 (org.apache.zookeeper.server.ZooKeeperServer)
[2020-02-25 18:37:35,581] INFO Reading snapshot /tmp/zookeeper/version-2/snapshot.16d (org.apache.zookeeper.server.persistence.FileSnap)
[2020-02-25 18:37:35,617] INFO Snapshotting: 0x28d to /tmp/zookeeper/version-2/snapshot.28d (org.apache.zookeeper.server.persistence.FileTxnSnapLog)
[2020-02-25 18:37:35,637] INFO Using checkIntervalMs=60000 maxPerMinute=10000 (org.apache.zookeeper.server.ContainerManager)
[2020-02-25 18:37:39,389] INFO Creating new log file: log.28e (org.apache.zookeeper.server.persistence.FileTxnLog)

[2020-02-25 18:37:52,188] INFO [GroupMetadataManager brokerId=0] Finished loading offsets and group metadata from __consumer_offsets-48 in 21 milliseconds. (kafka.coordinator.group.GroupMetadataManager)
[2020-02-25 18:38:07,128] INFO [TransactionCoordinator id=0] Initialized transactionalId my-transactional-id with producerId 10000 and producer epoch 5 on partition __transaction_state-46 (kafka.coordinator.transaction.TransactionCoordinator)
[2020-02-25 18:38:09,701] INFO [TransactionCoordinator id=0] Initialized transactionalId my-transactional-id with producerId 10000 and producer epoch 6 on partition __transaction_state-46 (kafka.coordinator.transaction.TransactionCoordinator)
[2020-02-25 18:38:20,790] INFO [GroupCoordinator 0]: Preparing to rebalance group p test in state PreparingRebalance with old generation 2 (__consumer_offsets-48) (reason: Adding new member consumer-test-1-a867b457-1d2c-4339-88e5-7013323f4998 with group instanceId None) (kafka.coordinator.group.GroupCoordinator)
[2020-02-25 18:38:20,794] INFO [GroupCoordinator 0]: Stabilized group test generation 3 (__consumer_offsets-48) (kafka.coordinator.group.GroupCoordinator)
[2020-02-25 18:38:20,800] INFO [GroupCoordinator 0]: Assignment received from leader for group test for generation 3 (kafka.coordinator.group.GroupCoordinator)
[2020-02-25 18:38:42,452] INFO [TransactionCoordinator id=0] Initialized transactionalId my-transactional-id with producerId 10000 and producer epoch 7 on partition __transaction_state-46 (kafka.coordinator.transaction.TransactionCoordinator)
```


- Consumer

`@SpringBootApplication`

```
public class ConsumerApplication {
```

```
    public static void main(String[] args) {  
        SpringApplication.run(ConsumerApplication.class, args);  
    }
```

`@Bean`

```
public NewTopic topic() {  
    return TopicBuilder.name("topic1")  
        .partitions(10)  
        .replicas(1)  
        .build();  
}
```

`@KafkaListener(id = "myId", topics = "topic1")`

```
public void listen(String in) {  
    System.out.println(in);  
}  
}
```

- Producer

`@SpringBootApplication`

```
public class KafkaProducerApplication {
```

```
    public static void main(String[] args) {
```

```
        SpringApplication.run(ProducerApplication.class, args);
    }
```

`@Bean`

```
public NewTopic topic() {
```

```
    return TopicBuilder.name("topic1")
```

```
        .partitions(10)
```

```
        .replicas(1)
```

```
        .build();
```

```
}
```

`@Bean`

```
public ApplicationRunner runner(KafkaTemplate<String, String> template) {
```

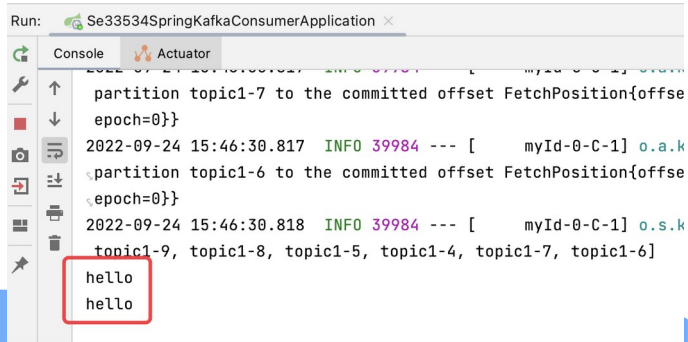
```
    return args -> {
```

```
        template.send("topic1", "hello");
```

```
    };
```

```
}
```

```
}
```



- BankController

@RestController

```
public class BankController {
```

@Autowired

```
private KafkaTemplate<String, String> kafkaTemplate;
```

@RequestMapping("/send")

```
public void send() {
```

```
    String data = "Tom,Jerry,80";
```

```
    kafkaTemplate.send("topic1", "key", data);
```

```
    System.out.println(data);
```

```
}
```

```
}
```

- BankListener

@Component

```
public class BankListener {
```

@Autowired

```
private BankService bankService;
```

```
@KafkaListener(topics = "topic1", groupId = "group_topic_test")
```

```
public void topicListener(ConsumerRecord<String, String> record) {
```

```
    String[] value = record.value().split(",");
```

```
    bankService.transfer(value[0], value[1], Integer.valueOf(value[2]));
```

```
}
```

```
}
```

- SQL Script

```
DROP TABLE IF EXISTS `bank`;  
CREATE TABLE `bank` (  
  `account` varchar(45) NOT NULL,  
  `balance` int DEFAULT NULL,  
  PRIMARY KEY (`account`)  
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci;
```

- Apache Kafka Quick Start
 - <https://kafka.apache.org/quickstart>
- Spring for Apache Kafka
 - <https://docs.spring.io/spring-kafka/docs/current/reference/html/#spring-boot-consumer-app>
- 如何轻松在 SpringBoot 中正确配置并运行 Kafka
 - https://blog.csdn.net/Eternal_Blue/article/details/125293622
- 详细解决 zookeeper 启动占用 8080 端口方法
 - <https://blog.csdn.net/tianynnb/article/details/128617174>



Thank You!