

Architecture of Enterprise Applications 20

Neo4J & Graph Computing

Haopeng Chen

REliable, ***IN***telligent and ***SC***alable Systems Group (***REINS***)

Shanghai Jiao Tong University

Shanghai, China

<http://reins.se.sjtu.edu.cn/~chenhp>

e-mail: chen-hp@sjtu.edu.cn

- Contents

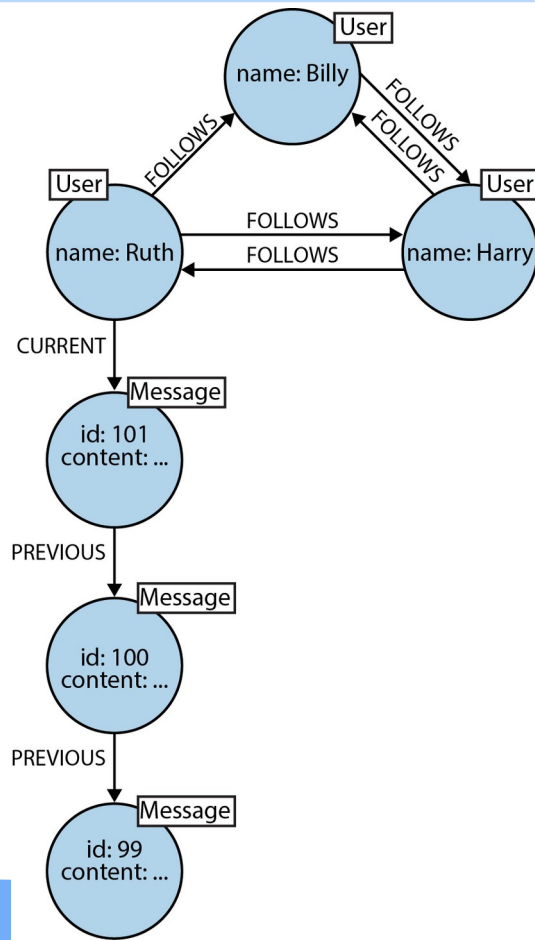
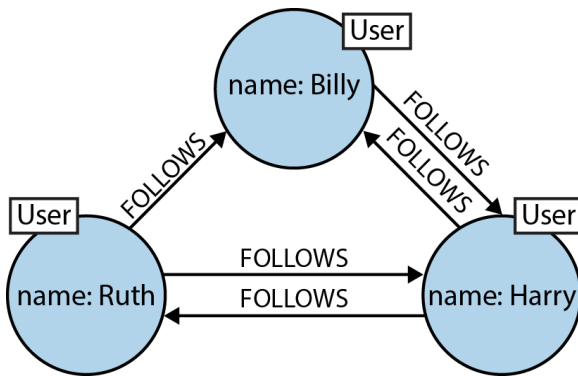
- Neo4J
- Graph Database
- Options for Storing Connected Data
- Data Modeling with Graphs
- Building a Graph Database Application
- Graph Database Internals

- Objectives

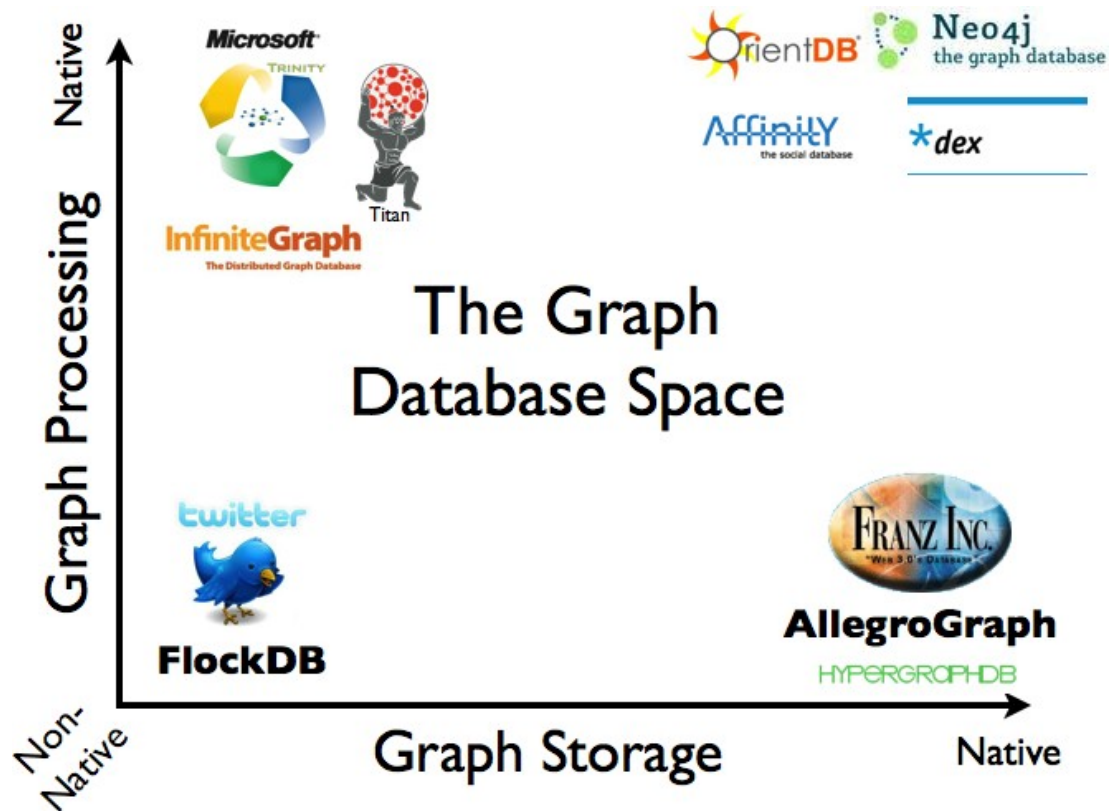
- 能够根据数据特性和数据访问模式，识别适合图数据库存储的数据，设计并实现其在图数据库中的存储和访问方案

What is a Graph?

- Formally, a graph is just a collection of *vertices* and *edges*
 - or, in less intimidating language, a set of *nodes* and the *relationships* that connect them.
 - Graphs represent entities as nodes and the ways in which those entities relate to the world as relationships.



- A *graph database management system* (henceforth, a *graph database*) is
 - an online database management system with Create, Read, Update, and Delete (CRUD) methods that expose a graph data model.
 - Graph databases are generally built for use with transactional (OLTP) systems.
 - Accordingly, they are normally optimized for transactional performance, and engineered with transactional integrity and operational availability in mind.
- There are two properties of graph databases we should consider when investigating graph database technologies:
 - *The underlying storage*
 - *The processing engine*



- **Relational Databases Lack Relationships**


User					
UserID	User	Address	Phone	Email	Alternate
1	Alice	123 Foo St.	12345678	alice@example.org	alice@neo4j.org
2	Bob	456 Bar Ave.		bob@example.org	
...
99	Zach	99 South St.		zach@example.org	

Order	
OrderID	UserID
1234	1
5678	1
...	...
5588	99

LineItem		
OrderID	ProductID	Quantity
1234	765	2
1234	987	1
...
5588	765	1

Product		
ProductID	Description	Handling
321	strawberry ice cream	freezer
765	potatoes	
...	...	
987	dried spaghetti	

- **Relational Databases Lack Relationships**



Example 3. Alice's friends-of-friends

```
SELECT p1.Person AS PERSON,  
       p2.Person AS FRIEND_OF_FRIEND  
FROM PersonFriend pf1 JOIN Person p1  
    ON pf1.PersonID = p1.ID  
JOIN PersonFriend pf2  
    ON pf2.PersonID = pf1.FriendID  
JOIN Person p2  
    ON pf2.FriendID = p2.ID  
WHERE p1.Person = 'Alice' AND pf2.FriendID <> p1.ID
```

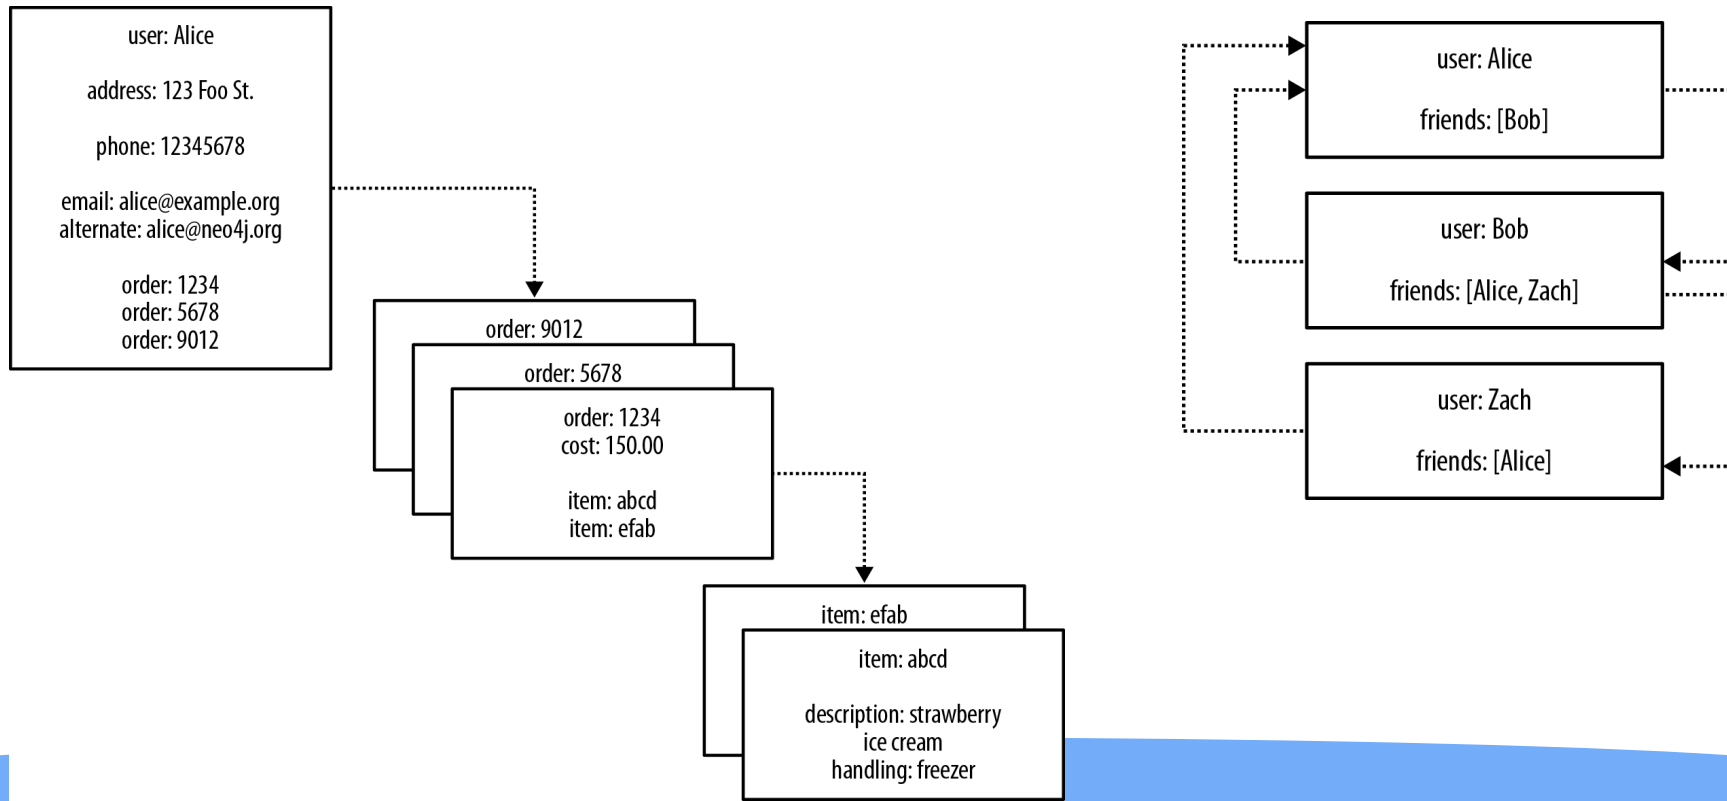
Example 1. Bob's friends

```
SELECT p1.Person  
FROM Person p1 JOIN PersonFriend  
    ON PersonFriend.FriendID = p1.ID  
JOIN Person p2  
    ON PersonFriend.PersonID = p2.ID  
WHERE p2.Person = 'Bob'
```

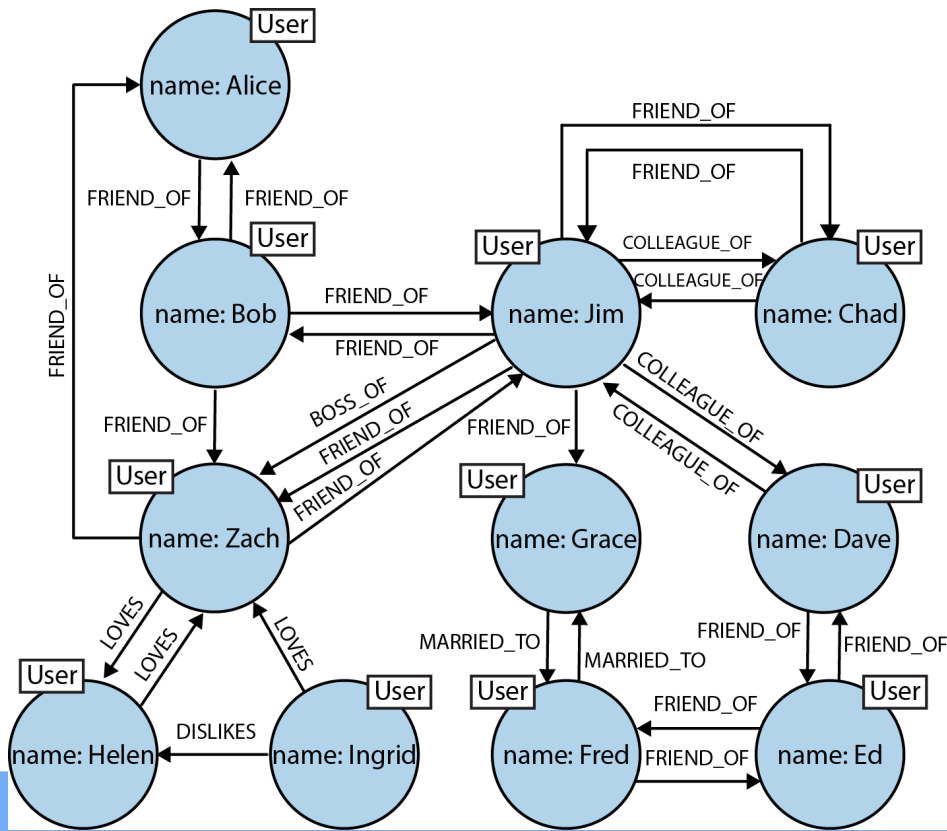
Example 2-2. Who is friends with Bob?

```
SELECT p1.Person  
FROM Person p1 JOIN PersonFriend  
    ON PersonFriend.PersonID = p1.ID  
JOIN Person p2  
    ON PersonFriend.FriendID = p2.ID  
WHERE p2.Person = 'Bob'
```

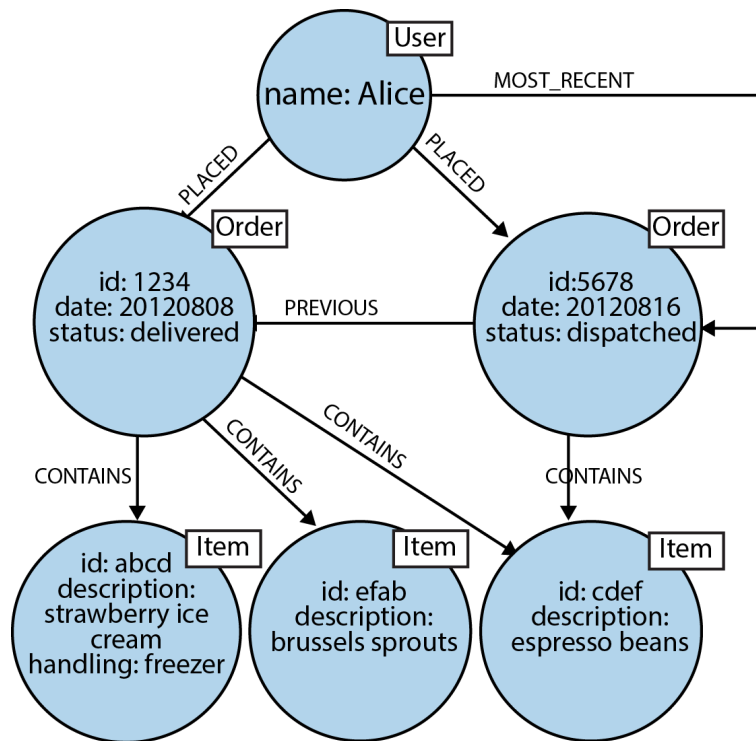
- NOSQL Databases Also Lack Relationships**



- Graph Databases Embrace Relationships**



- Graph Databases Embrace Relationships**




- **The Labeled Property Graph Model**

- A *labeled property graph* is made up of *nodes*, *relationships*, *properties*, and *labels*.
 - Nodes contain properties. Think of nodes as documents that store properties in the form of arbitrary key-value pairs. In Neo4j, the keys are strings and the values are the Java string and primitive data types, plus arrays of these types.
 - Nodes can be tagged with one or more labels. Labels group nodes together, and indicate the roles they play within the dataset.
 - Relationships connect nodes and structure the graph. A relationship always has a direction, a single name, and a *start node* and an *end node*—there are no dangling relationships. Together, a relationship's direction and name add semantic clarity to the structuring of nodes.

- **Querying Graphs: Cypher**

- Cypher is an expressive (yet compact) graph database **query** language.




- Here's the equivalent ASCII art representation in Cypher:

```
(emil:Person {name:'Emil'}) <-[:KNOWS]-(jim:Person {name:'Jim'}) -[:KNOWS]->(ian:Person {name:'Ian'}) -[:KNOWS]->(emil)
```

- **Querying Graphs: Cypher**

- Cypher is an expressive (yet compact) graph database **query** language.



- Match:

MATCH (a:Person)-[:KNOWS]->(b)-[:KNOWS]->(c), (a)-[:KNOWS]->(c)

WHERE a.name = 'Jim'

RETURN b, c

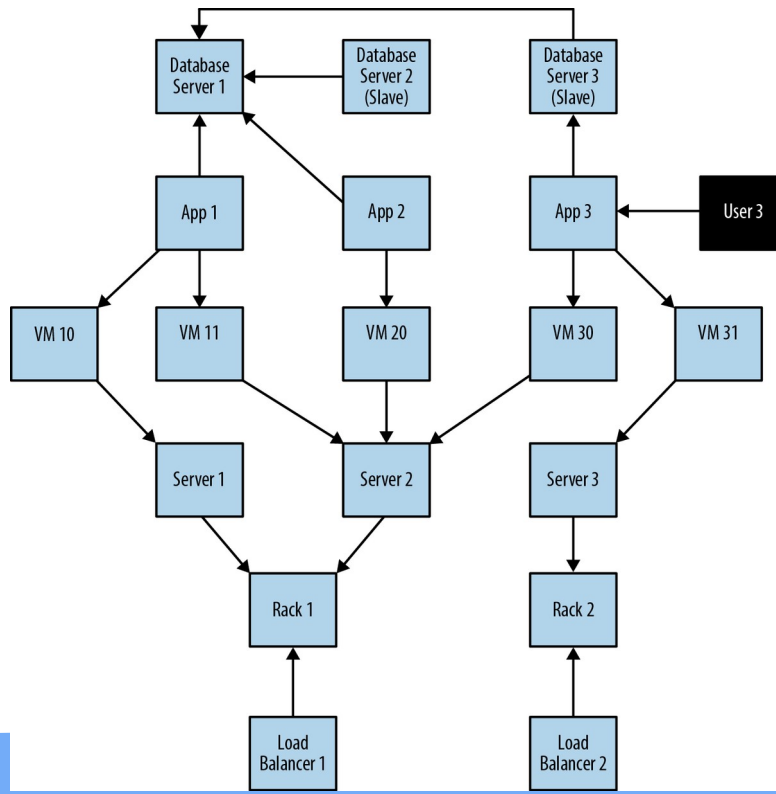
- **Querying Graphs: Cypher**

- The other clauses we can use in a Cypher query include:

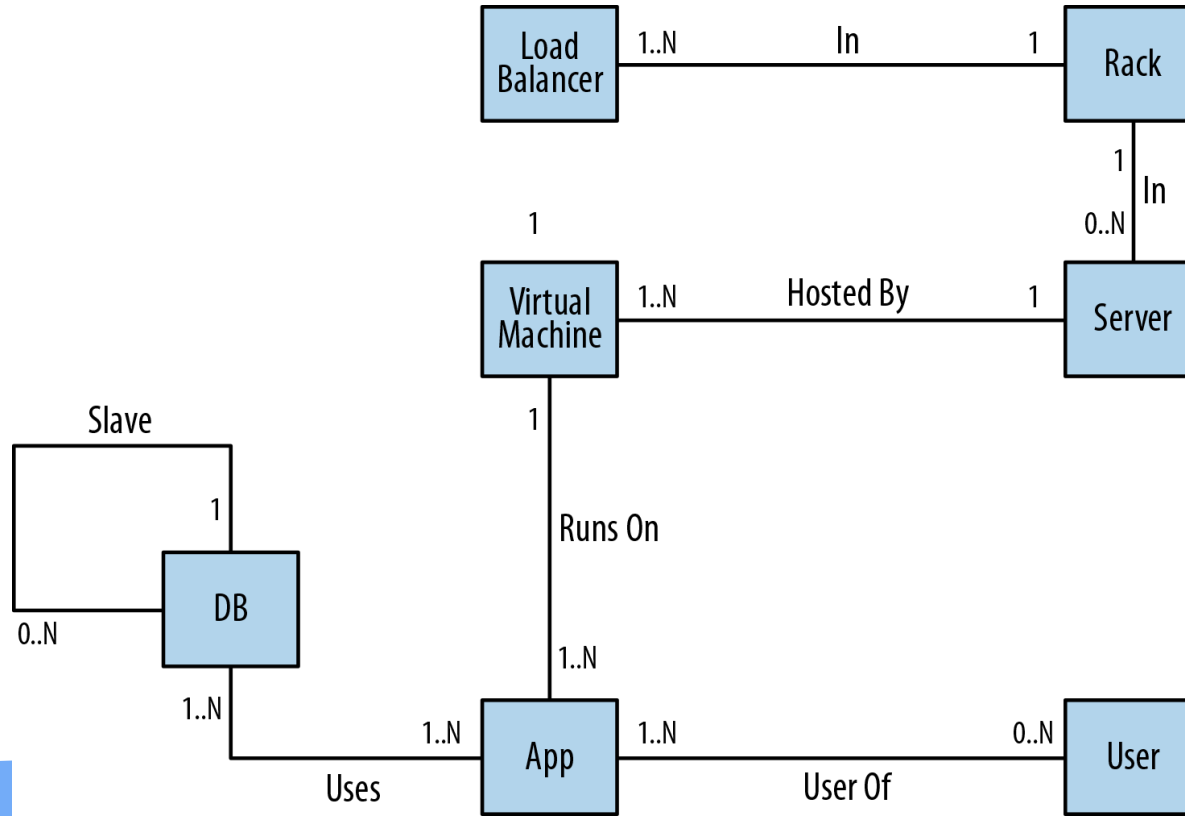
- WHERE
 - Provides criteria for filtering pattern matching results.
- CREATE and CREATE UNIQUE
 - Create nodes and relationships.
- MERGE
 - Ensures that the supplied pattern exists in the graph, either by reusing existing nodes and relationships that match the supplied predicates, or by creating new nodes and relationships.
- DELETE
 - Removes nodes, relationships, and properties.
- SET
 - Sets property values
- FOREACH
 - Performs an updating action for each element in a list.
- UNION
 - Merges results from two or more queries.
- WITH
 - Chains subsequent query parts and forwards results from one to the next. Similar to piping commands in Unix.
- START
 - Specifies one or more explicit starting points—nodes or relationships—in the graph.

- **A Comparison of Relational and Graph Modeling**

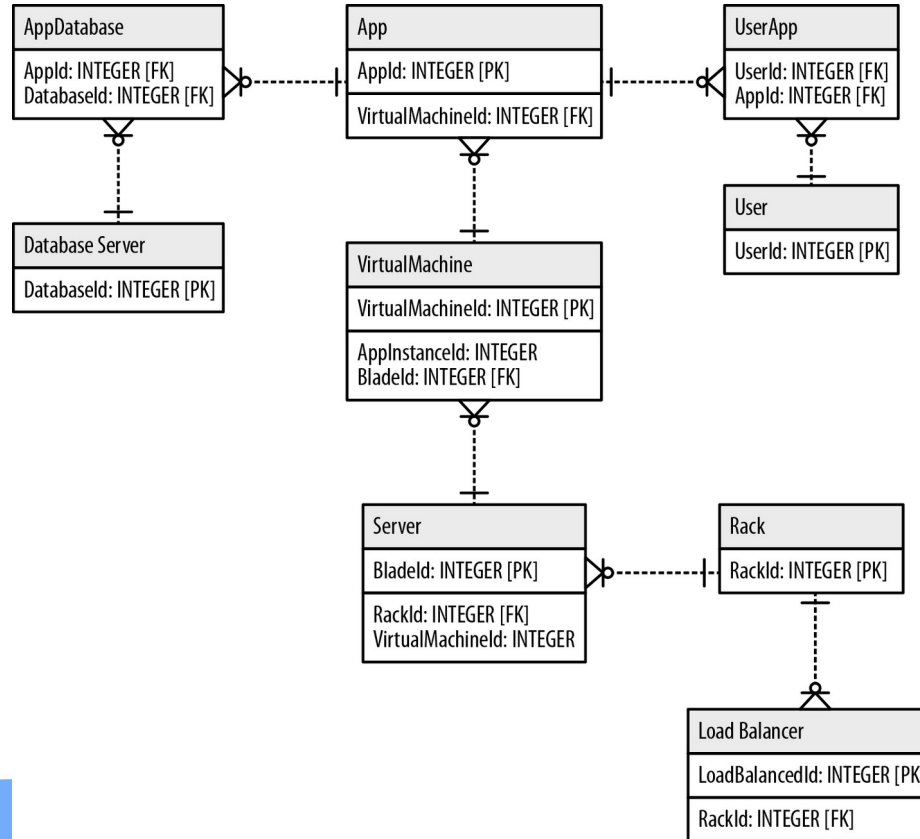
- To facilitate this comparison, we'll examine a simple data center management domain.



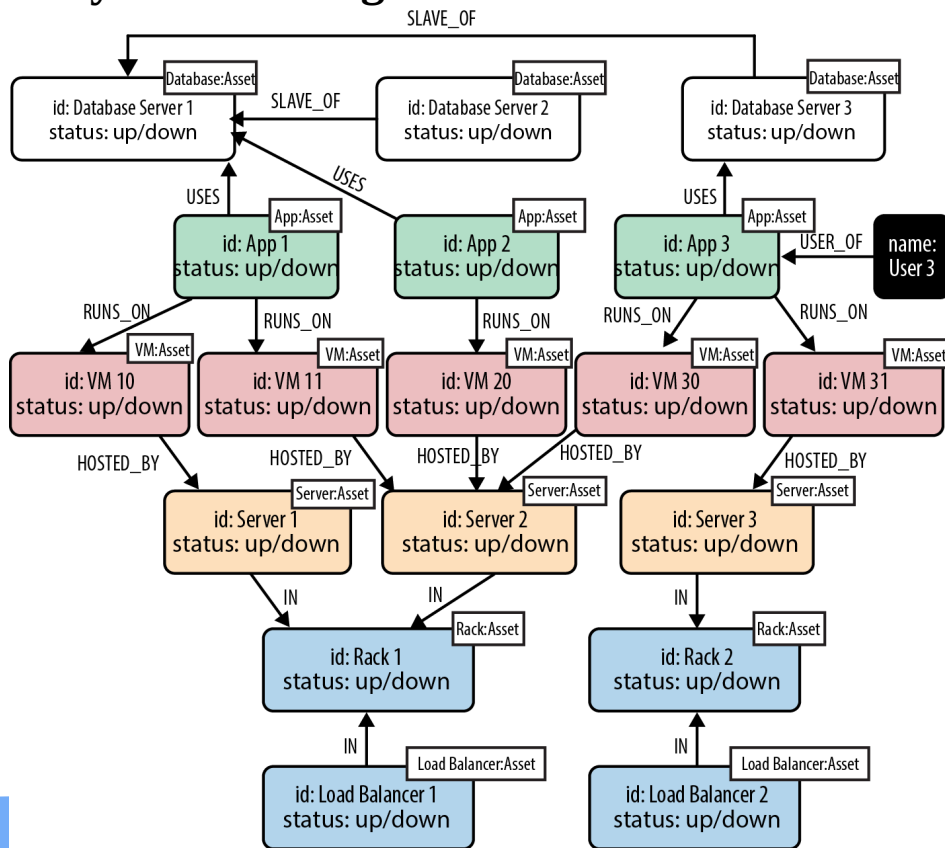
- Relational Modeling in a Systems Management Domain**



- Relational Modeling in a Systems Management Domain**




- Graph Modeling in a Systems Management Domain**



- **Test Graph Model**

- When a user reports a problem, we can limit the physical fault-finding to problematic network elements between the user and the application and the application and its dependencies.
- In our graph we can find the faulty equipment with the following query:
 - **MATCH** (user:User)-[*1..5]-(asset:Asset)
WHERE user.name = 'User 3' **AND** asset.status = 'down' **RETURN DISTINCT** asset
- This allows us to match paths such as:
 - (user)-[:USER_OF]->(app)
 - (user)-[:USER_OF]->(app)-[:USES]->(database)
 - (user)-[:USER_OF]->(app)-[:USES]->(database)-[:SLAVE_OF]->(another-database)
 - (user)-[:USER_OF]->(app)-[:RUNS_ON]->(vm)
 - (user)-[:USER_OF]->(app)-[:RUNS_ON]->(vm)-[:HOSTED_BY]->(server)
 - (user)-[:USER_OF]->(app)-[:RUNS_ON]->(vm)-[:HOSTED_BY]->(server)-[:IN]->(rack)
 - (user)-[:USER_OF]->(app)-[:RUNS_ON]->(vm)-[:HOSTED_BY]->(server)-[:IN]->(rack) <-[:IN]-(load-balancer)

- **Describe the Model in Terms of the Application's Needs**
- Here's an example of a user story for a book review web application:
 - **AS A** reader who likes a book, **I WANT** to know which books other readers who like the same book have liked, **SO THAT** I can find other books to read.



- since *Alice likes Dune*, find books that others who like Dune have enjoyed:
MATCH (:Reader {name:'Alice'})-[:LIKES]->(:Book {title:'Dune'})<-[:LIKES]-(:Reader)-[:LIKES]->(books:Book)
RETURN books.title

- **Nodes for Things, Relationships for Structure**

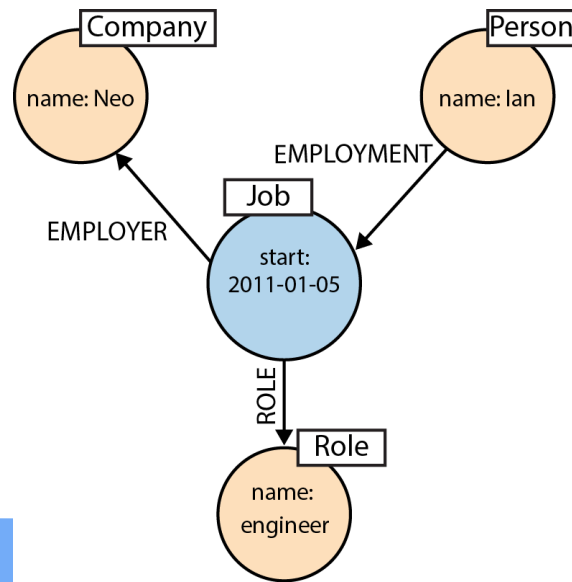
- Use nodes to represent entities—that is, the *things* in our domain that are of interest to us, and which can be labeled and grouped.
- Use relationships both to express the *connections* between entities and to establish semantic context for each entity, thereby structuring the domain.
- Use relationship direction to further clarify relationship semantics. Many relationships are asymmetrical, which is why relationships in a property graph are always directed. For bidirectional relationships, we should make our queries ignore direction, rather than using two relationships.
- Use node properties to represent entity *attributes*, plus any necessary entity metadata, such as timestamps, version numbers, etc.
- Use relationship properties to express the strength, weight, or quality of a relationship, plus any necessary relationship metadata, such as timestamps, version numbers, etc.

- **Model Facts as Nodes**

- **Employment**

- The Figure shows how the fact of Ian being employed by Neo Technology in the role of engineer can be represented in the graph.

```
CREATE (:Person {name:'Ian'})-[:EMPLOYMENT]-> (employment:Job {start_date:'2011-01-05'}) -  
[:EMPLOYER]->(:Company {name:'Neo'}),  
(employment)-[:ROLE]->(:Role {name:'engineer'})
```

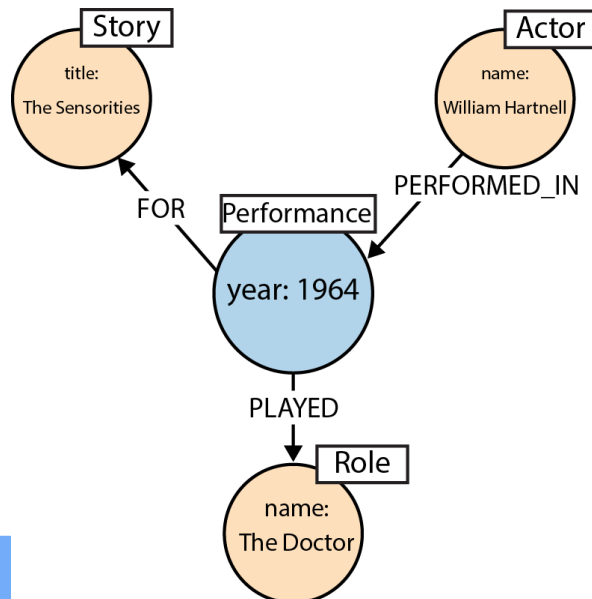


- **Model Facts as Nodes**

- **Performance**

- The Figure shows how the fact that William Hartnell played The Doctor in the story *The Sensorities* can be represented in the graph.

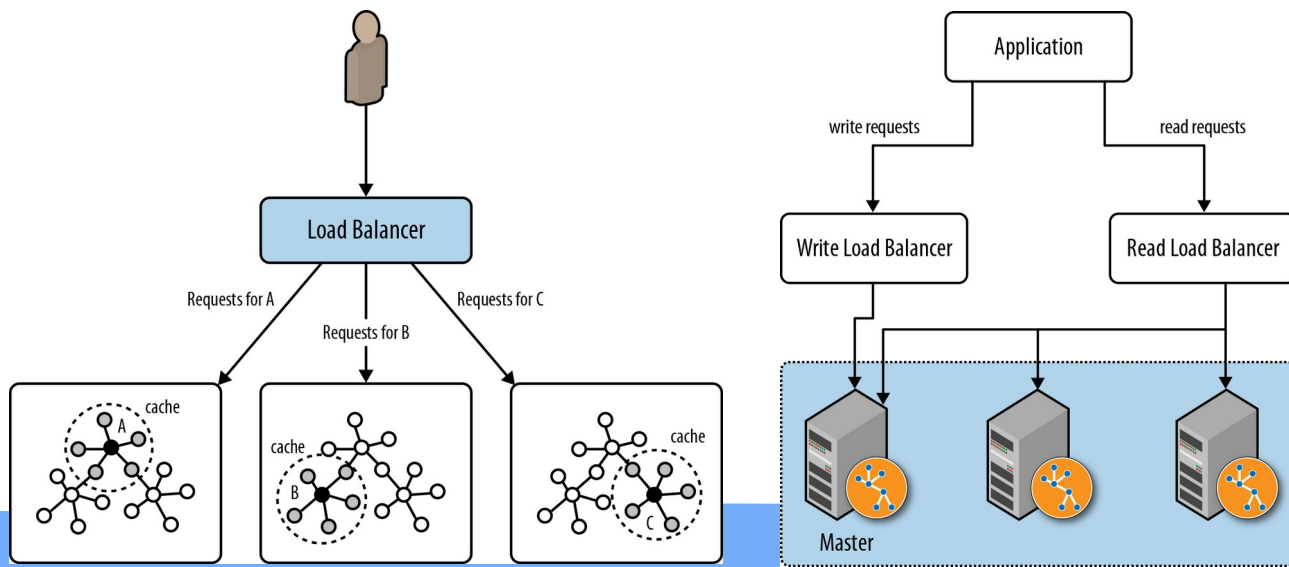
```
CREATE (:Actor {name:'William Hartnell'})-[:PERFORMED_IN]-> (performance:Performance {year:1964})-[:PLAYED]-> (:Role {name:'The Doctor'}),  
(performance)-[:FOR]->(:Story {title:'The Sensorities'})
```



- **Embedded versus Server**

- Most databases today run as a server that is accessed through a client library. Neo4j is somewhat unusual in that it can be run in embedded as well as server mode.
- In embedded mode, Neo4j runs in the same process as our application.
- Running Neo4j in server mode is the most common means of deploying the database today.

- **Cluster**




- Neo4j stores graph data in a number of different *store files*.
 - Each store file contains the data for a specific part of the graph (e.g., there are separate stores for nodes, relation- ships, labels, and properties).


Node (15 bytes)



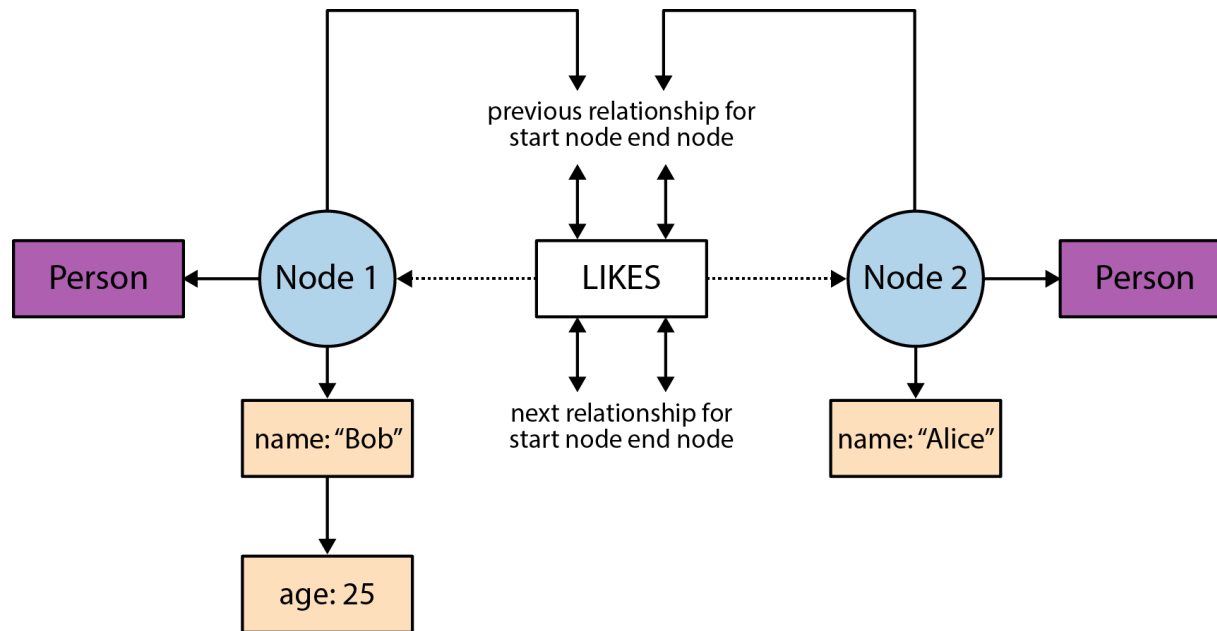
inUse




Relationship (34 bytes)




- Neo4j stores graph data in a number of different *store files*.
 - Doubly Linked Lists in the Relationship Store




- Neo4j stores graph data



- Neo4j stores graph data
 - Node

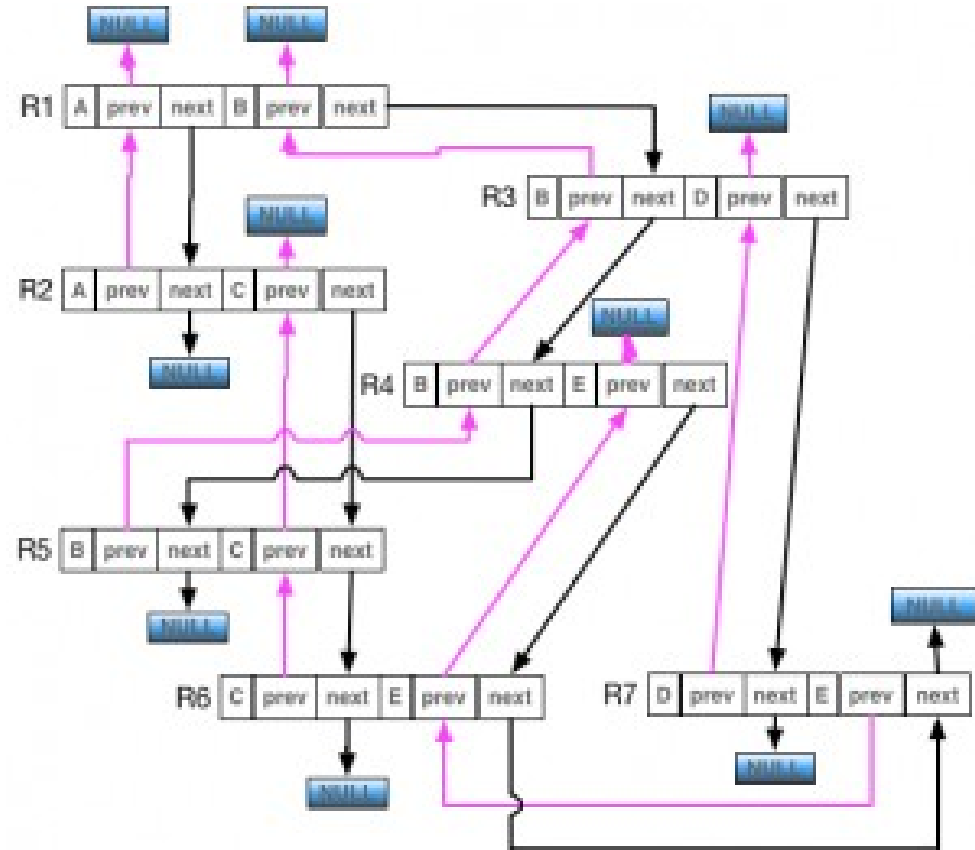



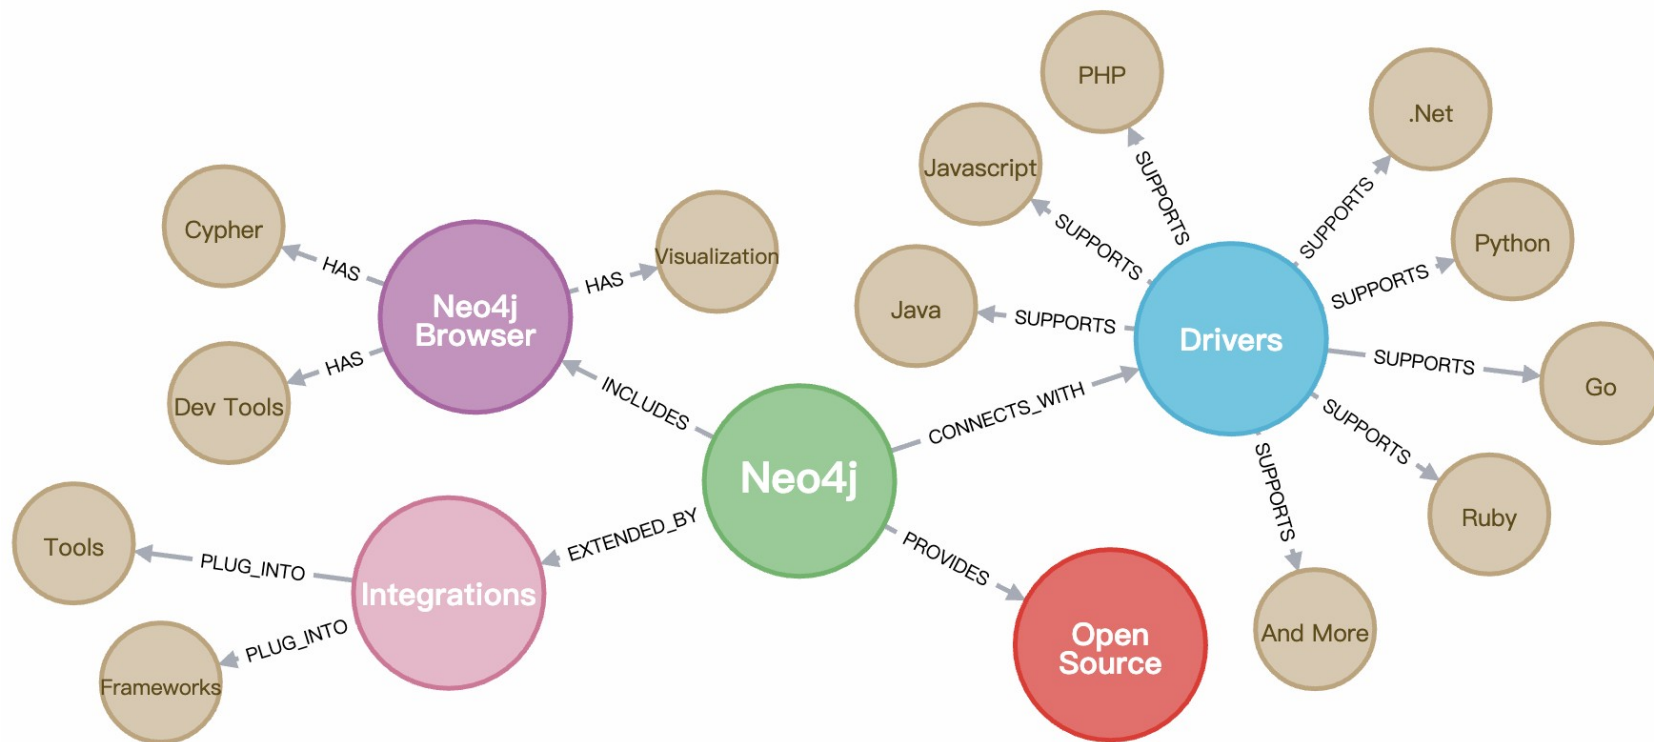
inUse nextRelId nextPropId



A:	1	R1	P1
B:	1	R1	P3
C:	1	R2	P5
D:	1	R3	P7
E:	1	R4	P9

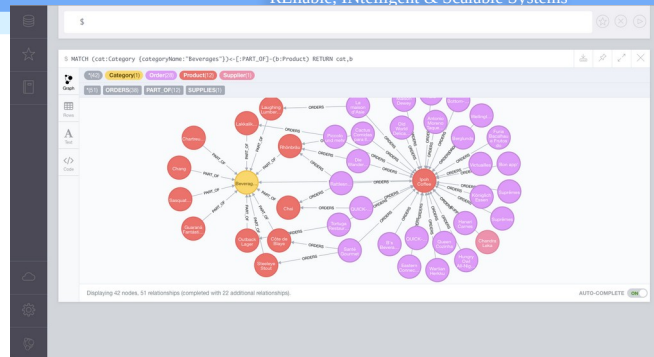
- Neo4j stores graph data
 - relationship





Install and Run Neo4j

1. Open up your terminal/shell.
2. Extract the contents of the archive, using:
`tar -xf <filecode>.`
For example,
`tar -xf neo4j-community-4.3.3-unix.tar.gz`



3. Place the extracted files in a permanent home on your server. The top level directory is referred to as NEO4J_HOME.
 - To run Neo4j as a console application, use:
`<NEO4J_HOME>/bin/neo4j console`
 - To run Neo4j in a background process, use:
`<NEO4J_HOME>/bin/neo4j start`
4. Visit <http://localhost:7474> in your web browser.
5. Connect using the username 'neo4j' with default password 'neo4j'. You'll then be prompted to change the password.

- pom.xml

```
<dependencies>
  <dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-data-neo4j</artifactId>
  </dependency>

  <dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-test</artifactId>
    <scope>test</scope>
    <exclusions>
      <exclusion>
        <groupId>org.junit.vintage</groupId>
        <artifactId>junit-vintage-engine</artifactId>
      </exclusion>
    </exclusions>
  </dependency>
</dependencies>
```


- Person.java

@NodeEntity

```
public class Person {
```

```
    @Id @GeneratedValue
```

```
    private Long id;
```

```
    private String name;
```

```
    private Person() {
```

```
        // Empty constructor required as of Neo4j API 2.0.5
```

```
    };
```

```
    public Person(String name) {
```

```
        this.name = name;
```

```
    }
```

```
    @Relationship(type = "TEAMMATE", direction = Relationship.UNDIRECTED)
```

```
    public Set<Person> teammates;
```

```
    public void worksWith(Person person) {
```

```
        if (teammates == null) {
```

```
            teammates = new HashSet<>();
```

```
        }
```

```
        teammates.add(person);
```

```
    }
```

- Person.java

```
public String toString() {  
  
    return this.name + "'s teammates => "  
        + Optional.ofNullable(this.teammates).orElse(  
            Collections.emptySet()).stream()  
            .map(Person::getName)  
            .collect(Collectors.toList());  
}  
  
public String getName() {  
    return name;  
}  
  
public void setName(String name) {  
    this.name = name;  
}  
}
```

- PersonRepository.java

```
package com.example.accessingdataneoj;
```

```
import java.util.List;
```

```
import org.springframework.data.neo4j.repository.Neo4jRepository;
```

```
public interface PersonRepository extends Neo4jRepository<Person, Long> {
```

```
    Person findByName(String name);
```

```
    List<Person> findByTeammatesName(String name);
```

```
}
```

- AccessingDataNeo4jApplication.java

@SpringBootApplication

@EnableNeo4jRepositories

public class AccessingDataNeo4jApplication {

private final static Logger log = LoggerFactory.getLogger(AccessingDataNeo4jApplication.class);

public static void main(String[] args) throws Exception {

SpringApplication.run(AccessingDataNeo4jApplication.class, args);

System.exit(0);

}

@Bean

CommandLineRunner demo(PersonRepository personRepository) {

return args -> {

personRepository.deleteAll();

Person greg = new Person("Greg");

Person roy = new Person("Roy");

Person craig = new Person("Craig");

List<Person> team = Arrays.asList(greg, roy, craig);

- AccessingDataNeo4jApplication.java

```
List<Person> team = Arrays.asList(greg, roy, craig);
```

```
log.info("Before linking up with Neo4j...");
```

```
team.stream().forEach(person -> log.info("\t" + person.toString()));
```

```
personRepository.save(greg);
```

```
personRepository.save(roy);
```

```
personRepository.save(craig);
```

```
greg = personRepository.findByName(greg.getName());
```

```
greg.worksWith(roy);
```

```
greg.worksWith(craig);
```

```
personRepository.save(greg);
```

```
roy = personRepository.findByName(roy.getName());
```

```
roy.worksWith(craig);
```

```
// We already know that roy works with greg
```

```
personRepository.save(roy);
```

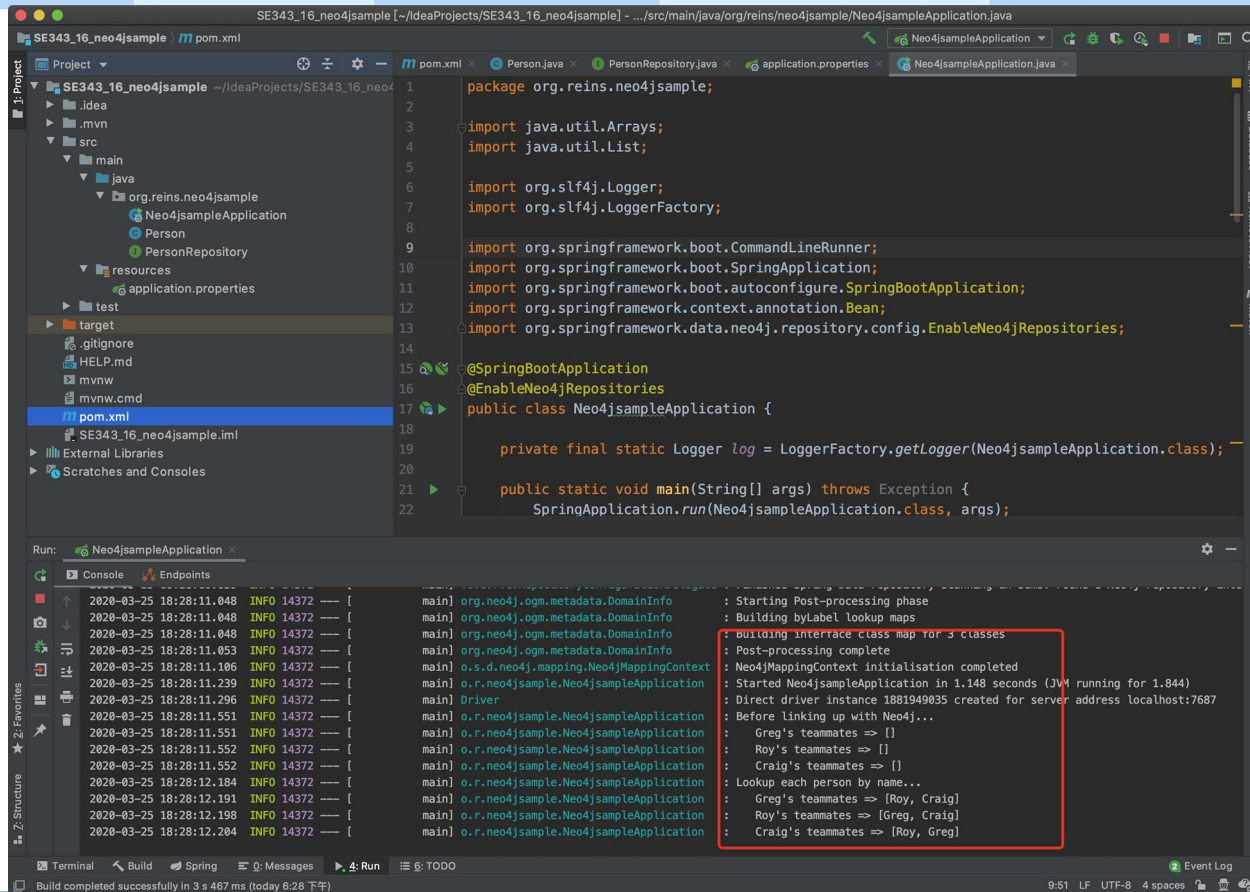
- AccessingDataNeo4jApplication.java

```
// We already know craig works with roy and greg
```

```
log.info("Lookup each person by name...");  
team.stream().forEach(person -> log.info(  
    "\t" + personRepository.findByName(person.getName()).toString()));
```

```
List<Person> teammates = personRepository.findByTeammatesName(greg.getName());  
log.info("The following have Greg as a teammate...");  
teammates.stream().forEach(person -> log.info("\t" + person.getName()));
```

```
};  
}  
  
}
```



```

package org.reins.neo4jsample;

import java.util.Arrays;
import java.util.List;

import org.slf4j.Logger;
import org.slf4j.LoggerFactory;

import org.springframework.boot.CommandLineRunner;
import org.springframework.boot.SpringApplication;
import org.springframework.boot.autoconfigure.SpringBootApplication;
import org.springframework.context.annotation.Bean;
import org.springframework.data.neo4j.repository.config.EnableNeo4jRepositories;

@SpringBootApplication
@EnableNeo4jRepositories
public class Neo4jsampleApplication {

    private final static Logger log = LoggerFactory.getLogger(Neo4jsampleApplication.class);

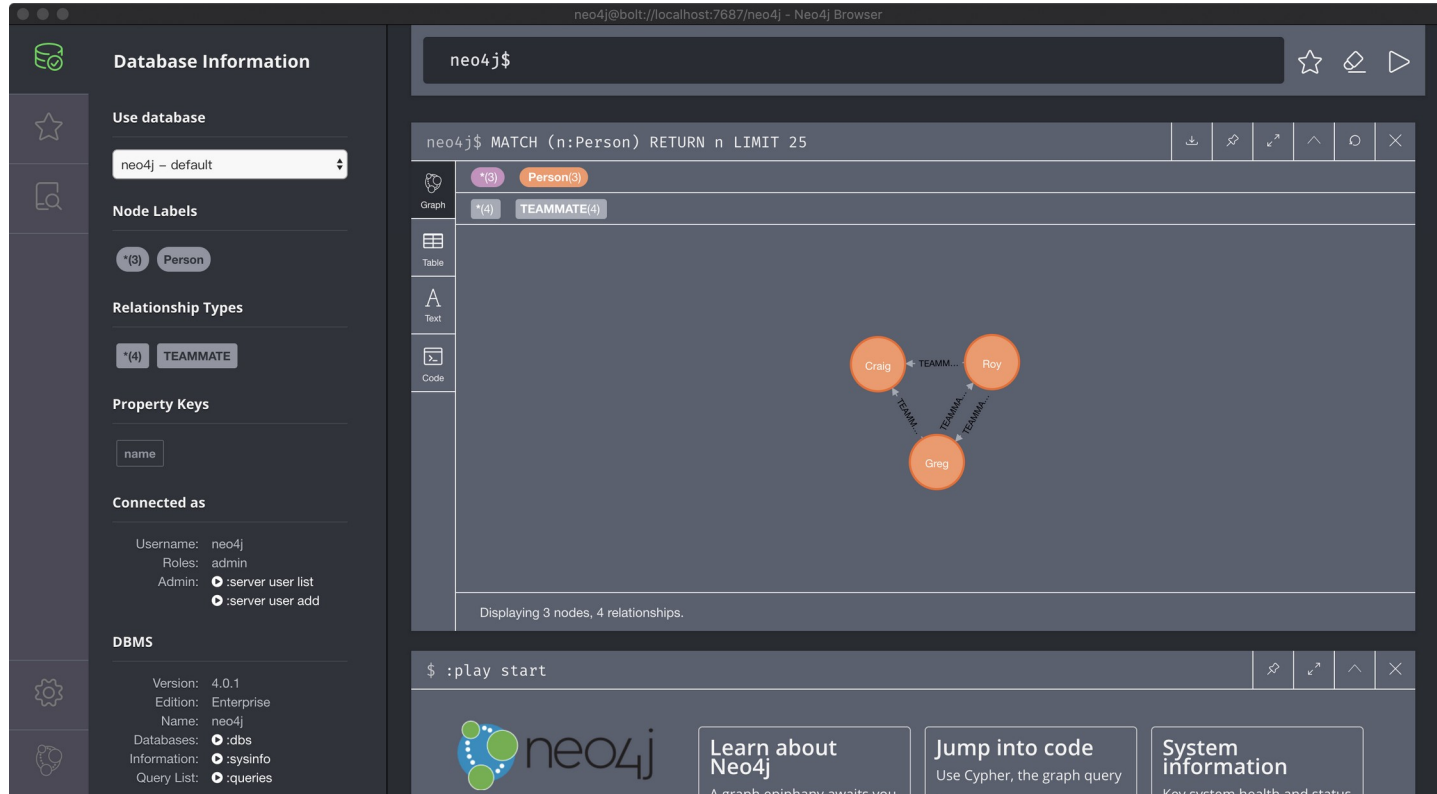
    public static void main(String[] args) throws Exception {
        SpringApplication.run(Neo4jsampleApplication.class, args);
    }
}

```

```

2020-03-25 18:28:11.048 INFO 14372 --- [main] org.neo4j.ogm.metadata.DomainInfo : Starting Post-processing phase
2020-03-25 18:28:11.048 INFO 14372 --- [main] org.neo4j.ogm.metadata.DomainInfo : Building byLabel lookup maps
2020-03-25 18:28:11.048 INFO 14372 --- [main] org.neo4j.ogm.metadata.DomainInfo : Building interface class map for 3 classes
2020-03-25 18:28:11.053 INFO 14372 --- [main] org.neo4j.ogm.metadata.DomainInfo : Post-processing complete
2020-03-25 18:28:11.106 INFO 14372 --- [main] o.s.d.neo4j.mapping.Neo4jMappingContext : Neo4jMappingContext initialisation completed
2020-03-25 18:28:11.239 INFO 14372 --- [main] o.r.neo4jsample.Neo4jsampleApplication : Started Neo4jsampleApplication in 1.148 seconds (JVM running for 1.844)
2020-03-25 18:28:11.296 INFO 14372 --- [main] Driver : Direct driver instance 1881949035 created for server address localhost:7687
2020-03-25 18:28:11.551 INFO 14372 --- [main] o.r.neo4jsample.Neo4jsampleApplication : Before linking up with Neo4j...
2020-03-25 18:28:11.551 INFO 14372 --- [main] o.r.neo4jsample.Neo4jsampleApplication : Greg's teammates => []
2020-03-25 18:28:11.552 INFO 14372 --- [main] o.r.neo4jsample.Neo4jsampleApplication : Roy's teammates => []
2020-03-25 18:28:11.552 INFO 14372 --- [main] o.r.neo4jsample.Neo4jsampleApplication : Craig's teammates => []
2020-03-25 18:28:11.552 INFO 14372 --- [main] o.r.neo4jsample.Neo4jsampleApplication : Lookup each person by name...
2020-03-25 18:28:12.191 INFO 14372 --- [main] o.r.neo4jsample.Neo4jsampleApplication : Greg's teammates => [Roy, Craig]
2020-03-25 18:28:12.198 INFO 14372 --- [main] o.r.neo4jsample.Neo4jsampleApplication : Roy's teammates => [Greg, Craig]
2020-03-25 18:28:12.204 INFO 14372 --- [main] o.r.neo4jsample.Neo4jsampleApplication : Craig's teammates => [Roy, Greg]

```



The screenshot displays the Neo4j Browser interface. On the left, the 'Database Information' sidebar shows details for the 'neo4j' database, including node labels (*3) Person, relationship types (*4) TEAMMATE, and property keys (name). The main area shows a Cypher query: `neo4j$ MATCH (n:Person) RETURN n LIMIT 25`. Below the query, a graph visualization displays three nodes: Craig, Roy, and Greg, connected by TEAMMATE relationships. The bottom status bar indicates 'Displaying 3 nodes, 4 relationships.' The footer includes the Neo4j logo and links for 'Learn about Neo4j', 'Jump into code', and 'System information'.

Database Information

Use database

neo4j – default

Node Labels

*3 Person

Relationship Types

*4 TEAMMATE

Property Keys

name

Connected as

Username: neo4j
Roles: admin
Admin: ☒ :server user list
☐ :server user add

DBMS

Version: 4.0.1
Edition: Enterprise
Name: neo4j
Databases: ☒ :dbs
Information: ☒ :sysinfo
Query List: ☒ :queries

neo4j\$

neo4j\$ MATCH (n:Person) RETURN n LIMIT 25

Graph

Table

Text

Code

Craig Roy Greg

TEAMMATE TEAMMATE TEAMMATE

Displaying 3 nodes, 4 relationships.

\$:play start

neo4j

Learn about Neo4j
A graph epiphany awaits you.

Jump into code
Use Cypher, the graph query

System information
Key system health and status

➤ 图数据上的分类问题


- 图粒度 vs. 节点粒度
- 节点特征 vs. 边特征

➤ 应用场景：

- 应用对象：社交网络、交易网络（边特征敏感）
- 应用任务：用户分类、异常行为检测…

➤ 图神经网络：

- 传统图神经网络未考虑边特征对于分类结果的影响
 - 谱图理论：Graph Convolutional Network (GCN)
 - 图注意力：Graph Attention Network (GAT)
- 仅有少数研究考虑了 GNN 在边特征上的拓展
 - 仅可处理标签化的边特征（R-GCN）
 - 仅将边特征视为权重进行处理（EGNN）



- 边特征的集合形式 \rightarrow 邻接形式

$$\vec{e}_p \rightarrow \vec{e}_{ij}$$

- 边映射矩阵 M_E

- 维度大小为 $N \times N \times M$
- 第三维度: one-hot 编码
- 可以预先构造

- 高维矩阵乘法




图 3-3 图上的边特征映射变换举例

➤ 节点特征的聚合（参与后续迭代）

$$\vec{h}'_i = \sigma(\sum_{j \in \mathcal{N}_i} \alpha_{ij} \vec{h}_j)$$

$$\sum_{j \in \mathcal{N}_i} \alpha_{ij} = 1$$

➤ 带有边特征的聚合（仅于融合层使用）

$$\vec{m}_i = \sigma(\sum_{j \in \mathcal{N}_i} \alpha_{ij} (\vec{h}_j \parallel \vec{e}_{ij}))$$




图 3-5 节点注意力模块中的节点特征聚合举例

- 带有边特征的注意力权重计算

$$w_{ij} = a(\vec{h}_i, \vec{h}_j, \vec{e}_{ij})$$

- 权重标准化 (softmax)

$$\alpha_{ij} = \text{softmax}_j(w_{ij}) = \frac{\exp(w_{ij})}{\sum_{k \in \mathcal{N}_i} \exp(w_{ik})}$$

- 展开

$$\alpha_{ij} = \frac{\exp(\text{LeakyReLU}(\vec{a}^T [\vec{h}_i \parallel \vec{h}_j \parallel \vec{e}_{ij}]))}{\sum_{k \in \mathcal{N}_i} \exp(\text{LeakyReLU}(\vec{a}^T [\vec{h}_i \parallel \vec{h}_k \parallel \vec{e}_{ik}]))}$$




图 3-4 标准化注意力权重计算方式举例

- 边特征迭代的必要性
 - 节点与边特征在迭代上的一致性
 - 模型的对称性

点边互换策略

- 反转图
- 边邻接矩阵、节点映射矩阵

$$\beta_{pq} = \frac{\exp(\text{LeakyReLU}(\vec{\mathbf{b}}^T [\vec{e}_p \| \vec{e}_q \| \vec{h}_{pq}]))}{\sum_{k \in \mathcal{N}_p} \exp(\text{LeakyReLU}(\vec{\mathbf{b}}^T [\vec{e}_p \| \vec{e}_k \| \vec{h}_{pk}]))}$$

$$\vec{e}'_p = \sigma(\sum_{q \in \mathcal{N}_p} \beta_{pq} \vec{e}_q)$$

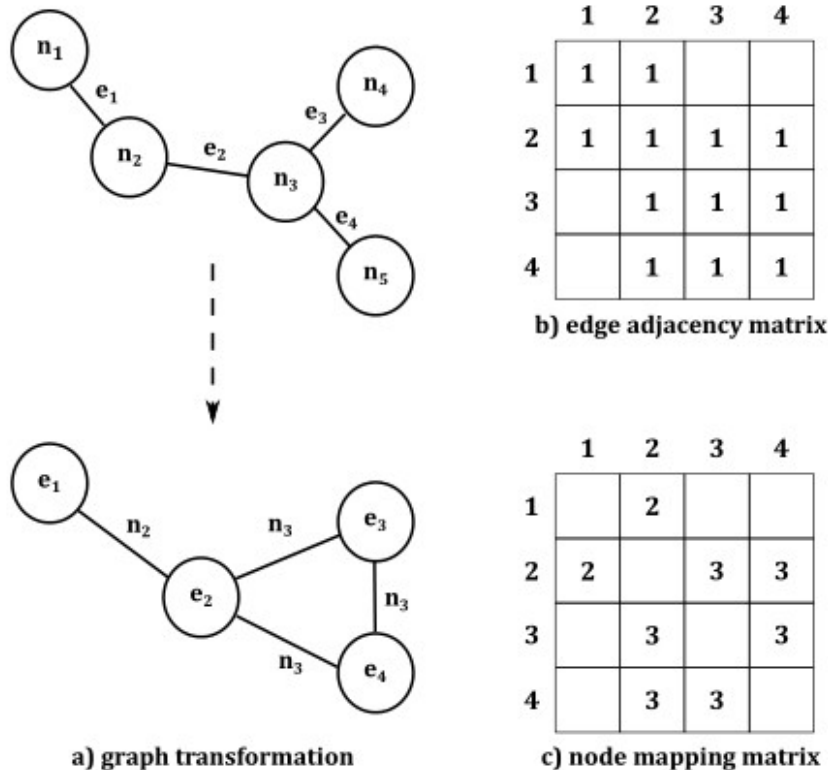


图 3-7 边特征迭代相关的图转换过程及相关矩阵图示

- 多尺度的特征融合策略

$$\vec{h}_i^* = \parallel_{l=1}^L m_i^l$$

- 粗粒度的 multi-head attention

$$\vec{h}_i^* = \parallel_{k=1}^K (\parallel_{l=1}^L m_i^{l,k})$$

- 维度压缩

$$\vec{h}_i^f = \text{LeakyReLU}(\mathbf{W}\vec{h}_i^*)$$

- softmax




图 3-8 EGAT 融合层及与其相关的输入输出实例

- 邻接矩阵 & 映射矩阵
 - 维度大小 ~ 图规模关系的非线性
 - 内存与显存的局限性 & 计算的复杂性

- 稀疏矩阵表示
 - 空元素数量 >> 非空元素数量
 - 高维映射矩阵的初次压缩

	0	1	2	3
0				1
1			1	
2		1		1
3	1		1	

0	1	2	2	3	3
---	---	---	---	---	---

3	2	1	3	0	2
---	---	---	---	---	---

1	1	1	1	1	1
---	---	---	---	---	---

Row Indices

Column Indices

Values

图 4-1 COO 格式下稀疏矩阵表示形式举例


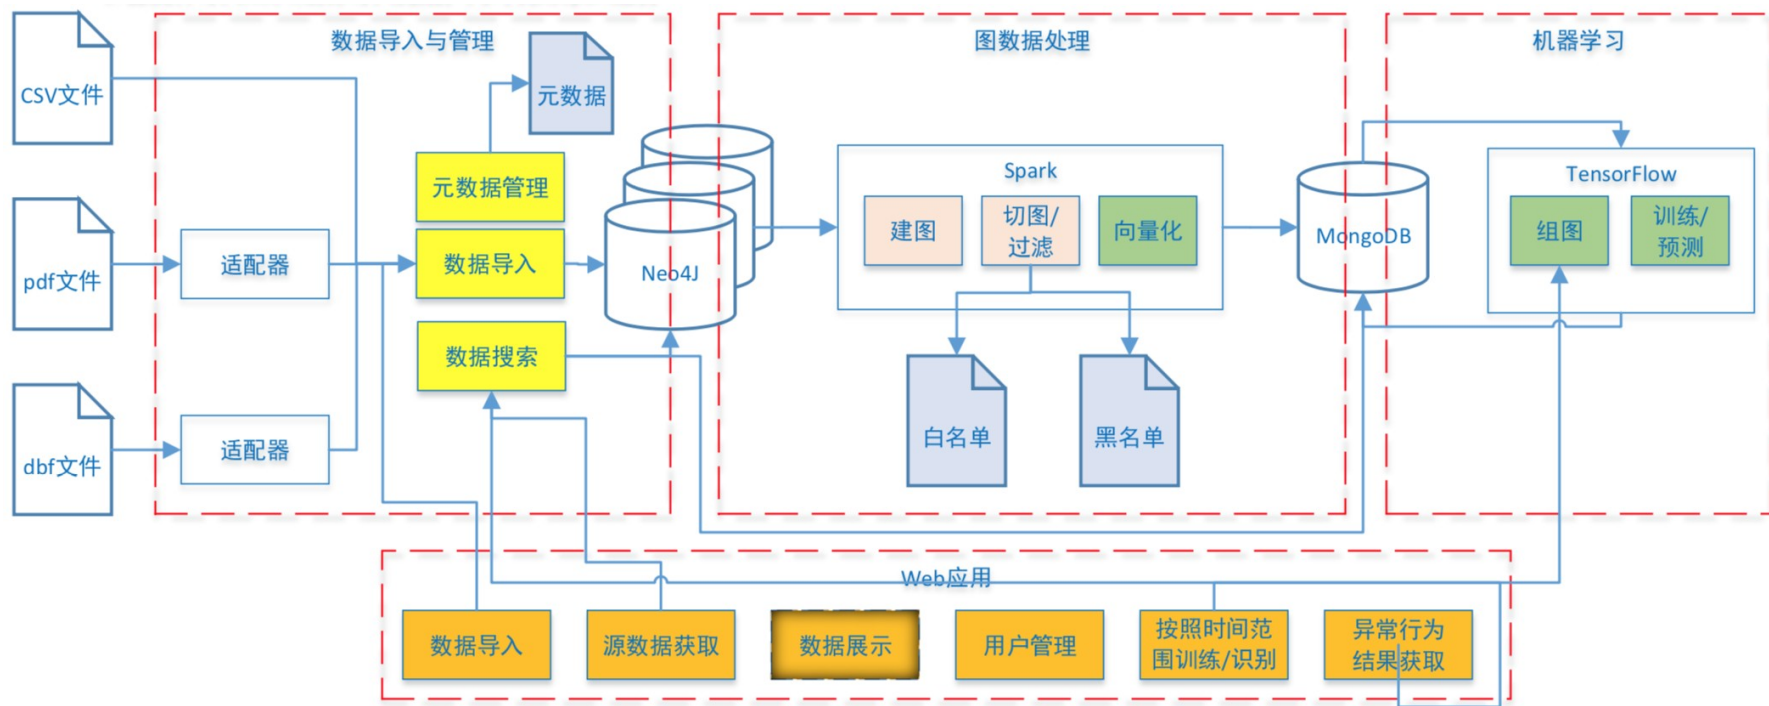



图 4-2 映射矩阵维度压缩示意图


- 根据实体关联数据构建图 -> 根据样例学习图中的模式 -> 发现符合预期模式的数据，进行预警




- Multi-Graph 多边合并



- 边属性 -> 节点属性



One-Hot Encoding



Aggregation(Sum)

	In	Out	Self-loop
①	(60,0,1)	(50,1,1)	(100,1,0)
②	(20,0,1)	(35,2,0)	(0,0,0)
③	(40,1,1)	(60,0,1)	(0,0,0)
④	(25,1,0)	(30,0,1)	(0,0,0)
⑤	(30,1,0)	(0,0,0)	(0,0,0)

- Graph -> Vector

```
data = {
  time: 2019-01,
  type: train,
  data: [
    {
      id: 1,
      label: 1,
      neighbors: [{
        id: 1,
        attr: {
          layer: 0,
          origin: [1.0, 1.0],
          summable: [1.0, 2.0]
        },
        ...
      }]
    },
    ...
  ]
}
```

时间

数据类型

数据

中心节点

ID 中心节点标签

邻居节点

邻居节点


ID

距中心节点距离

邻居节点原始属性

邻居节点可加属性
(边合并属性)

- GraphX (Pregel 迭代计算)
- 设置传播层数上限
- 固定的邻居数量 (Fixed-Size)
- 结果存储到 MongoDB 中 (Hiding Processing Latency)



- Graph Databases: NEW OPPORTUNITIES FOR CONNECTED DATA
 - http://graphdatabases.com/?ref=blog&_ga=2.114598147.542116335.1524369358-1583480045.1524187154
- Neo4j 图数据库简介和底层原理
 - <https://blog.csdn.net/bluejoe2000/article/details/72978478>
- Neo4j install
 - <https://neo4j.com/download-thanks/?edition=community&release=4.3.3&flavour=unix>
- Accessing Data with Neo4j
 - <https://spring.io/guides/gs/accessing-data-neo4j/>



Thank You!