

Architecture of Enterprise Applications 28

HDFS

Haopeng Chen

REliable, INtelligent and Scalable Systems Group (REINS)

Shanghai Jiao Tong University

Shanghai, China

<http://reins.se.sjtu.edu.cn/~chenhp>

e-mail: chen-hp@sjtu.edu.cn

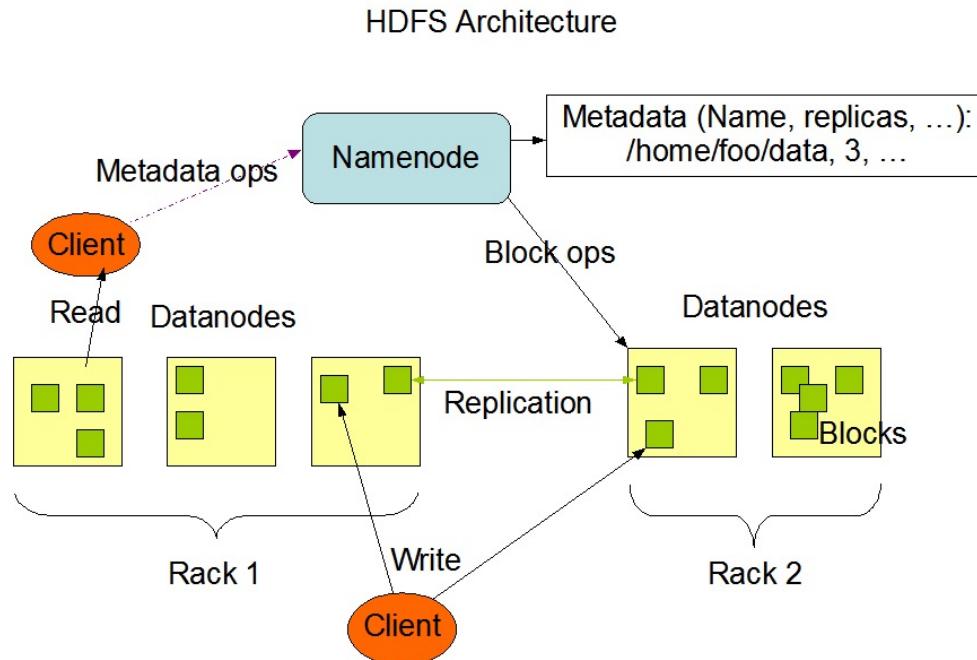
- HDFS
 - Basic Concepts
 - Namenodes and Datanodes
 - Other Issues
- Objectives
 - 能够根据大尺寸数据文件的存储需求，设计使用HDFS存储的方案，设计并实现对HDFS中存储文件的读写操作

- HDFS
 - is a file system designed for storing very large files with streaming data access patterns, running on clusters of commodity hardware.
- Very large files
 - “Very large” in this context means files that are **hundreds of megabytes, gigabytes, or terabytes** in size.
- Streaming data access
 - HDFS is built around the idea that the most efficient data processing pattern is a **write-once, read-many-times** pattern.
- Commodity hardware
 - Hadoop **doesn't** require expensive, highly reliable hardware to run on.

- These are areas where HDFS is **not** a good fit today:
 - Low-latency data access
 - Remember, HDFS is optimized for delivering a high throughput of data, and this may be at the expense of latency.
 - Lots of small files
 - Since the namenode holds file system metadata in memory, the limit to the number of files in a file system is governed by the amount of memory on the namenode.
 - Multiple writers, arbitrary file modifications
 - Files in HDFS may be written to by a single writer. Writes are always made at the end of the file. There is no support for multiple writers, or for modifications at arbitrary offsets in the file.

NameNode and DataNodes

- HDFS has a master/slave architecture.



- Blocks
 - A disk has a block size, which is the minimum amount of data that it can read or write.
 - It is a much larger unit—128 MB by default.
- HDFS has a master/slave architecture.
 - The **master** server that manages the file system namespace and regulates access to files by clients.
 - In addition, there are a number of **DataNodes**, usually **one per node in the cluster**, which manage storage attached to the nodes that they run on.

- HDFS exposes a file system namespace and allows user data to be stored in files.
 - Internally, a file is split into **one or more blocks** and these blocks are stored in a set of DataNodes.
 - The NameNode executes file system namespace operations like **opening, closing, and renaming files and directories**. It also determines the **mapping of blocks to DataNodes**.
 - The DataNodes are responsible for **serving read and write requests from the file system's clients**.
 - The DataNodes also perform **block creation, deletion, and replication** upon instruction from the NameNode.
- The existence of a single NameNode in a cluster greatly simplifies the architecture of the system.
 - The NameNode is the **arbitrator** and **repository** for **all HDFS metadata**.
 - The system is designed in such a way that user data **never flows through the NameNode**.

- The NameNode and DataNode are pieces of software designed to run on commodity machines.
 - These machines typically run a **GNU/Linux operating system (OS)**.
 - HDFS is built using the **Java** language; any machine that supports Java can run the NameNode or the DataNode software. Usage of the highly portable Java language means that HDFS can be deployed on a wide range of machines.
 - A typical deployment has **a dedicated machine that runs only the NameNode software**. Each of the other machines in the cluster runs **one instance of the DataNode software**.
 - The architecture does **not** preclude running multiple DataNodes on the same machine but in a real deployment that is rarely the case.

The File System Namespace

- HDFS supports a traditional hierarchical file organization.
 - A user or an application can create **directories** and store files inside these directories.
 - The file system namespace hierarchy is similar to most other existing file systems; one can create and remove files, move a file from one directory to another, or rename a file.
 - HDFS supports [user quotas](#) and [access permissions](#).
 - HDFS does **not** support hard links or soft links. However, the HDFS architecture does **not** preclude implementing these features.

The File System Namespace

- While HDFS follows [naming convention of the FileSystem](#), some paths and names (e.g. ./reserved and .snapshot) are reserved.
 - Features such as [transparent encryption](#) and [snapshot](#) use reserved paths.
- The **NameNode** maintains the file system namespace.
 - Any change to the file system namespace or its properties is recorded by the NameNode.
 - An application can specify **the number of replicas** of a file that should be maintained by HDFS.
 - The number of copies of a file is called the **replication factor** of that file. This information is stored by the NameNode.

- HDFS is designed to reliably store very large files **across machines** in a large cluster.
 - It stores each file as **a sequence of blocks**.
 - The blocks of a file are replicated for fault tolerance.
 - The block size and replication factor are **configurable per file**.
- All blocks in a file except the last block are the same size,
 - while users can start a new block without filling out the last block to the configured block size after the support for variable length block was added to append and hsync.
- An application can specify **the number of replicas of a file**.
 - The replication factor can be specified at file creation time and can be changed later.
 - Files in HDFS are write-once (except for appends and truncates) and have **strictly one writer** at any time.

Data Replication

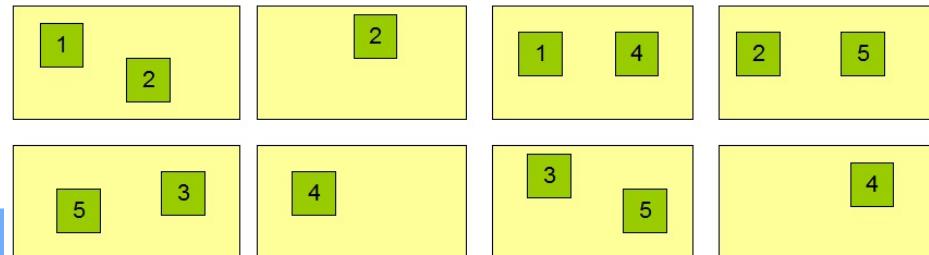
- The NameNode

- makes all decisions regarding replication of blocks.
- It periodically receives a Heartbeat and a Blockreport from each of the DataNodes in the cluster.
- Receipt of a Heartbeat implies that the DataNode is functioning properly.
- A Blockreport contains a list of all blocks on a DataNode.

Block Replication

```
Namenode (Filename, numReplicas, block-ids, ...)  
/users/sameerp/data/part-0, r:2, {1,3}, ...  
/users/sameerp/data/part-1, r:3, {2,4,5}, ...
```

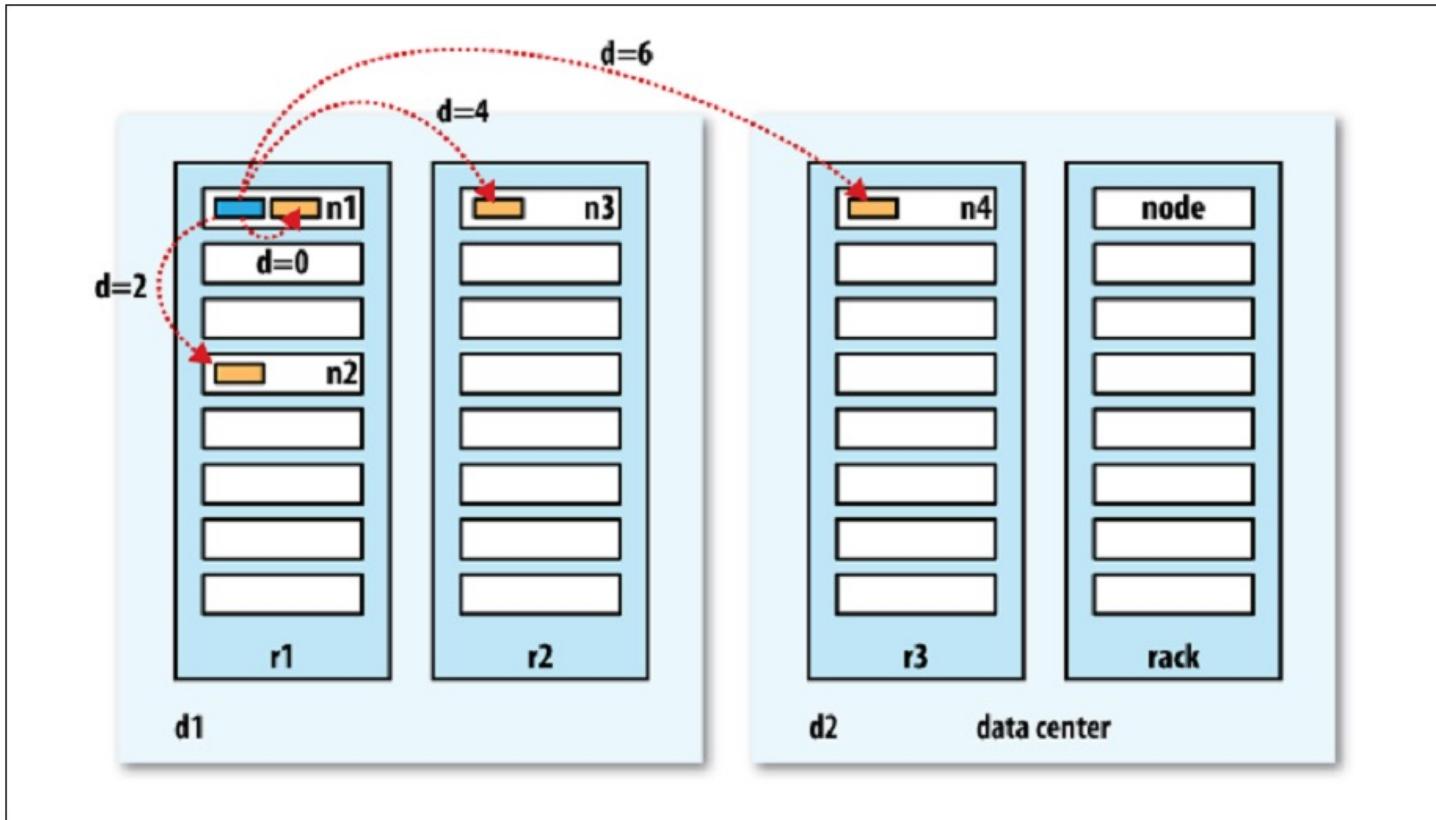
Datanodes



Replica Placement: The First Baby Steps

- The placement of replicas is critical to HDFS reliability and performance.
 - The purpose of a **rack-aware replica placement policy** is to improve data reliability, availability, and network bandwidth utilization.
- Large HDFS instances run on a cluster of computers that commonly spread across many racks.
 - Communication between two nodes in different racks has to go through switches.
 - In most cases, network bandwidth between machines in the same rack is greater than network bandwidth between machines in different racks.
- The NameNode determines the rack id each DataNode belongs to via the process outlined in [Hadoop Rack Awareness](#).
 - A simple but non-optimal policy is to place replicas on unique racks.
 - This policy **evenly distributes replicas** in the cluster which makes it easy to balance load on component failure.
 - However, this policy **increases the cost of writes** because a write needs to transfer blocks to multiple racks.

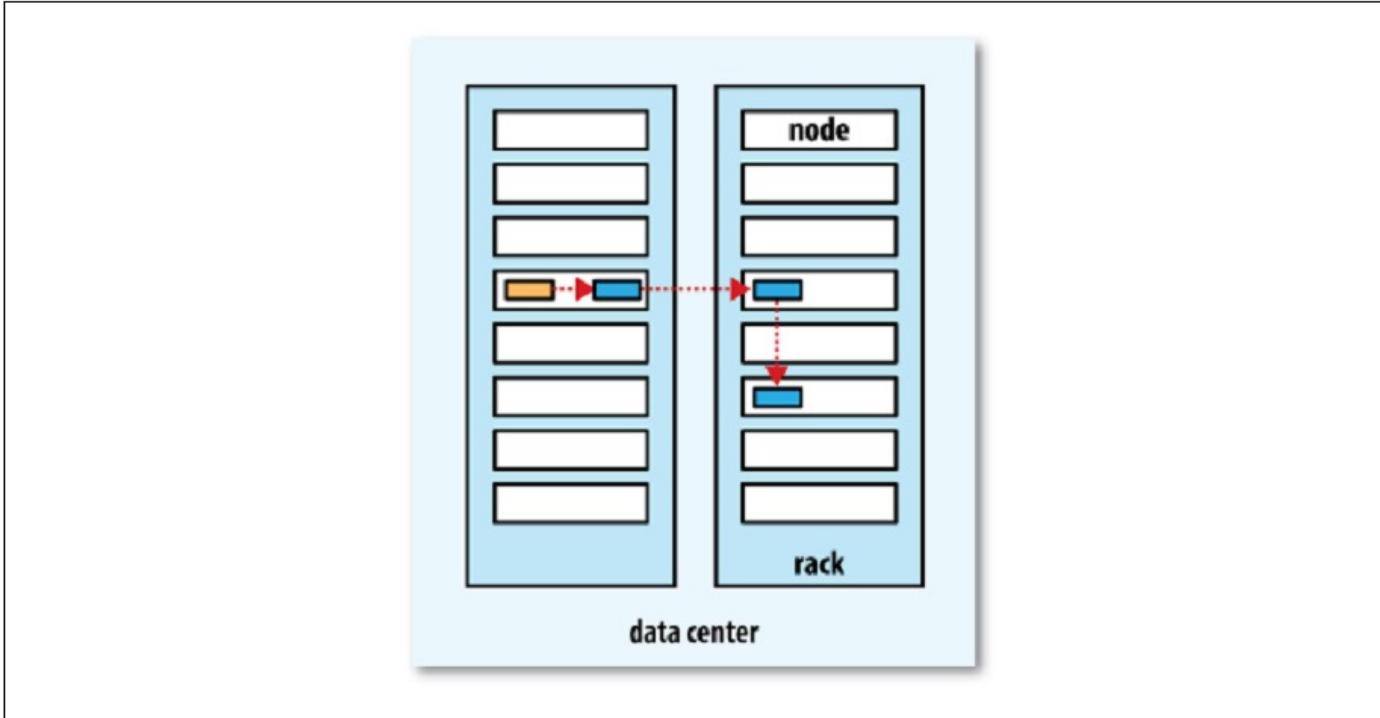
HDFS: network distance



Replica Placement: The First Baby Steps

- For the common case, when the replication factor is **three**,
 - HDFS's placement policy is to put **one replica on the local machine if the writer is on a datanode**, otherwise on a random datanode in the same rack as that of the writer,
 - **another replica** on a node in a different (remote) rack, and **the last** on a different node in the same remote rack.
- If the replication factor is **greater than 3**,
 - the placement of the **4th and following replicas** are determined **randomly** while keeping the number of replicas per rack below **the upper limit** (which is basically $(\text{replicas} - 1) / \text{racks} + 2$).

HDFS: Replicas



- To minimize global bandwidth consumption and read latency,
 - HDFS tries to satisfy a **read request** from a replica that is **closest to the reader**.
 - If there exists a replica on the **same rack** as the reader node, then that replica is preferred to satisfy the read request.
 - If HDFS cluster spans multiple data centers, then a replica that is resident in the **local data center is preferred over any remote replica**.

- On startup, the **NameNode** enters a special state called **Safemode**.
 - Replication of data blocks does **not** occur when the NameNode is in the Safemode state.
 - The NameNode receives **Heartbeat** and **Blockreport** messages from the DataNodes.
 - A Blockreport contains the list of data blocks that a DataNode is hosting.
 - Each block has a specified **minimum number of replicas**.
 - A block is considered safely replicated when the minimum number of replicas of that data block has checked in with the NameNode.
 - After **a configurable percentage of safely replicated data blocks** checks in with the NameNode (plus an additional 30 seconds), the NameNode exits the Safemode state.
 - It then determines the list of data blocks (if any) that still have fewer than the specified number of replicas.
 - The NameNode then replicates these blocks to other DataNodes.

The Persistence of File System Metadata

- The HDFS namespace is stored by the NameNode.
 - The NameNode uses a transaction log called the **EditLog** to persistently record every change that occurs to file system metadata.
 - For example, **creating a new file in HDFS** causes the NameNode to **insert a record** into the EditLog indicating this.
 - Similarly, **changing the replication factor of a file** causes a **new record** to be inserted into the EditLog.
 - The NameNode uses a file in its **local host OS file system** to store the EditLog.
 - The entire file system namespace, including the mapping of blocks to files and file system properties, is stored in a file called the **FsImage**.
 - The FsImage is stored as a file in the NameNode's **local file system** too.

The Persistence of File System Metadata

- The NameNode keeps an image of the entire file system namespace and file Blockmap in memory.
 - When the NameNode starts up, or a checkpoint is triggered by a configurable threshold, it reads the FsImage and EditLog from disk, applies all the transactions from the EditLog to the in-memory representation of the FsImage, and flushes out this new version into a new FsImage on disk.
 - It can then truncate the old EditLog because its transactions have been applied to the persistent FsImage.
 - This process is called a checkpoint.
 - The purpose of a checkpoint is to make sure that HDFS has a consistent view of the file system metadata by taking a snapshot of the file system metadata and saving it to FsImage.
 - A checkpoint can be triggered at a given time interval (dfs.namenode.checkpoint.period) expressed in seconds, or after a given number of filesystem transactions have accumulated (dfs.namenode.checkpoint.txns).
 - If both of these properties are set, the first threshold to be reached triggers a checkpoint.

The Persistence of File System Metadata

- The DataNode stores HDFS data in files in its local file system.
 - The DataNode has **no** knowledge about HDFS files.
 - It stores each block of HDFS data in **a separate file in its local file system**.
 - The DataNode does **not** create all files in the same directory.
 - Instead, it uses a **heuristic** to determine the **optimal number of files per directory** and **creates subdirectories appropriately**.
 - It is **not optimal** to create all local files in **the same directory** because the local file system might not be able to efficiently support a huge number of files in a single directory.
 - When a DataNode starts up, it scans through its local file system, generates a list of all HDFS data blocks that correspond to each of these local files, and sends this report to the NameNode.
 - The report is called the ***Blockreport***.

The Communication Protocols

- All HDFS communication protocols are layered on top of the **TCP/IP** protocol.
 - A **client** establishes a connection to a **configurable TCP port** on the NameNode machine.
 - It talks the **Client Protocol** with the NameNode.
 - The DataNodes talk to the NameNode using the **DataNode Protocol**.
 - A Remote Procedure Call (**RPC**) abstraction wraps both the Client Protocol and the DataNode Protocol.
 - By design, the NameNode **never** initiates any RPCs. Instead, it **only responds** to RPC requests issued by DataNodes or clients.

- The primary objective of HDFS is to store data reliably even in the presence of failures.
 - The three common types of failures are **NameNode** failures, **DataNode** failures and **network partitions**.
- **Data Disk Failure, Heartbeats and Re-Replication**
 - Each DataNode sends a **Heartbeat** message to the NameNode **periodically**.
 - A **network partition** can cause a subset of DataNodes to lose connectivity with the NameNode.
 - The NameNode detects this condition by the absence of a Heartbeat message.
 - The NameNode marks DataNodes **without recent Heartbeats** as **dead** and does **not** forward any new IO requests to them. Any data that was registered to a **dead DataNode** is **not** available to HDFS any more.

- The primary objective of HDFS is to store data reliably even in the presence of failures.
 - The three common types of failures are **NameNode** failures, **DataNode** failures and **network** partitions.
- **Data Disk Failure, Heartbeats and Re-Replication**
 - DataNode death may cause the **replication factor** of some blocks to fall **below** their specified value.
 - The NameNode **constantly tracks** which blocks need to be replicated and initiates replication whenever necessary. The necessity for re-replication may arise due to many reasons:
 - a DataNode may become unavailable,
 - a replica may become corrupted,
 - a hard disk on a DataNode may fail,
 - or the replication factor of a file may be increased.
 - The time-out to mark DataNodes dead is conservatively long (over 10 minutes by default) in order to avoid replication storm caused by state flapping of DataNodes.

- **Cluster Rebalancing**
- The HDFS architecture is **compatible** with **data rebalancing schemes**.
 - A scheme might automatically move data from one DataNode to another if the free space on a DataNode falls below a certain threshold.
 - In the event of a **sudden high demand** for a particular file, a scheme might **dynamically create additional replicas** and rebalance other data in the cluster.
 - These types of data rebalancing schemes are not yet implemented.

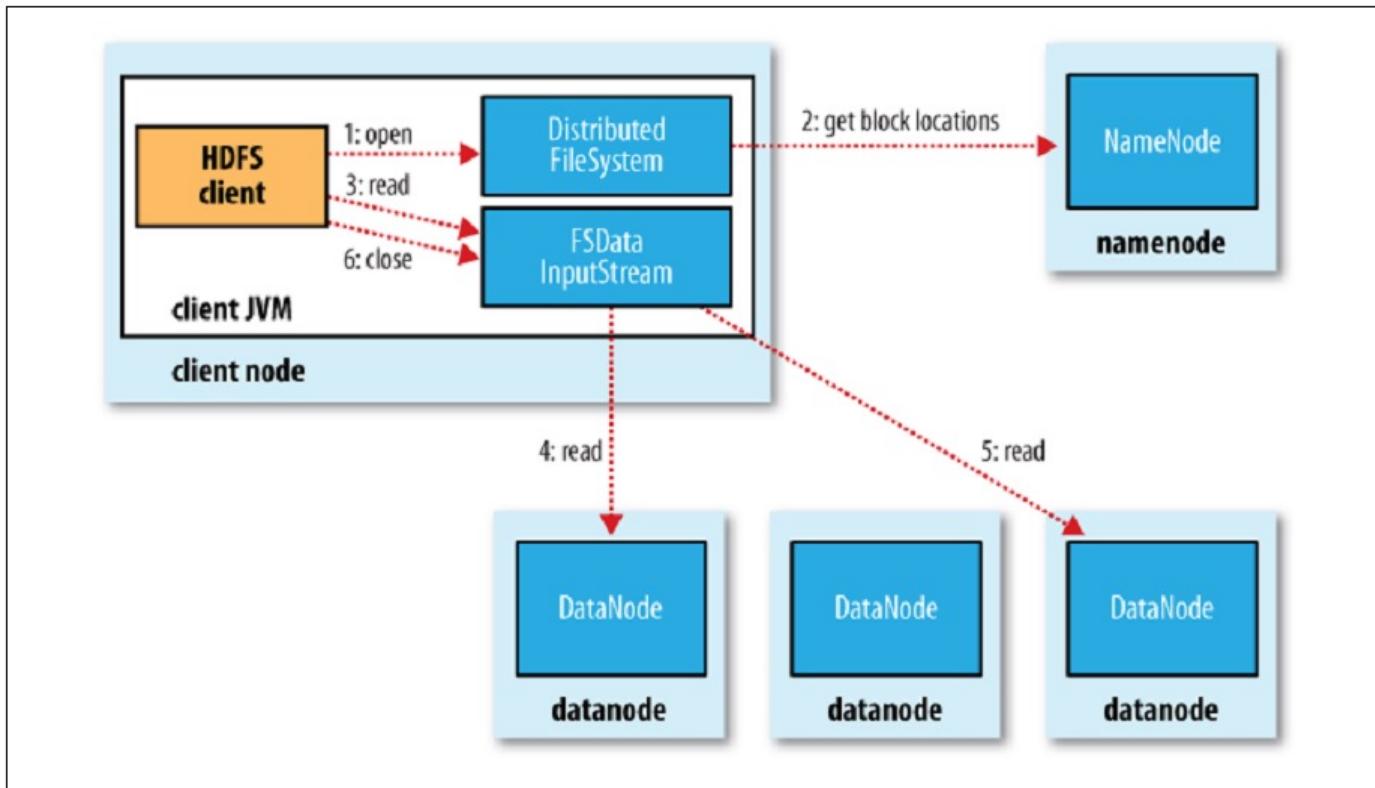
- **Data Integrity**
- It is possible that a block of data fetched from a DataNode arrives **corrupted**.
 - This corruption can occur because of faults in **a storage device, network faults, or buggy software**.
 - The HDFS client software implements **checksum checking** on the contents of HDFS files.
 - When a client creates an HDFS file, it computes **a checksum of each block of the file** and **stores these checksums in a separate hidden file** in the same HDFS namespace.
 - When a client retrieves file contents it verifies that the data it received from each DataNode matches the checksum stored in the associated checksum file.
 - If **not**, then the client can opt to retrieve that block from **another DataNode** that has a replica of that block.

- **Metadata Disk Failure**
- The **FsImage** and the **EditLog** are central data structures of HDFS.
 - A corruption of these files can cause the HDFS instance to be non-functional.
 - For this reason, the NameNode can be configured to support **maintaining multiple copies of the FsImage and EditLog**.
 - Any update to either the FsImage or EditLog causes each of the FsImages and EditLogs to get updated **synchronously**.
 - When a NameNode restarts, it selects the latest consistent FsImage and EditLog to use.
- **Snapshots**
 - Snapshots support storing a copy of data at a particular instant of time. One usage of the snapshot feature may be to roll back a corrupted HDFS instance to a previously known good point in time.

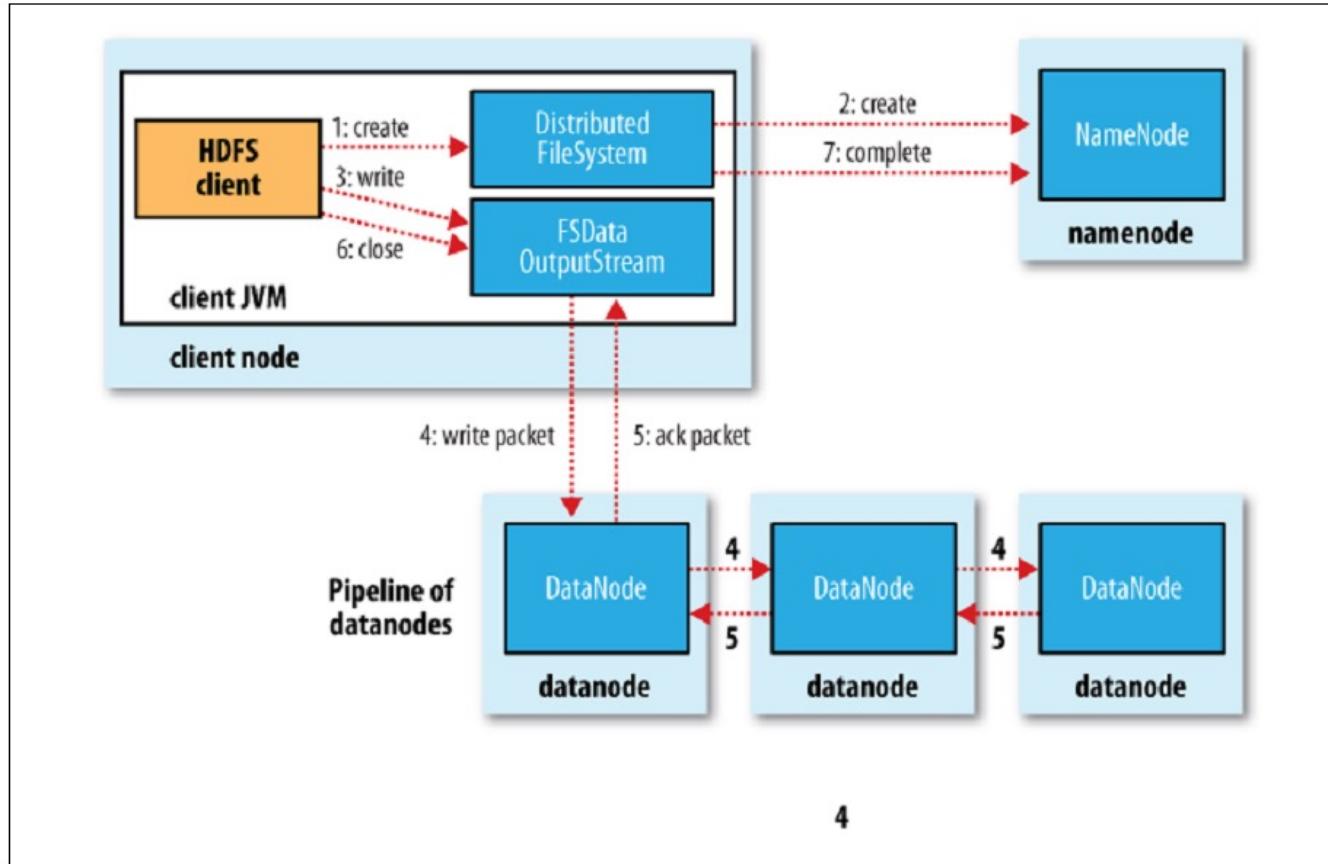
- **Data Blocks**
- HDFS is designed to support very large files.
 - Applications that are compatible with HDFS are those that deal with large data sets.
 - These applications **write their data only once** but they **read it one or more times** and require these reads to be satisfied at streaming speeds.
 - HDFS supports **write-once-read-many** semantics on files.
 - A typical block size used by HDFS is **128 MB**. Thus, an HDFS file is chopped up into 128 MB chunks, and if possible, **each chunk will reside on a different DataNode**.

- **Replication Pipelining**
- When a client is writing data to an HDFS file with a **replication factor of three**,
 - the NameNode retrieves a list of DataNodes using a **replication target choosing algorithm**.
 - This list contains the DataNodes that will host a replica of that block.
 - The client then writes to the **first** DataNode.
 - The **first** DataNode starts receiving the data in portions, writes each portion to its local repository and transfers that portion to the **second** DataNode in the list. The **second** DataNode, in turn starts receiving each portion of the data block, writes that portion to its repository and then flushes that portion to the **third** DataNode. Finally, the **third** DataNode writes the data to its local repository.
 - Thus, a DataNode can be receiving data from the **previous one** in the pipeline and at the same time forwarding data to the **next one** in the pipeline.
 - Thus, the data is **pipelined from one DataNode to the next**.

HDFS: Reading data



HDFS: Writing data



- HDFS can be accessed from applications in many different ways.
 - Natively, HDFS provides a [FileSystem Java API](#) for applications to use.
 - A [C language wrapper for this Java API](#) and [REST API](#) is also available.
 - In addition, an HTTP browser and can also be used to browse the files of an HDFS instance.
 - By using [NFS gateway](#), HDFS can be mounted as part of the client's local file system.
- **Browser Interface**
 - A typical HDFS install configures a **web server** to expose the HDFS namespace through a configurable TCP port.
 - This allows a user to navigate the HDFS namespace and view the contents of its files using a web browser.

Overview 'localhost:9000' (active)

Started:	Fri Dec 17 15:19:01 +0800 2021
Version:	3.2.2, r7a3bc90b05f257c8ace2f76d74264906f0f7a932
Compiled:	Sun Jan 03 17:26:00 +0800 2021 by hexiaoqiao from branch-3.2.2
Cluster ID:	CID-c85b6ff1-3dac-423a-91d8-32ab8eb79c79
Block Pool ID:	BP-391551921-127.0.0.1-1639725536253

Summary

Security is off.

Safemode is off.

1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).

Heap Memory used 93.89 MB of 161 MB Heap Memory. Max Heap Memory is 4 GB.

Non Heap Memory used 46.19 MB of 50 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

Configured Capacity:	0 B
Configured Remote Capacity:	0 B
DFS Used:	0 B (100%)
Non DFS Used:	0 B

- **FS Shell**
- HDFS allows user data to be organized in the form of files and directories.
 - It provides a **commandline** interface called **FS shell** that lets a user interact with the data in HDFS.
 - The syntax of this command set is similar to other shells (e.g. bash, csh) that users are already familiar with. Here are some sample action/command pairs:

Action	Command
Create a directory named /foodir	bin/hadoop dfs -mkdir /foodir
Remove a directory named /foodir	bin/hadoop fs -rm -R /foodir
View the contents of a file named /foodir/myfile.txt	bin/hadoop dfs -cat /foodir/myfile.txt

- **DFSAdmin**
- The **DFSAdmin** command set is used for administering an HDFS cluster.
 - These are commands that are used **only by an HDFS administrator**.
 - Here are some sample action/command pairs:

Action	Command
Put the cluster in Safemode	bin/hdfs dfsadmin -safemode enter
Generate a list of DataNodes	bin/hdfs dfsadmin -report
Recommission or decommission DataNode(s)	bin/hdfs dfsadmin -refreshNodes

- **File Deletes and Undeletes**

- If trash configuration is enabled, files removed by [FS Shell](#) is **not** immediately removed from HDFS.
- Instead, HDFS moves it to a **trash directory** (each user has its own trash directory under /user/<username>/.Trash).
- The file can be restored quickly as long as it remains in trash.
- Most recent deleted files are moved to the current trash directory (/user/<username>/.Trash/Current), and in a **configurable interval**, HDFS creates checkpoints (under /user/<username>/.Trash/<date>) for files in current trash directory and deletes old checkpoints when they are expired.
- After the **expiry** of its life in trash, the NameNode deletes the file from the HDFS namespace.
- The deletion of a file causes **the blocks associated with the file to be freed**.
- Note that there could be an appreciable time delay between the time a file is deleted by a user and the time of the corresponding increase in free space in HDFS.

- Following is an example which will show how the files are deleted from HDFS by FS Shell.

- We created 2 files (test1 & test2) under the directory delete

```
$ hadoop fs -mkdir -p delete/test1
```

```
$ hadoop fs -mkdir -p delete/test2
```

```
$ hadoop fs -ls delete/
```

```
Found 2 items
```

```
drwxr-xr-x - hadoop hadoop 0 2015-05-08 12:39 delete/test1
```

```
drwxr-xr-x - hadoop hadoop 0 2015-05-08 12:40 delete/test2
```

- We are going to remove the file **test1**. The comment below shows that the file has been moved to Trash directory.

```
$ hadoop fs -rm -r delete/test1
```

```
Moved: hdfs://localhost:8020/user/hadoop/delete/test1 to trash at:  
hdfs://localhost:8020/user/hadoop/.Trash/Current
```

- now we are going to remove the file with **skipTrash** option, which will **not** send the file to Trash. It will be completely removed from HDFS.

```
$ hadoop fs -rm -r -skipTrash delete/test2
```

```
Deleted delete/test2
```

- We can see now that the Trash directory contains **only file test1**.

```
$ hadoop fs -ls .Trash/Current/user/hadoop/delete/  
Found 1 items\  
drwxr-xr-x - hadoop hadoop 0 2015-05-08 12:39  
.Trash/Current/user/hadoop/delete/test1
```

- So file test1 goes to Trash and file test2 is deleted permanently.

- **Decrease Replication Factor**
- When the replication factor of a file is **reduced**, the NameNode selects excess replicas that can be deleted.
 - The next Heartbeat transfers this information to the DataNode.
 - The DataNode then removes the corresponding blocks and the corresponding free space appears in the cluster.
 - Once again, there might be a **time delay** between the completion of the setReplication API call and the appearance of free space in the cluster.

- Format the filesystem:
`$ bin/hdfs namenode -format`
- Start NameNode daemon and DataNode daemon:
`$ sbin/start-dfs.sh`
- Browse the web interface for the NameNode; by default it is available at:
 - NameNode - <http://localhost:9870/>

- TestHDFS.java

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;

public class TestHDFS {
    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();

        conf.set("dfs.defaultFS", "hdfs://hadoop:9000");
        FileSystem fs = FileSystem.get(conf);
        fs.create(new Path("./test/helloByJava"));
        fs.close();
    }
}
```

- HDFSApp.java

```
public class HDFSApp {  
    public static final String HDFS_PATH = "hdfs://hadoop:9000";  
    //文件系统  
    FileSystem fileSystem = null;  
    //配置类  
    Configuration configuration = null;  
  
    //Before适用于类加载之前  
    @Test  
    /**  
     * 创建HDFS系统  
     */  
    public void mkdir() throws Exception {  
        fileSystem.mkdirs(new Path("./output"));  
    }  
}
```

- HDFSApp.java

```
@Test
public void create() throws Exception {
    FSDataOutputStream output = fileSystem.create(new Path("./output/a.txt"));
    output.write("hello hadoop".getBytes());
    output.flush();
    output.close();
}

@Test
/**
 * 查看HDFS文件上的内容
 */
public void cat() throws Exception {
    FSDataInputStream in = fileSystem.open(new Path("./output/b.txt"));
    IOutils.copyBytes(in, System.out, 1024);
    in.close();
}
```

- HDFSApp.java

```
@Before
public void setUp() throws Exception {
    System.out.printf("HDFSapp.setup");
    configuration = new Configuration();
    configuration.set("dfs.defaultFS", "hdfs://hadoop:9000");
    fileSystem = FileSystem.get(configuration);
}

//关闭资源用的这个
@After
public void tearDown() throws Exception {
    //释放资源
    configuration = null;
    fileSystem = null;
    System.out.printf("HDFSAPP.tearDown");
}
```

- HDFSApp.java

```
/**  
 * 重命名文件  
 */  
  
@Test  
public void rename() throws Exception {  
    Path oldPath = new Path("./output/a.txt");  
    Path newPath = new Path("./output/b.txt");  
    fileSystem.rename(oldPath, newPath);  
}
```

- HDFSApp.java

```
/**  
 * 上传一个文件  
 *  
 * @throws Exception  
 */  
@Test  
public void copyFromLocalFile() throws Exception {  
    Path localPath = new Path("./test/helloByJava");  
    Path hdfsPath = new Path("./output/helloByJava");  
    fileSystem.copyFromLocalFile(localPath, hdfsPath);  
}
```

- HDFSApp.java

```
/*
 * 上传一个大文件
 * @throws Exception
 */
@Test
public void copyFromLocalBigFile() throws Exception {
    InputStream in = new BufferedInputStream(
        new FileInputStream(
            new File("./resources/hadoop-hdfs-3.2.1.jar")));
    FSDataOutputStream output = fileSystem.create(
        new Path("./output/hadoop-hdfs-3.2.1.jar"),
        new Progressable() {
            public void progress() {
                System.out.print(".");
                //带进度提醒信息
            }
        });
    IOUtils.copyBytes(in, output, 4096);
}
```

- HDFSApp.java

```
/*
 * 下载HDFS文件
 */
@Test
public void copyToLocalFile() throws Exception{
    Path localPath = new Path("./test/b.txt");
    Path hdfsPath = new Path("./output/b.txt");
    fileSystem.copyToLocalFile(false,hdfsPath,localPath,true);
}
```

- HDFSApp.java

```
@Test
public void listFiles() throws Exception {
    FileStatus[] fileStatuses = fileSystem.listStatus(new Path("./output"));

    for (FileStatus fileStatus : fileStatuses) {
        String isDir = fileStatus.isDirectory() ? "文件夹" : "文件";
        //副本
        short replication = fileStatus.getReplication();
        //大小
        long len = fileStatus.getLength();
        //路径
        String path = fileStatus.getPath().toString();

        System.out.println(isDir + "\t" + replication + "\t" + len + "\t" + path);
    }
}
```

- HDFSApp.java

```
@Test
    public void delete() throws Exception{
        fileSystem.delete(new Path("./test/"),true);
    }

}
```

Samples

- pom.xml

```
<dependency>
    <groupId>org.apache.hadoop</groupId>
    <artifactId>hadoop-common</artifactId>
    <version>3.2.1</version>
</dependency>
<dependency>
    <groupId>org.apache.hadoop</groupId>
    <artifactId>hadoop-hdfs</artifactId>
    <version>3.2.1</version>
</dependency>
<dependency>
    <groupId>org.apache.hadoop</groupId>
    <artifactId>hadoop-hdfs-client</artifactId>
    <version>3.2.1</version>
    <scope>provided</scope>
</dependency>
```

```
<dependency>
    <groupId>junit</groupId>
    <artifactId>junit</artifactId>
    <version>4.13</version>
    <scope>test</scope>
</dependency>
<dependency>
    <groupId>junit</groupId>
    <artifactId>junit</artifactId>
    <version>4.13</version>
    <scope>compile</scope>
</dependency>
```

Samples

The screenshot shows an IDE interface with the following details:

- Project Structure:** The project is named "SE343_13_HDFSSample". It contains a "output" directory with files like ".b.txt.crc", ".hadoop-hdfs-3.2.1.jar.crc", ".helloByJava.crc", "b.txt", and "hadoop-hdfs-3.2.1.jar". There are also "resources" and "src" directories.
- pom.xml Content:**

```
<!-- https://mvnrepository.com/artifact/org.apache.hadoop/hadoop-common -->
<dependency>
    <groupId>org.apache.hadoop</groupId>
    <artifactId>hadoop-common</artifactId>
    <version>3.2.1</version>
</dependency>
<!-- https://mvnrepository.com/artifact/org.apache.hadoop/hadoop-hdfs -->
<dependency>
    <groupId>org.apache.hadoop</groupId>
    <artifactId>hadoop-hdfs</artifactId>
    <version>3.2.1</version>
```
- Terminal Log Output:**

```
2022-03-24 14:02:55,427 INFO common.Storage: Storage directory /tmp/hadoop-hadoop/haopengchen/namenode has been successfully formatted.
2022-03-24 14:02:55,489 INFO namenode.FSImageFormatProtobuf: Saving image file /tmp/hadoop-haopengchen/dfs/namenode/current/fsimage.ckpt_00000000000000000000 using no compression
2022-03-24 14:02:55,578 INFO namenode.FSImageFormatProtobuf: Image file /tmp/hadoop-haopengchen/dfs/namenode/current/fsimage.ckpt_00000000000000000000 of size 406 by test
2022-03-24 14:02:55,697 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid > 0
2022-03-24 14:02:55,697 INFO namenode.FSImageSaver clean checkpoint: tx id@ when meet shutdown.
2022-03-24 14:02:55,697 INFO namenode.Namenode: SHUTDOWN_MSG:
*****Shutdown Message*****  
SHUTDOWN_MSG: Shutting down Namenode at HAOPENNodeMac.local/127.0.0.1
*****Shutdown Message*****
```
- Run Tab:**
 - Test results for "HDFSApp": 8 tests passed in 461 ms.
 - Test cases include: create, listFiles, copyFromLocalBigFile, rename, cat, mkdir, copyToLocalFile, and copyFromLocalFile.
 - Log output for the tests includes: "No appenders could be found for logger (org.apache.hadoop.util.Shell)".

- HDFS Architecture
 - <https://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-hdfs/HdfsDesign.html>
- Hadoop: The Definitive Guide
 - By Tom White
 - O'Reilly Publishing
- Hadoop入门IntelliJ编写Maven创建HDFS文件
 - https://blog.csdn.net/qq_39622065/article/details/86600322
- IDEA 创建HDFS项目 JAVA api
 - <https://www.cnblogs.com/chenligeng/p/9348120.html>



Thank You!