

# Architecture of Enterprise Applications 30

## Hive

Haopeng Chen

*REliable, INtelligent and Scalable Systems Group (REINS)*

Shanghai Jiao Tong University

Shanghai, China

<http://reins.se.sjtu.edu.cn/~chenhp>

e-mail: chen-hp@sjtu.edu.cn

- Hive
  - Basic Concepts
  - Hive DDL and DML
  - Other Issues
- Objectives
  - 能够根据数据仓库的基础要求，设计基于Hive的数据分析方案

- Hive, a framework for **data warehousing on top of Hadoop**.
  - Hive grew from a need to manage and learn from the **huge volumes of data** that Facebook was producing every day from its burgeoning social network.
  - After trying a few different systems, the team chose Hadoop for storage and processing, since it was cost-effective and met their scalability needs.
- Hive was created to make it possible for analysts with strong **SQL** skills (but meager Java programming skills) to run queries on the huge volumes of data that Facebook stored in **HDFS**.
  - Today, Hive is a successful **Apache** project used by many organizations as a general-purpose, scalable data processing platform.
- <https://hive.apache.org/>



- In normal use, Hive runs on your workstation and converts your **SQL** query into a series of **MapReduce** jobs for execution on a Hadoop cluster.
  - Hive organizes data into **tables**, which provide a means for attaching structure to data stored in HDFS.
  - Metadata—such as table schemas—is stored in a database called the **metastore**.

- Hive uses Hadoop, so:
  - you must have Hadoop in your path OR
  - `$ export HADOOP_HOME=<hadoop-install-dir>`
- Start NameNode daemon and DataNode daemon:
  - `$ sbin/start-dfs.sh`
- In addition,
  - you must use below HDFS commands to create `/tmp` and `/user/hive/warehouse`
  - (aka `hive.metastore.warehouse.dir`) and set them `chmod g+w` before you can create a table in Hive.
  - `$ HADOOP_HOME/bin/hadoop fs -mkdir /tmp`
  - `$ HADOOP_HOME/bin/hadoop fs -mkdir /user/hive/warehouse`
  - `$ HADOOP_HOME/bin/hadoop fs -chmod g+w /tmp`
  - `$ HADOOP_HOME/bin/hadoop fs -chmod g+w /user/hive/warehouse`

- To use the Hive [command line interface](#) (CLI) from the shell:
  - `$ HIVE_HOME/bin/hive`
- **Running HiveServer2 and Beeline**
  - Starting from Hive 2.1, we need to run the **schematool** command below as an initialization step. For example, we can use "**derby**" as db type.
  - `$ HIVE_HOME/bin/schematool -dbType <db type> -initSchema`
- HiveServer2 (introduced in Hive 0.11) has its own CLI called [Beeline](#).
  - HiveCLI is now deprecated in favor of Beeline, as it lacks the multi-user, security, and other capabilities of HiveServer2. To run HiveServer2 and Beeline from shell:
    - `$ HIVE_HOME/bin/hiveserver2`
    - `$ HIVE_HOME/bin/beeline -u jdbc:hive2://$HS2_HOST:$HS2_PORT`

# Configuration Management Overview

- Hive by default gets its configuration from <install-dir>/conf/hive-default.xml
- Hive configuration is an overlay on top of Hadoop
  - it inherits the Hadoop configuration variables by default.
- Hive configuration can be manipulated by:
  - Editing **hive-site.xml** and defining any desired variables (including Hadoop variables) in it
  - Using the set command
  - Invoking **Hive** (deprecated), **Beeline** or **HiveServer2** using the syntax:
    - \$ bin/hive --hiveconf x1=y1 --hiveconf x2=y2 //this sets the variables x1 and x2 to y1 and y2 respectively
    - \$ bin/hiveserver2 --hiveconf x1=y1 --hiveconf x2=y2 //this sets server-side variables x1 and x2 to y1 and y2 respectively
    - \$ bin/beeline --hiveconf x1=y1 --hiveconf x2=y2 //this sets client-side variables x1 and x2 to y1 and y2 respectively.
  - Setting the **HIVE\_OPTS** environment variable to "--hiveconf x1=y1 --hiveconf x2=y2" which does the same as above.

# Hive, Map-Reduce and Local-Mode

- Hive compiler generates **map-reduce jobs for most queries**.
  - These jobs are then submitted to the Map-Reduce cluster indicated by the variable:
  - **mapred.job.tracker**
  - While this usually points to a map-reduce **cluster** with multiple nodes, Hadoop also offers a nifty option to run map-reduce jobs **locally** on the user's workstation.
  - This can be very useful to run queries over **small data sets** – in such cases local mode execution is usually significantly faster than submitting jobs to a large cluster.
  - Data is accessed transparently from **HDFS**. Conversely, local mode only runs with **one reducer** and can be very slow processing larger data sets.

# Hive, Map-Reduce and Local-Mode

- Starting with release 0.7, Hive fully supports local mode execution.
  - To enable this, the user can enable the following option:
  - `hive> SET mapreduce.framework.name=local;`
- In addition,
  - `mapred.local.dir` should point to a path that's valid on the local machine (for example `/tmp/<username>/mapred/local`).
  - (Otherwise, the user will get an exception allocating local disk space.)

- Starting with release 0.7, Hive also supports a mode to run map-reduce jobs in local-mode automatically.
  - The relevant options are `hive.exec.mode.local.auto`, `hive.exec.mode.local.auto.inputbytes.max`, and `hive.exec.mode.local.auto.tasks.max`:
  - `hive> SET hive.exec.mode.local.auto=false;`
- Note that this feature is *disabled* by default.
  - If enabled, Hive analyzes the size of each map-reduce job in a query and may run it locally if the following thresholds are satisfied:
    - The total input size of the job is lower than: `hive.exec.mode.local.auto.inputbytes.max` (128MB by default)
    - The total number of map-tasks is less than: `hive.exec.mode.local.auto.tasks.max` (4 by default)
    - The total number of reduce tasks required is 1 or 0.
  - So for queries over small data sets, or for queries with multiple map-reduce jobs where the input to subsequent jobs is substantially smaller (because of reduction/filtering in the prior job), jobs may be run **locally**.

- The Hive DDL operations are documented in [Hive Data Definition Language](#).
- **Creating Hive Tables**
  - `hive> CREATE TABLE pokes (foo INT, bar STRING);`
  - creates a table called pokes with two columns, the first being an integer and the other a string.
  - `hive> CREATE TABLE invites (foo INT, bar STRING) PARTITIONED BY (ds STRING);`
  - creates a table called invites with two columns and a **partition column** called ds.
  - The partition column is a **virtual** column. It is **not** part of the data itself but is derived from the partition that a particular dataset is loaded into.
  - By default, tables are assumed to be of text input format and the **delimiters** are assumed to be **`^A(ctrl-a)`**.

## Browse Directory

/user/hive/warehouse

Go!



Show 25 entries

Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	drwxr-xr-x	chenhaopeng	supergroup	0 B	Dec 24 22:09	0	0 B	invites	
<input type="checkbox"/>	drwxr-xr-x	chenhaopeng	supergroup	0 B	Dec 24 22:06	0	0 B	pokes	
<input type="checkbox"/>	drwxr-xr-x	chenhaopeng	supergroup	0 B	Dec 24 22:15	0	0 B	u_data	
<input type="checkbox"/>	drwxr-xr-x	chenhaopeng	supergroup	0 B	Dec 24 22:27	0	0 B	u_data_new	

Showing 1 to 4 of 4 entries

Previous **1** Next

- **Browsing through Tables**

- `hive> SHOW TABLES;`
- lists all the tables.
- `hive> SHOW TABLES '.*s';`
- lists all the table that end with 's'. The pattern matching follows Java regular expressions.

```
– hive> DESCRIBE invites;
– OK
– foo          int
– bar          string
– ds           string
–
– # Partition Information
– # col_name      data_type      comment
– ds             string
– Time taken: 0.181 seconds, Fetched: 7 row(s)
```

- **Altering and Dropping Tables**

- Table names can be changed and columns can be added or replaced:
  - `hive> ALTER TABLE events RENAME TO 3koobecaf;`
  - `hive> ALTER TABLE pokes ADD COLUMNS (new_col INT);`
  - `hive> ALTER TABLE invites ADD COLUMNS (new_col2 INT COMMENT 'a comment');`
  - `hive> ALTER TABLE invites REPLACE COLUMNS (foo INT, bar STRING, baz INT COMMENT 'baz replaces new_col2');`
- Note that REPLACE COLUMNS replaces all existing columns and **only** changes the table's schema, **not** the data.
- The table must use a native SerDe. REPLACE COLUMNS can also be used to drop columns from the table's schema:
  - `hive> ALTER TABLE invites REPLACE COLUMNS (foo INT COMMENT 'only keep the first column');`
- Dropping tables:
  - `hive> DROP TABLE pokes;`

- **Metadata Store**

- Metadata is in an **embedded Derby database** whose disk storage location is determined by the Hive configuration variable named **javax.jdo.option.ConnectionURL**.
  - By default this location is **./metastore\_db** (see conf/hive-default.xml).
- Right now, in the default configuration, this metadata can only be seen by one user at a time.
- Metastore can be stored in any database that is supported by **JPOX**.
- The location and the type of the RDBMS can be controlled by the two variables **javax.jdo.option.ConnectionURL** and **javax.jdo.option.ConnectionDriverName**.
- In the future, the metastore itself can be a standalone server.
- If you want to run the metastore as a network server so it can be accessed from multiple nodes, see [Hive Using Derby in Server Mode](#).

- The Hive DML operations are documented in [Hive Data Manipulation Language](#).
  - Loading data from flat files into Hive:
  - `hive> LOAD DATA LOCAL INPATH './examples/files/kv1.txt' OVERWRITE INTO TABLE pokes;`
  - Loads a file that contains two columns separated by **ctrl-a** into pokes table.
  - '**LOCAL**' signifies that the input file is on the local file system. If '**LOCAL**' is omitted then it looks for the file in HDFS.
  - The keyword '**OVERWRITE**' signifies that existing data in the table is deleted. If the '**OVERWRITE**' keyword is omitted, data files are appended to existing data sets.
- NOTES:
  - **NO** verification of data against the schema is performed by the load command.
  - If the file is in hdfs, it is moved into the Hive-controlled file system namespace.
  - The root of the Hive directory is specified by the option **hive.metastore.warehouse.dir** in **hive-default.xml**. We advise users to create this directory before trying to create tables via Hive.

# DML Operations

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities

## Browse Directory

/user/hive/warehouse/pokes

Show 25 entries

Permission	Owner
-rw-r--r--	chenhai

Showing 1 to 1 of 1 entries

Hadoop, 2021.

File information - kv1.txt

Download Head the file (first 32K) Tail the file (last 32K)

Block information -- Block 0

Block ID: 1073741825  
Block Pool ID: BP-591783681-127.0.0.1-1640354044142  
Generation Stamp: 1001  
Size: 5812  
Availability:

- localhost

File contents

```
238val_238
86val_86
311val_311
27val_27
165val_165
409val_409
255val_255
278val_278
```

Close

- Data Load

- ```
hive> LOAD DATA LOCAL INPATH './examples/files/kv2.txt' OVERWRITE INTO TABLE invites PARTITION (ds='2008-08-15');
```
- ```
hive> LOAD DATA LOCAL INPATH './examples/files/kv3.txt' OVERWRITE INTO TABLE invites PARTITION (ds='2008-08-08');
```
- The two LOAD statements above load data into **two different partitions** of the table **invites**.
- Table invites must be created as partitioned by the key ds for this to succeed.
  
- ```
hive> LOAD DATA INPATH '/user/myname/kv2.txt' OVERWRITE INTO TABLE invites PARTITION (ds='2008-08-15');
```
- The above command will load data from an HDFS file/directory to the table.
- Note that loading data from HDFS will result in moving the file/directory. As a result, the operation is almost instantaneous.

## Browse Directory

| /user/hive/warehouse/invites |            |             |            |      |               |             |            |               | Go! |         |  |  |  |
|------------------------------|------------|-------------|------------|------|---------------|-------------|------------|---------------|-----|---------|--|--|--|
| Show 25 entries              |            |             |            |      |               |             |            |               |     | Search: |  |  |  |
| <input type="checkbox"/>     | Permission | Owner       | Group      | Size | Last Modified | Replication | Block Size | Name          |     |         |  |  |  |
| <input type="checkbox"/>     | drwxr-xr-x | chenhaopeng | supergroup | 0 B  | Dec 24 22:09  | 0           | 0 B        | ds=2008-08-08 |     |         |  |  |  |
| <input type="checkbox"/>     | drwxr-xr-x | chenhaopeng | supergroup | 0 B  | Dec 24 22:09  | 0           | 0 B        | ds=2008-08-15 |     |         |  |  |  |

Showing 1 to 2 of 2 entries

Previous 1 Next

Hadoop, 2021.

- The Hive query operations are documented in [Select](#).
- **Example Queries**
  - Some example queries are shown below. They are available in [build/dist/examples/queries](#). More are available in the Hive sources at [ql/src/test/queries/positive](#).
- **SELECTS and FILTERS**
  - `hive> SELECT a.foo FROM invites a WHERE a.ds='2008-08-15';`
  - selects column 'foo' from all rows of partition ds=2008-08-15 of the invites table. The results are not stored anywhere, but are displayed on the console.



```
[hive> SELECT a.foo FROM invites a WHERE a.ds='2008-08-15';
OK
474
281
179
291
62
271
217
135
167]
```

- Note that in all the examples that follow, **INSERT** (into a Hive table, local directory or HDFS directory) is optional.
  - `hive> INSERT OVERWRITE DIRECTORY '/tmp/hdfs_out' SELECT a.* FROM invites a WHERE a.ds='2008-08-15';`
  - selects all rows from partition ds=2008-08-15 of the invites table **into an HDFS directory**.
  - The result data is **in files (depending on the number of mappers)** in that directory.  
NOTE: partition columns if any are selected by the use of \*. They can also be specified in the projection clauses.
  - Partitioned tables must always have a partition selected in the **WHERE** clause of the statement.

# SQL Operations

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

## Browse Directory

Show 25 entries

| Permission | Owner   |
|------------|---------|
| -rw-r--r-- | chenhao |

Showing 1 to 1 of 1 entries

Hadoop, 2021.

/tmp/hdfs\_out

File information - 000000\_0

Download Head the file (first 32K) Tail the file (last 32K)

Block information -- Block 0

Block ID: 1073741832  
Block Pool ID: BP-591783681-127.0.0.1-1640354044142  
Generation Stamp: 1008  
Size: 11291  
Availability:

- localhost

File contents

```
474val_4752008-08-15
281val_2822008-08-15
179val_1802008-08-15
291val_2922008-08-15
62val_632008-08-15
271val_2722008-08-15
217val_2182008-08-15
135val_1362008-08-15
```

**Close**

lock Size Name  
28 MB 000000\_0

Previous 1 Next

- **SELECTS and FILTERS**

- `hive> INSERT OVERWRITE LOCAL DIRECTORY '/tmp/local_out' SELECT a.* FROM pokes a;`
- selects all rows from pokes table into a local directory.
  
- `hive> INSERT OVERWRITE TABLE events SELECT a.* FROM profiles a;`
- `hive> INSERT OVERWRITE TABLE events SELECT a.* FROM profiles a WHERE a.key < 100;`
- `hive> INSERT OVERWRITE LOCAL DIRECTORY '/tmp/reg_3' SELECT a.* FROM events a;`
- `hive> INSERT OVERWRITE DIRECTORY '/tmp/reg_4' select a.invites, a.pokes FROM profiles a;`
- `hive> INSERT OVERWRITE DIRECTORY '/tmp/reg_5' SELECT COUNT(*) FROM invites a WHERE a.ds='2008-08-15';`
- `hive> INSERT OVERWRITE DIRECTORY '/tmp/reg_5' SELECT a.foo, a.bar FROM invites a;`
- `hive> INSERT OVERWRITE LOCAL DIRECTORY '/tmp/sum' SELECT SUM(a.pc) FROM pc1 a;`
- selects the **sum** of a column. The **avg**, **min**, or **max** can also be used.

- **GROUP BY**

- ```
hive> FROM invites a INSERT OVERWRITE TABLE events SELECT a.bar, count(*) WHERE a.foo > 0 GROUP BY a.bar;
```
- ```
hive> INSERT OVERWRITE TABLE events SELECT a.bar, count(*) FROM invites a WHERE a.foo > 0 GROUP BY a.bar;
```
- Note that for versions of Hive which don't include [HIVE-287](#), you'll need to use COUNT(1) in place of COUNT(\*).

- **JOIN**

- ```
hive> FROM pokes t1 JOIN invites t2 ON (t1.bar = t2.bar) INSERT OVERWRITE TABLE events SELECT t1.bar, t1.foo, t2.foo;
```

- **MULTITABLE INSERT**

- FROM src
- INSERT OVERWRITE TABLE dest1 SELECT src.\* WHERE src.key < 100
- INSERT OVERWRITE TABLE dest2 SELECT src.key, src.value WHERE src.key >= 100 and src.key < 200
- INSERT OVERWRITE TABLE dest3 PARTITION(ds='2008-04-08', hr='12') SELECT src.key WHERE src.key >= 200 and src.key < 300
- INSERT OVERWRITE LOCAL DIRECTORY '/tmp/dest4.out' SELECT src.value WHERE src.key >= 300;

- **STREAMING**

- hive> FROM invites a INSERT OVERWRITE TABLE events SELECT TRANSFORM(a.foo, a.bar) AS (oof, rab) USING '/bin/cat' WHERE a.ds > '2008-08-09';
- This streams the data in the map phase through the script /bin/cat (like Hadoop streaming).  
Similarly – streaming can be used on the reduce side (please see the [Hive Tutorial](#) for examples).

- **MovieLens User Ratings**

- First, create a table with tab-delimited text file format:
- ```
CREATE TABLE u_data (
    - userid INT,
    - movieid INT,
    - rating INT,
    - unixtime STRING)
    - ROW FORMAT DELIMITED
    - FIELDS TERMINATED BY '\t'
    - STORED AS TEXTFILE;
```

- **MovieLens User Ratings**

- Then, download the data files from **MovieLens 100k** on the [GroupLens datasets](#) page (which also has a README.txt file and index of unzipped files):
  - `wget http://files.grouplens.org/datasets/movielens/ml-100k.zip`
  - or:
  - `curl --remote-name http://files.grouplens.org/datasets/movielens/ml-100k.zip`
- Unzip the data files:
  - `unzip ml-100k.zip`
- And load u.data into the table that was just created:
  - `LOAD DATA LOCAL INPATH '<path>/u.data' OVERWRITE INTO TABLE u_data;`

# Simple Example Use Cases

The screenshot shows a Hadoop file browser interface. At the top, there's a navigation bar with tabs: Hadoop, Overview, Datanodes, Datanode Volume Failures, Snapshot, Startup Progress, and Utilities. Below the navigation bar, the main area has a title "Browse Directory" and a path bar showing "/user/hive/warehouse/u\_data". A red box highlights this path bar.

Below the path bar, there are filters for "Show 25 entries" and sorting options for "Permission" and "Owner". The results table shows one entry: a file named "u.data" owned by "chenhadu" with permission "-rw-r--r--". A red box highlights the file name "u.data" in the table.

A modal window titled "File information - u.data" is open. It contains sections for "Download", "Head the file (first 32K)", and "Tail the file (last 32K)". Under "Block information", it shows Block ID: 1073741829, Block Pool ID: BP-591783681-127.0.0.1-1640354044142, Generation Stamp: 1005, Size: 1979173, and Availability: "localhost". A red box highlights the "Block information" section.

At the bottom of the modal, there's a "File contents" section containing the following data:

| File contents                                                                                                                                                                      |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 196 242 3 881250949<br>186 302 3 891717742<br>22 377 1 878887116<br>244 51 2 880606923<br>166 346 1 886397596<br>298 474 4 884182806<br>115 265 2 881171488<br>253 465 5 891628467 |

A red box highlights the "File contents" table. At the bottom right of the modal, there's a "Close" button.

# Simple Example Use Cases

- Count the number of rows in table u\_data:

```
SELECT COUNT(*) FROM u_data;
```

```
hive> SELECT COUNT(*) FROM u_data;
Query ID = chenhaopeng_20211225212223_69a61947-9a84-455e-a009-85a50cbcb817
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (local Hadoop)
2021-12-25 21:22:24,783 Stage-1 map = 100%,  reduce = 100%
Ended Job = job_local1383216318_0004
MapReduce Jobs Launched:
Stage-Stage-1:  HDFS Read: 10286786 HDFS Write: 2381094 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
100000
Time taken: 1.58 seconds, Fetched: 1 row(s)
hive> □
```

# Simple Example Use Cases

- Now we can do some complex data analysis on the table u\_data:
- Create `weekday_mapper.py`:

```
import sys
import datetime
for line in sys.stdin:
    line = line.strip()
    userid, movieid, rating, unixtime = line.split('\t')
    weekday = datetime.datetime.fromtimestamp(float(unixtime)).isoweekday()
    print ('\t'.join([userid, movieid, rating, str(weekday)]))
```

# Simple Example Use Cases

- Use the mapper script:

```
CREATE TABLE u_data_new (
    userid INT,
    movieid INT,
    rating INT,
    weekday INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t';

add FILE weekday_mapper.py;

INSERT OVERWRITE TABLE u_data_new
SELECT
    TRANSFORM (userid, movieid, rating, unixtime)
    USING 'python weekday_mapper.py'
    AS (userid, movieid, rating, weekday)
FROM u_data;

SELECT weekday, COUNT(*)
FROM u_data_new
GROUP BY weekday;
```

# Simple Example Use Cases

```
apache-hive-3.1.2-bin — java -Dproc_jar -Djava.library.path=/Users/chen...  
hive>  
    > SELECT weekday, COUNT(*)  
    > FROM u_data_new  
    > GROUP BY weekday;  
[ [ ]  
Query ID = chenhaopeng_20211224222723_d9018a35-1b4d-47f9-abd8-e9b333b9ccd0  
Total jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks not specified. Estimated from input data size: 1  
In order to change the average load for a reducer (in bytes):  
    set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
    set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
    set mapreduce.job.reduces=<number>  
Job running in-process (local Hadoop)  
2021-12-24 22:27:26,526 Stage-1 map = 100%,  reduce = 100%  
Ended Job = job_local1474112274_0002  
MapReduce Jobs Launched:  
Stage-Stage-1:  HDFS Read: 6316858 HDFS Write: 2358512 SUCCESS  
Total MapReduce CPU Time Spent: 0 msec  
OK  
1      12254  
2      13579  
3      14430  
4      15114  
5      14743  
6      18229  
7      11651  
Time taken: 2.97 seconds, Fetched: 7 row(s)
```

# An Example

- Let's see how to use Hive to run a query on the weather dataset.
- The first step is to load the data into Hive's managed storage.
- Just like an RDBMS, Hive organizes its data into tables. We create a table to hold the weather data using the **CREATE TABLE** statement:

**CREATE TABLE** records

(year STRING, temperature INT, quality INT)

ROW FORMAT DELIMITED

FIELDS TERMINATED BY '\t';

# An Example

- Next we can populate Hive with the data.

- This is just a small sample, for exploratory purposes:

```
LOAD DATA LOCAL INPATH 'input/ncdc/micro-tab/sample.txt'  
OVERWRITE INTO TABLE records;
```

- Running this command tells Hive to put the specified local file in its warehouse directory.
  - There is no attempt, for example, to parse the file and store it in an internal database format, since Hive does not mandate any particular file format.
  - Files are stored verbatim: they are not modified by Hive.

# An Example

- In this example, we are storing Hive tables on the local filesystem (`fs.default.name` is set to its default value of `file:///`).
  - Tables are stored as directories under Hive's warehouse directory, which is controlled by the `hive.metastore.warehouse.dir`, and defaults to `/user/hive/warehouse`.
  - Thus, the files for the records table are found in the `/user/hive/warehouse/records` directory on the local filesystem:  
`% ls /user/hive/warehouse/record/sample.txt`
- In this case, there is only one file, `sample.txt`, but in general there can be more, and Hive will read all of them when querying the table.

# An Example

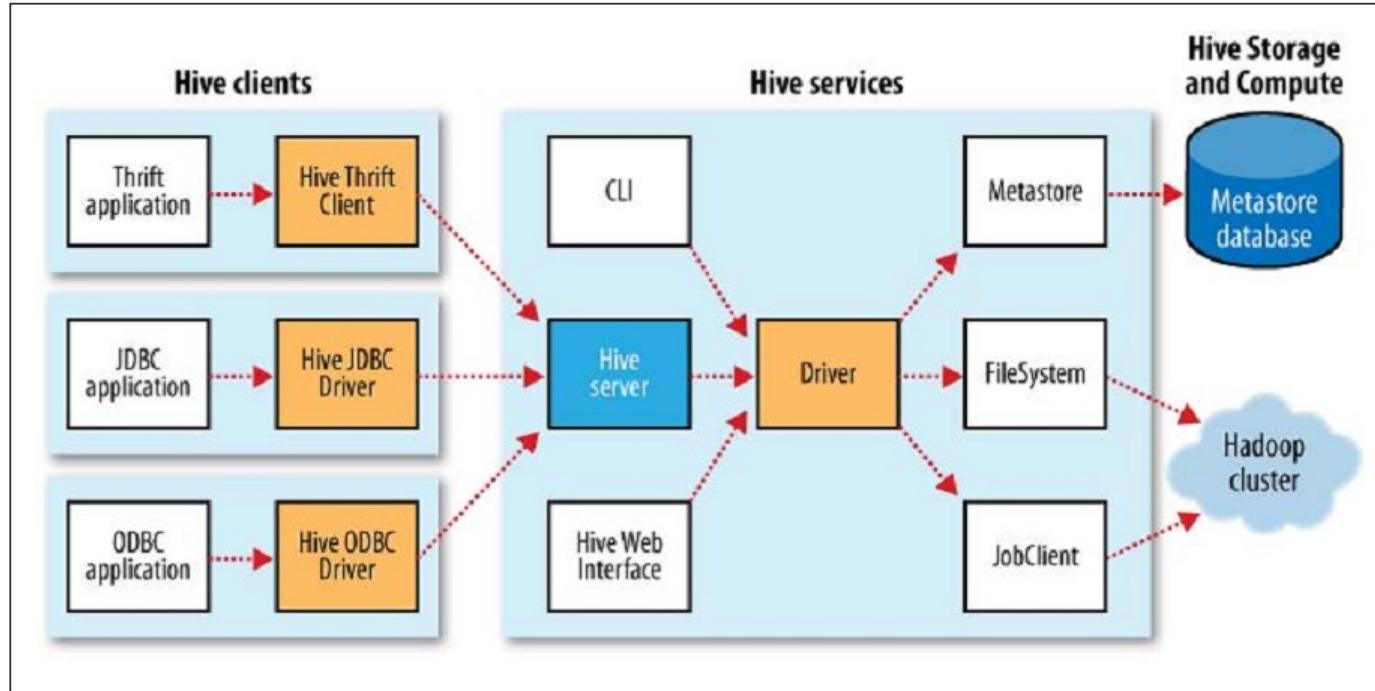
- Now that the data is in Hive, we can run a query against it:

```
hive> SELECT year, MAX(temperature)
    > FROM records
    > WHERE temperature != 9999
    > AND (quality = 0 OR quality = 1 OR quality = 4 OR quality = 5 OR quality = 9)
    > GROUP BY year;
```

1949 111

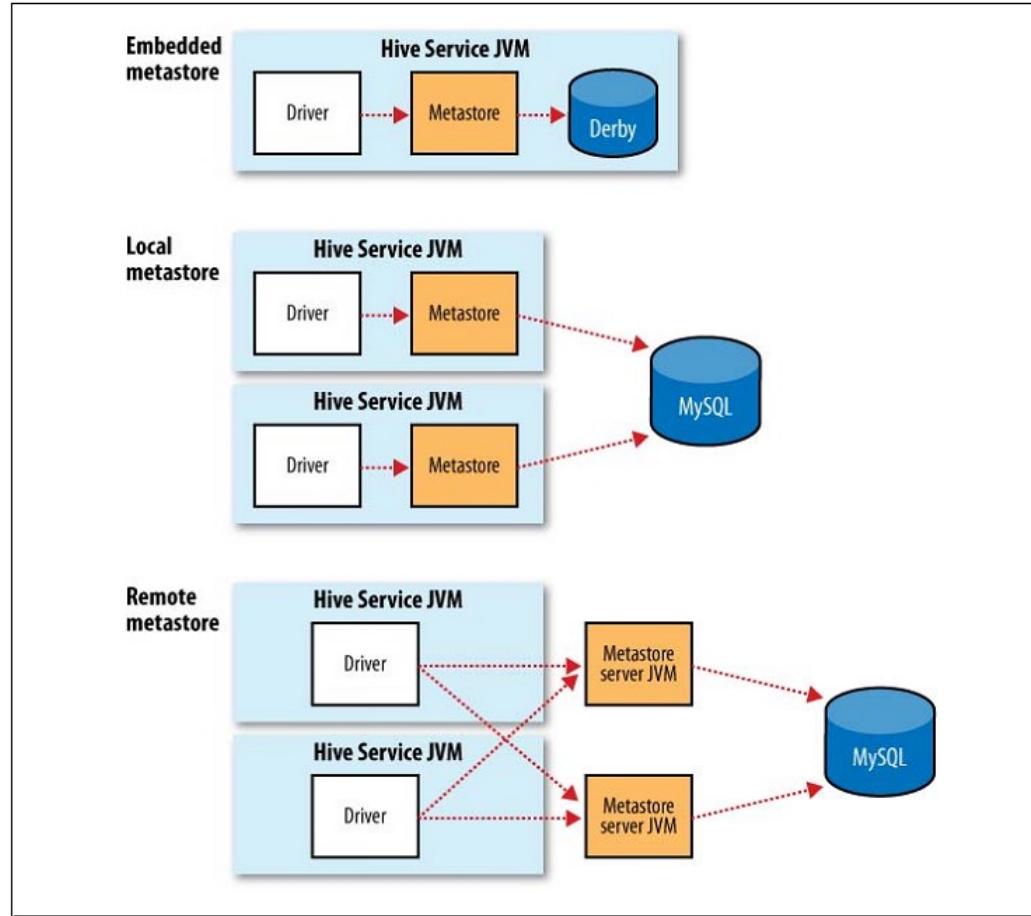
1950 22

- The Hive shell is only one of several services that you can run using the `hive` command.
- Type `hive -service help` to get a list of available service names; the most useful are described below.
  - `cli`
    - The command line interface to Hive (the shell). This is the default service.
  - `hiveserver`
    - Runs Hive as a server exposing a Thrift service, enabling access from a range of clients written in different languages.
  - `hwi`
    - The Hive Web Interface.
  - `jar`
    - The Hive equivalent to `hadoop jar`, a convenient way to run Java applications that includes both Hadoop and Hive classes on the classpath.
  - `metastore`
    - By default, the metastore is run in the same process as the Hive service. Using this service, it is possible to run the metastore as a standalone (remote) process.



- The metastore is the central repository of Hive metadata.
  - The metastore is divided into two pieces: a service and the backing store for the data.
  - By default, the metastore service runs in the **same JVM** as the Hive service and contains an embedded Derby database instance backed by the local disk.
  - It also supports multiple sessions (and therefore multiple users) is to use a standalone database. This configuration is referred to as a local metastore, since the metastore service still runs in the same process as the Hive service, but connects to a database running in **a separate process**, either on the same machine or on a remote machine.
  - Going a step further, there's another metastore configuration called a **remote metastore**, where one or more metastore servers run in separate processes to the Hive service.

# The Metastore



- Schema on Read Versus Schema on Write

- In a traditional database, a table's schema is enforced at data load time. If the data being loaded doesn't conform to the schema, then it is rejected.
  - This design is sometimes called **schema on write**, since the data is checked against the schema when it is written into the database.
- Hive, on the other hand, doesn't verify the data when it is loaded, but rather when a query is issued.
  - This is called **schema on read**.
- There are trade-offs between the two approaches.
  - Schema on read makes for a very fast initial load, since the data does not have to be read, parsed, and serialized to disk in the database's internal format.
  - Schema on write makes query time performance faster, since the database can index columns and perform compression on the data.

- Hive's SQL dialect, called HiveQL, does **not** support the full SQL-92 specification.

| Feature                   | SQL                                                                            | HiveQL                                                              | References                                                                                                                          |
|---------------------------|--------------------------------------------------------------------------------|---------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------|
| Updates                   | UPDATE, INSERT,<br>DELETE                                                      | INSERT OVERWRITE<br>TABLE (populates whole ta-<br>ble or partition) | <a href="#">"INSERT OVERWRITE TA-<br/>BLE" on page 392</a> , <a href="#">"Updates, Transac-<br/>tions, and Indexes" on page 376</a> |
| Transactions              | Supported                                                                      | Not supported                                                       |                                                                                                                                     |
| Indexes                   | Supported                                                                      | Not supported                                                       |                                                                                                                                     |
| Latency                   | Sub-second                                                                     | Minutes                                                             |                                                                                                                                     |
| Data types                | Integral, floating point,<br>fixed point, text and binary<br>strings, temporal | Integral, floating point, bo-<br>lean, string, array, map, struct   | <a href="#">"Data Types" on page 378</a>                                                                                            |
| Functions                 | Hundreds of built-in<br>functions                                              | Dozens of built-in functions                                        | <a href="#">"Operators and Functions"<br/>on page 380</a>                                                                           |
| Multitable inserts        | Not supported                                                                  | Supported                                                           | <a href="#">"Multitable insert" on page 393</a>                                                                                     |
| Create table as<br>select | Not valid SQL-92, but found<br>in some databases                               | Supported                                                           | <a href="#">"CREATE TABLE...AS SE-<br/>LECT" on page 394</a>                                                                        |

- Hive's SQL dialect, called HiveQL, does **not** support the full SQL-92 specification.

| Feature          | SQL                                                                                     | HiveQL                                                                                                                               | References                                                                                                |
|------------------|-----------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------|
| Select           | SQL-92                                                                                  | Single table or view in the FROM clause. SORT_BY for partial ordering. LIMIT to limit number of rows returned. HAVING not supported. | <a href="#">"Querying Data" on page 395</a>                                                               |
| Joins            | SQL-92 or variants (join tables in the FROM clause, join condition in the WHERE clause) | Inner joins, outer joins, semi joins, map joins. SQL-92 syntax, with hinting.                                                        | <a href="#">"Joins" on page 397</a>                                                                       |
| Subqueries       | In any clause. Correlated or noncorrelated.                                             | Only in the FROM clause. Correlated subqueries not supported                                                                         | <a href="#">"Subqueries" on page 400</a>                                                                  |
| Views            | Updatable. Materialized or nonmaterialized.                                             | Read-only. Materialized views not supported                                                                                          | <a href="#">"Views" on page 401</a>                                                                       |
| Extension points | User-defined functions. Stored procedures.                                              | User-defined functions. MapReduce scripts.                                                                                           | <a href="#">"User-Defined Functions" on page 402</a> ,<br><a href="#">"MapReduce Scripts" on page 396</a> |

- A Hive table is logically made up of the data being stored and the associated metadata describing the layout of the data in the table.
  - The data typically resides in HDFS, although it may reside in any Hadoop filesystem, including the local filesystem or S3.
  - Hive stores the metadata in a relational database—and not in HDFS.
- Managed Tables and External Tables
  - When you create a table in Hive, by default Hive will manage the data, which means that Hive moves the data into its warehouse directory.
  - Alternatively, you may create an external table, which tells Hive to refer to the data that is at an existing location outside the warehouse directory.

- Partitions and Buckets
  - Hive organizes tables into partitions, a way of dividing a table into coarse-grained parts based on the value of a partition column.
  - Partitions are defined at table creation time using the **PARTITIONED BY** clause, which takes a list of column definitions.
    - For the hypothetical log files example, we might define a table with records comprising a timestamp and the log line itself:

```
CREATE TABLE logs (ts BIGINT, line STRING)
PARTITIONED BY (dt STRING, country STRING);
```
  - When we load data into a partitioned table, the partition values are specified explicitly:

```
LOAD DATA LOCAL INPATH 'input/hive/partitions/file1'
INTO TABLE logs
PARTITION (dt='2001-01-01', country='GB');
```

- Partitions and Buckets
  - Tables or partitions may further be subdivided into buckets, to give extra structure to the data that may be used for more efficient queries.
  - There are two reasons why you might want to organize your tables (or partitions) into buckets.
    - The first is to enable more efficient queries. Bucketing imposes extra structure on the table, which Hive can take advantage of when performing certain queries.
    - The second reason to bucket a table is to make sampling more efficient. When working with large datasets, it is very convenient to try out queries on a fraction of your dataset while you are in the process of developing or refining them.

```
CREATE TABLE bucketed_users (id INT, name STRING)  
CLUSTERED BY (id) SORTED BY (id ASC)INTO 4 BUCKETS;
```

# Querying Data

- Sorting and Aggregating

```
hive> FROM records2
    > SELECT year, temperature
    > DISTRIBUT BY year
    > SORT BY year ASC, temperature DESC;
1949 111
1949 78
1950 22
1950 0
1950 -11
```

- MapReduce Scripts

- Using an approach like Hadoop Streaming, the **TRANSFORM**, **MAP**, and **REDUCE** clauses make it possible to invoke an external script or program from Hive.

```
FROM (
    FROM records2
    MAP year, temperature, quality
    USING 'is_good_quality.py'
    AS year, temperature) map_output
REDUCE year, temperature
USING 'max_temperature_reduce.py'
AS year, temperature;
```

- Hive supports several file formats:
  - Text File
  - SequenceFile
  - [RCFile](#)
  - [Avro Files](#)
  - [ORC Files](#)
  - [Parquet](#)
  - Custom INPUTFORMAT and OUTPUTFORMAT
- The [`hive.default.fileformat`](#) configuration parameter
  - determines the format to use if it is not specified in a [CREATE TABLE](#) or [ALTER TABLE](#) statement.
  - Text file is the parameter's **default** value.

- RCFfile

- RCFfile (**Record Columnar File**) is a data placement structure designed for MapReduce-based data warehouse systems. Hive added the RCFfile format in version 0.6.0.
- RCFfile stores table data in a flat file consisting of **binary key/value pairs**. It first partitions rows **horizontally** into **row splits**, and then it **vertically partitions each row split in a columnar way**.
- RCFfile stores the metadata of a row split as the key part of a record, and all the data of a row split as the value part.
- RCFfile combines the advantages of both **row-store** and **column-store** to satisfy the need for fast data loading and query processing, efficient use of storage space, and adaptability to highly dynamic workload patterns.
- As row-store, RCFfile guarantees that **data in the same row are located in the same node**.
- As column-store, RCFfile can exploit **column-wise data compression and skip unnecessary column reads**.

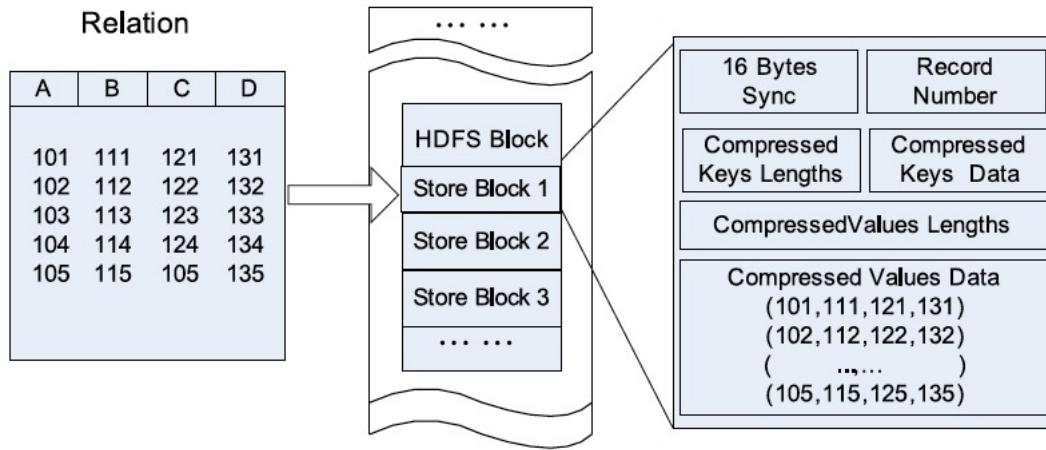


Fig. 1: An example of row-store in an HDFS block.

- 2011 ICDE conference paper "[RCFile: A Fast and Space-efficient Data Placement Structure in MapReduce-based Warehouse Systems](#)" by Yongqiang He, Rubao Lee, Yin Huai, Zheng Shao, Namit Jain, Xiaodong Zhang, and Zhiwei Xu

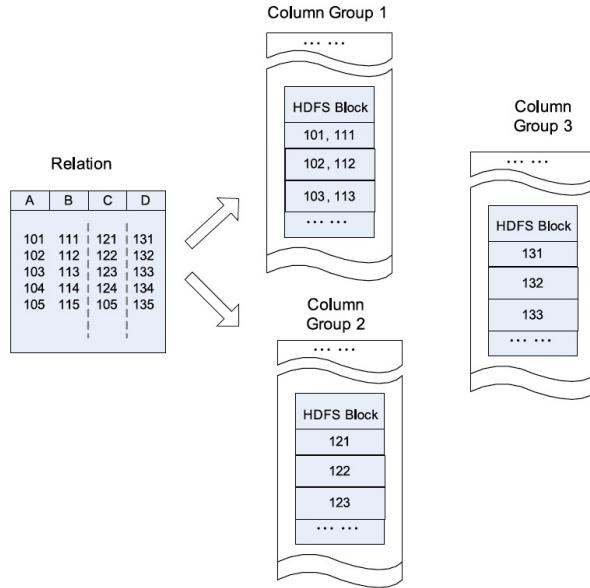


Fig. 2: An example of column-group in an HDFS block. The four columns are stored into three column groups, since column A and B are grouped in the first column group.

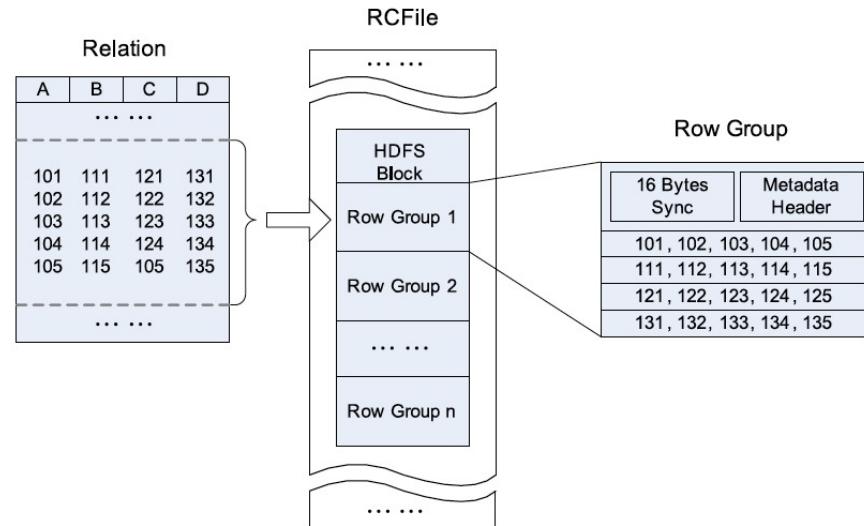
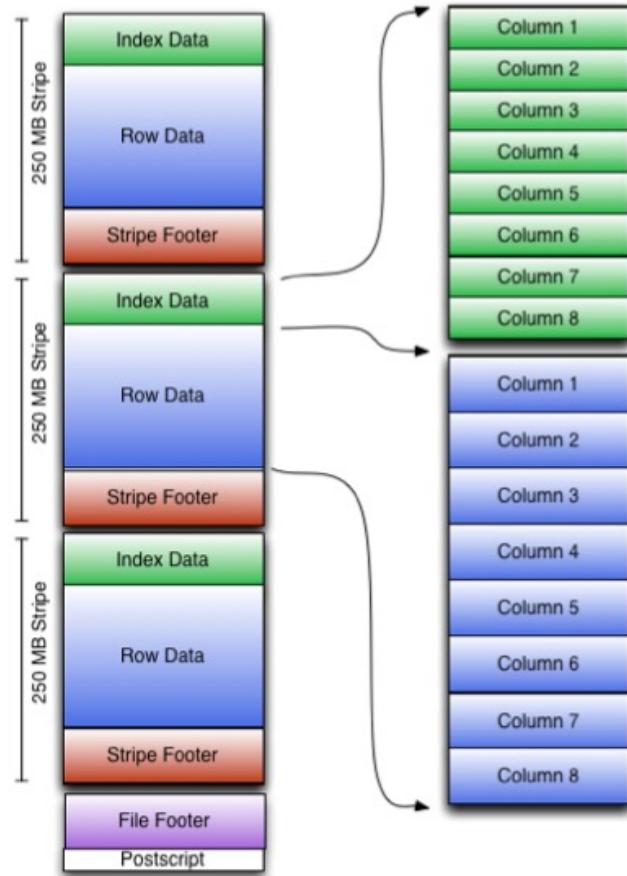


Fig. 3: An example to demonstrate the data layout of RCFfile in an HDFS block.

- 2011 ICDE conference paper "[RCFile: A Fast and Space-efficient Data Placement Structure in MapReduce-based Warehouse Systems](#)" by Yongqiang He, Rubao Lee, Yin Huai, Zheng Shao, Namit Jain, Xiaodong Zhang, and Zhiwei Xu

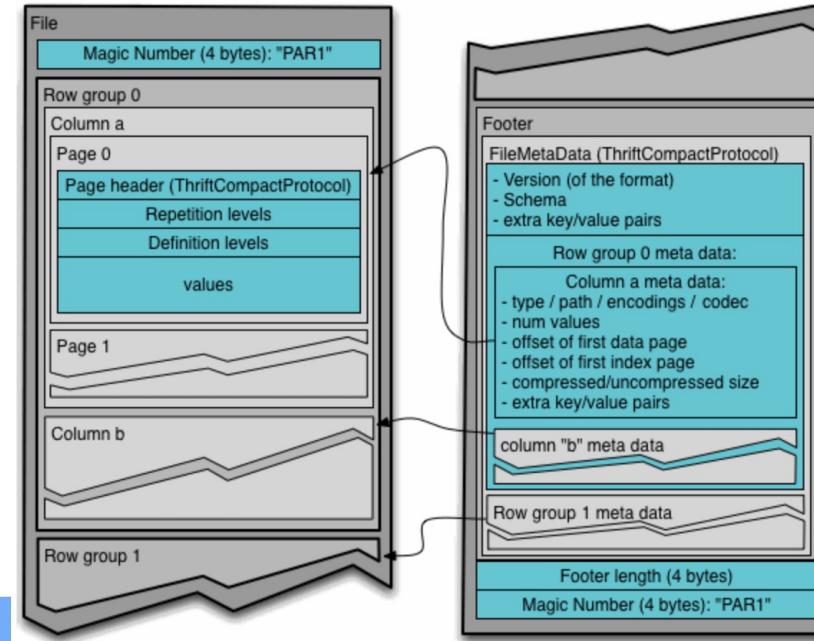
- The *Optimized Row Columnar* ([ORC](#)) file format provides a highly efficient way to store Hive data.
- An ORC file contains groups of row data called **stripes**, along with auxiliary information in a **file footer**.
  - At the end of the file a **postscript** holds compression parameters and the size of the compressed footer.
  - The default stripe size is **250 MB**. Large stripe sizes enable large, efficient reads from HDFS.
  - The file footer contains a list of stripes in the file, the number of rows per stripe, and each column's data type. It also contains column-level aggregates count, min, max, and sum.

- File Structure



- Apache Parquet

- is a [columnar storage](#) format available to any project in the Hadoop ecosystem, regardless of the choice of data processing framework, data model or programming language.
- <http://parquet.incubator.apache.org/>



- Keeping data **compressed** in Hive tables has, in some cases,
  - been known to give **better performance** than uncompressed storage; both in terms of **disk usage** and **query performance**.
- You can import text files compressed with **Gzip** or **Bzip2** directly into a table stored as **TextFile**.
  - The compression will be detected automatically and the file will be **decompressed on-the-fly during query execution**. For example:

```
CREATE TABLE raw (line STRING)
  ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY '\n';
```

```
LOAD DATA LOCAL INPATH '/tmp/weblogs/20090603-access.log.gz' INTO TABLE raw;
```

- The table 'raw' is stored as a **TextFile**, which is the **default** storage.
  - However, in this case Hadoop will **not be able to** split your file into chunks/blocks and run multiple maps in parallel. This can cause **underutilization** of your cluster's 'mapping' power.

# File Compression

- The **recommended** practice is to insert data into another table, which is stored as a **SequenceFile**.
  - A **SequenceFile** can be split by Hadoop and distributed across map jobs whereas a **GZIP file cannot** be. For example:

```
CREATE TABLE raw (line STRING)
```

```
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY '\n';
```

```
CREATE TABLE raw_sequence (line STRING)
```

```
STORED AS SEQUENCEFILE;
```

```
LOAD DATA LOCAL INPATH '/tmp/weblogs/20090603-access.log.gz' INTO TABLE raw;
```

```
SET hive.exec.compress.output=true;
```

```
SET io.seqfile.compression.type=BLOCK; -- NONE/RECORD/BLOCK (see below)
```

```
INSERT OVERWRITE TABLE raw_sequence SELECT * FROM raw;
```

- The value for **io.seqfile.compression.type** determines how the compression is performed.
  - Record compresses each value individually while **BLOCK** buffers up **1MB (default)** before doing compression.

# LZO Compression

- LZO
  - is a **lossless data compression** library that favors speed over compression ratio.  
See <http://www.oberhumer.com/opensource/lzo> and <http://www.lzop.org> for general information about LZO.
- Imagine a simple data file that has three columns
  - id
  - first name
  - last name
- Let's populate a data file containing 4 records:  
`19630001 john lennon`  
`19630002 paul mccartney`  
`19630003 george harrison`  
`19630004 ringo starr`

- Let's call the data file `/path/to/dir/names.txt`.
  - In order to make it into an LZO file, we can use the `lzop` utility and it will create a `names.txt.lzo` file.
  - Now copy the file `names.txt.lzo` to HDFS.
- Hive Queries**
- Option 1: Directly Create LZO Files**
  - Directly create LZO files as the output of the Hive query.
  - Use `lzop` command utility or your custom Java to generate `.lzo.index` for the `.lzo` files.
- Hive Query Parameters**

```
SET mapreduce.output.fileoutputformat.compress.codec=com.hadoop.compression.lzo.LzoCodec;
SET hive.exec.compress.output=true;
SET mapreduce.output.fileoutputformat.compress=true;
```
- For example:

```
hive -e "SET
mapreduce.output.fileoutputformat.compress.codec=com.hadoop.compression.lzo.LzoCodec;
SET hive.exec.compress.output=true;SET mapreduce.output.fileoutputformat.compress=true;
<query-string>"
```

- **Hive Queries**
- **Option 2: Write Custom Java to Create LZO Files**

1. Create **text** files as the output of the Hive query.
2. Write custom **Java** code to
  - a. convert Hive query generated text files to **.lzo** files
  - b. generate **.lzo.index** files for the **.lzo** files generated above

- **Hive Query Parameters**

```
SET hive.exec.compress.output=false
```

```
SET mapreduce.output.fileoutputformat.compress=false
```

- For example:

```
hive -e "SET hive.exec.compress.output=false;SET  
mapreduce.output.fileoutputformat.compress=false;<query-string>"
```

- Apache Hive
  - <https://hive.apache.org>
- hadoop : mkdir: 'input': No such file or directory 相关问题
  - [https://blog.csdn.net/ludonqin/article/details/51396187?depth\\_1-utm\\_source=distribute.pc\\_relevant.none-task&utm\\_source=distribute.pc\\_relevant.none-task](https://blog.csdn.net/ludonqin/article/details/51396187?depth_1-utm_source=distribute.pc_relevant.none-task&utm_source=distribute.pc_relevant.none-task)
- 启动hive报错： java.lang.NoSuchMethodError:  
[com.google.common.base.Preconditions.checkNotNull\(ZLjava/lang/String;Ljava/lang/Object;\)V \(已解决\)](com.google.common.base.Preconditions.checkNotNull(ZLjava/lang/String;Ljava/lang/Object;)V)
- Hadoop: The Definitive Guide
  - By Tom White
  - O'Reilly Publishing



# Thank You!