



Remote Procedure Call

IPADS, Shanghai Jiao Tong University

https://www.sjtu.edu.cn

Credits: Rong Chen@IPADS

Distributed system that support large-scale websites

Caching server Caching server Caching server Large-scale websites are composed of Caching Caching Caching different distributed systems Request processing, data storage Distributed caching Database server Database server **Database Database Users Application #1** user, price user, price generate the page Application server Distributed database Internet Load **Application #2** File server File server Balance File server add the order File: File: File: image image image Application server Distributed file system Message queue

Example: extend a single-node filesystem to a distributed one

Caching server Caching server Caching server Large-scale websites are composed of Caching Caching different distributed systems Request processing, data storage Distributed caching **How each system communicates?** Database server Database server **Database Database Users Application #1** user, price user, price generate the page Application server Distributed database Internet Load **Application #2** Balance File server File server add the order File: File: File: image Application server Message queue Distributed file system

Recall: The architecture of LAMP cannot scale!

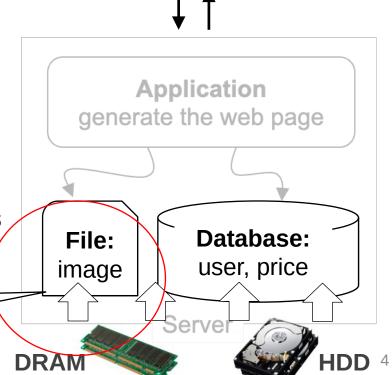
1. The **disk & memory** of one server cannot store massive amount of data

− DRAM: 64 – 256 GB

− HDD: 2 − 40 TB

 Facebook has more than 1 billion of images uploaded weekly, Baidu stores tens of billions of web pages

See previous lectures



Step #1 for scalability: disaggregating application & data

Application: handles application logic

Can be scaled with more CPUs

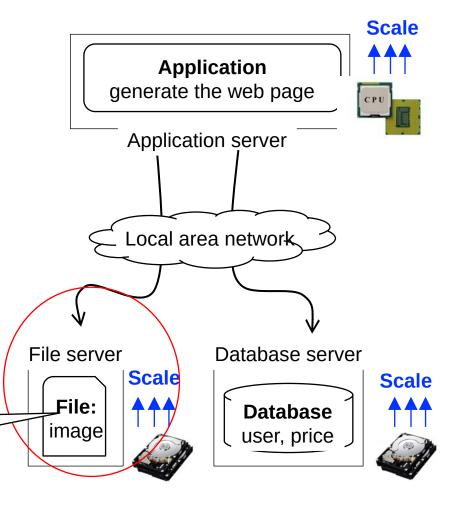
Database: requires reading/writing disk & cache

 Can be scaled with faster disks & larger memory

File system: store large bulks of data

- E.g., images, videos
- Can be scaled with faster disks

How to do so?



Remote Procedure Call (RPC)

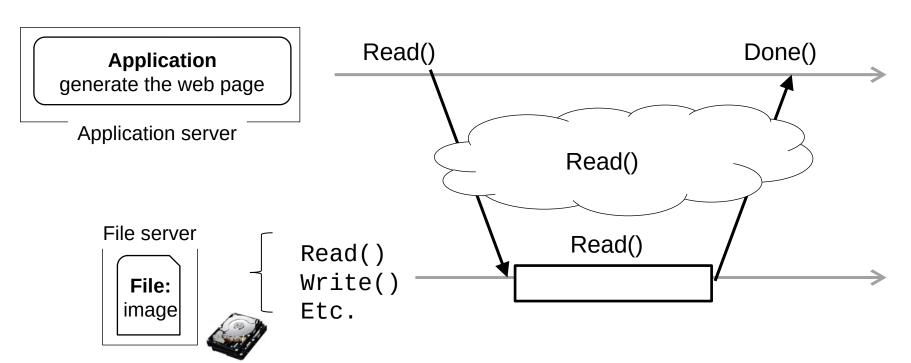
"Remote procedure calls (RPC) appear to be a useful paradigm."

— Birrel & Nelson, 1984

Andrew Birrell, Bruce Nelson, "Implementing Remote Procedure Calls". ACM Transactions on Computer Systems (TOCS), 2(1), 1984

filesystem + RPC, a form of distributed filesystem

Calling a function on a remote server like a local one!



RPC vs. Sockets API

The sockets interface forces a read/write mechanism

read(), write(), sendmsg(), etc.

Programming is often easier with a functional interface

To make distributed computing look more like centralized computing

Goal of RPC:

It should appear to the programmer that a normal call is taking place

Idea: build the RPC atop of the socket interface

 Hide the construction of messages and remote invocation logic from the developers

RPC: Remote Procedure Call

Allow a procedure to execute in another address space without coding the details for the remote interaction

RPC History

- Idea goes back in 1976
- Sun's RPC: first popular implementation on Unix
 - Used as the basis for NFS
- Many modern RPC frameworks: gRPC, bRPC, etc.
- RMI (Remote Method Invocation)
 - Object-oriented version of RPC, e.g. in Java







Example of RPC

```
1procedure MEASURE (func)1procedure GET_TIME (units)2start \leftarrow GET_TIME (SECONDS)2time \leftarrow CLOCK3func () // invoke the function3time \leftarrow CONVERT_TO_UNITS (time, units)4end \leftarrow GET_TIME (SECONDS)4return time5return end - start
```

The implementation of GET TIME.

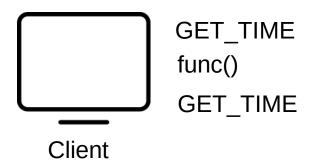
Suppose we want to measure the execution time of *func()*

Assumption:

Only the server has the implementation of GET_TIME

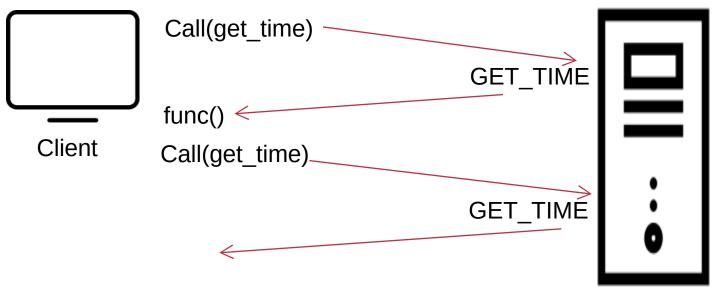
How can the client call server's **GET_TIME**?

Call GET_TIME in a single machine case



Single-machine call vs. distributed call (RPC)

```
 \begin{array}{lll} 1 & \textbf{procedure} \; \texttt{MEASURE} \; (\textit{func}) & 1 & \textbf{procedure} \; \texttt{GET\_TIME} \; (\textit{units}) \\ 2 & \textit{start} \leftarrow \texttt{GET\_TIME} \; (\texttt{SECONDS}) & 2 & \textit{time} \leftarrow \texttt{CLOCK} \\ 3 & \textit{func} \; () \; \textit{// invoke the function} & 3 & \textit{time} \leftarrow \texttt{CONVERT\_TO\_UNITS} \; (\textit{time, units}) \\ 4 & \textit{end} \leftarrow \texttt{GET\_TIME} \; (\texttt{SECONDS}) & 4 & \textbf{return} \; \textit{time} \\ 5 & \textbf{return} \; \textit{end} - \textit{start} \\ \end{array}
```



Handwritten code: client

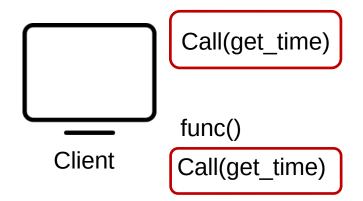
```
procedure MEASURE (func) Local

start ← GET_TIME (SECONDS)

func () // invoke the function

end ← GET_TIME (SECONDS)

return end – start
```





```
Client program RPC
```

```
procedure MEASURE (func)

SEND_MESSAGE (NameForTimeService, {"Get time", convert2external(seconds)})

response ← RECEIVE_MESSAGE (NameForClient)

start ← CONVERT2INTERNAL (response)

func () // invoke the function

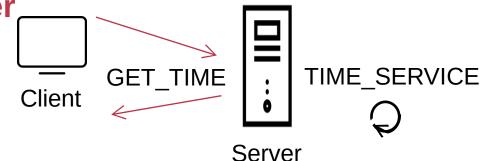
send_Message (NameForTimeService, {"Get time", convert2external(seconds)})

response ← RECEIVE_MESSAGE (NameForClient)

end ← CONVERT2INTERNAL (response)

return end – start
```

Handwritten code: server

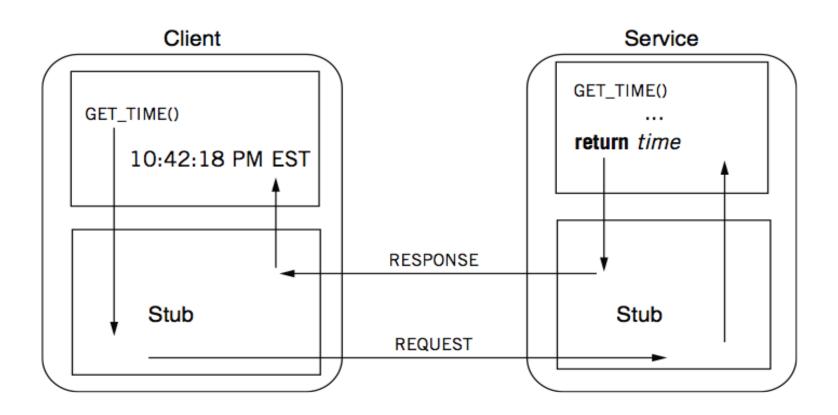


```
10
        procedure TIME_SERVICE ()
11
            do forever
12
                request ← RECEIVE_MESSAGE (NameForTimeService)
13
                opcode \leftarrow GET\_OPCODE (request)
14
                unit \leftarrow convert2internal(get\_argument (request))
15
                if opcode = "Get time" and (unit = SECONDS or unit = MINUTES) then
16
                    time \leftarrow convert_{to_units} (clock, unit)
17
                    response \leftarrow \{"OK", convert2external (time)\}
18
                else
19
                    response ←{"Bad request"}
20
                SEND_MESSAGE (NameForClient, response)
```

RPC simplifies the implementation of remote calls

```
Client program
Abstracts away the common
                                                     SEND_MESSAGE (NameForTimeService, {"Get time", CONVERT2EXTERNAL(SECONDS)})
                                                     response ← RECEIVE MESSAGE (NameForClient)
parts with stub
                                                     start ← CONVERT2INTERNAL (response)
                                                             // invoke the function
                                                     SEND_MESSAGE (NameForTimeService, {"Get time", convert2external(seconds)}
                                                     response ← RECEIVE_MESSAGE (NameForClient)
               Provided in
                                                     end ← CONVERT2INTERNAL (response)
               RPC's stub
                                                     return end – start
10
         procedure TIME SERVICE ()
11
             do forever
12
                  request ← RECEIVE_MESSAGE (NameForTimeService)
13
                  opcode \leftarrow GET\_OPCODE (request)
                  unit ← CONVERT2INTERNAL(GET_ARGUMENT (request))
14
15
                  if opcode = "Get time" and (unit = SECONDS or unit = MINUTES) then
16
                       time \leftarrow convert_{to}units (clock, unit)
17
                       response \leftarrow \{"OK", convert2external (time)\}
18
                  else
19
                       response ←{"Bad request"}
20
                  SEND_MESSAGE (NameForClient, response)
```

RPC: a complete calling process



RPC stub

Client stub

- Put the arguments into a request
- Send the request to the server
- Wait for a response

Stub: hide communication details from up-level code, so that up-level code does not change.

Service stub

- Wait for a message
- Get the parameters from the request
- Call a procedure according to the parameters (e.g. GET_TIME)
- Put the result into a response
- Send the response to the client

Client Program using RPC

```
procedure MEASURE (func)
    start <- GET_TIME(SECONDS)
    func()
    end <- GET_TIME(SECONDS)
    return end - start</pre>
```

 Note: code is not changed comparing with the single-machine case

This is the **stub** of client

```
procedure GET_TIME (units)
    SEND_MESSAGE(ServerName, {"Get time",

CONVERT2EXTERNAL(units)})
    response <- RECEIVE_MESSAGE(ClientName)
    if GET_RETCODE(response) != "OK"
        HANDLE_ERROR(response)
    else
        return CONVERT2INTERNAL(GET_ARGUMENT(response))</pre>
```

Server Program using RPC

procedure GET_TIME (units)

```
time <- CLOCK
    time <- CONVERT_TO_UNITS(time,
units)
                                                This is the stub of server
    return time
procedure TIME_SERVICE ()
    do forever
         request <- RECEIVE_MESSAGE(ServerName)</pre>
        opcode <- GET_OPCODE(request)</pre>
        arg <- CONVERT2INTERNAL(GET ARGUMENT(request))</pre>
         if opcode = "Get time" and (arg = SECONDS or arg = MINUTES)
then
             retval <- GET_TIME(arg)</pre>
             response <- {"OK", CONVERT2EXTERNAL(retval)}
        else
             response <- {"Bad request"}</pre>
        SEND_MESSAGE(ClientName, response)
```

← Note: this code is not changed

Question: what is inside a message?

```
procedure TIME_SERVICE ()
    do forever
        request <- RECEIVE_MESSAGE(ServerName)
        opcode <- GET_OPCODE(request)
        arg <- CONVERT2INTERNAL(GET_ARGUMENT(request))
        if opcode = "Get time" and (arg = SECONDS or arg = MINUTES) then
            retval <- GET_TIME(arg)
            response <- {"OK", CONVERT2EXTERNAL(retval)}
        else
            response <- {"Bad request"}
        SEND_MESSAGE(ClientName, response)</pre>
```

Question: what is inside a message?

A message may contain:

- Service ID (e.g., function ID)
- Service parameter (e.g., function parameter)
- Using marshal / unmarshal

RPC request message

RPC request:

- Xid
- call/reply
- rpc version
- program #
- program version
- procedure #
- auth stuff
- arguments

X is short for "transaction"

Client reply dispatch uses xid

Client remembers the xid of each call



Server dispatch uses prog#, proc#

RPC reply message

RPC reply:

- Xid
- call/reply
- accepted? (Yes, or No due to bad RPC version, auth failure, etc.)
- auth stuff
- success? (Yes, or No due to bad prog/proc #, etc.)
- results

Binding: find the server

Can implement with other network name services

- E.g., <u>192.168.10.233:8888</u> + function ID

Example: gRPC

```
gRPC client
grpc::CreateChannel("localhost:50051",
grpc::InsecureChannelCredentials());

gRPC server

std::stringserver_address("0.0.0.0:50051");
...
... AddListeningPort(server_address, ...);
```

How to pass the data between client & server?

Parameter passing

Pass by value?

Easy: just copy data to network message

Pass by reference?

Makes no sense without shared memory

Needs a conversion between data used in a program vs. data that can be transferred through the network

- Client converts data structure into pointerless representation
- Client transmits data to the server
- Server reconstructs structure with local pointers

Parameter passing is challenging across machines

Distributed systems have the **incompatibility** problem,

which does not exist on a single machine

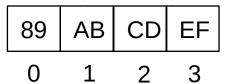
For example, remote machine may have **different**

- byte ordering,
- sizes of integers and other types,
- floating point representations,
- character sets,
- alignment requirements
- etc.

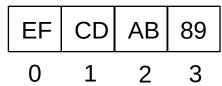
Represent 0x89abcdef



Big endian, e.g., Power processor



Little endian, e.g., X86



Parameter passing is more challenging

Application is changing over time

- What would happen if a server upgrades while the client is intact?
- E.g., server upgrades and adds a priority field to its message, what would happen if the client does not change?

Evolvability: we should built systems that are easy to adapt to changes!

- Backward compatibility: newer code can read data that was written by older code
- Forward compatibility: older code can read that was written by newer code

Representing data: encoding

Need standard **encoding** to enable communication between heterogeneous systems & different versions of software

- Sun's RPC uses XDR (eXternal Data Representation)
- ASN.1 (ISO Abstract Syntax Notation)
- JSON
- Google Protocol Buffers
- W3C XML Schema Language







Why not using language-specific formats?

Many languages have **built-in support** for encoding in-memory objects

- Java's java.io.Serializable
- Python's pickle





Drawbacks:

- The encoding is tied to a particular programming language
 - E.g., it's challenging to use a Java client to call a python server
- No versioning -> no forward and backward compatibility

System requirements for encoding/decoding

Transfer objects through the network

Correctly encode and decode a object to a byte stream

Compatibility

Support multiple language, multiple versions of program

Efficiency

- Reduce the traffic transferred from the network
- Network bandwidth is a scarce resource

Standardized encoding: JSON, XML & etc.

Independent to a specific programming language

Textual formats:

- JSON
- XML
- CSV

Logic client request format

- Xid
- call/reply
- rpc version
- program #
- program version
- procedure #
- auth stuff
- arguments

JSON representation

```
{ "xid" : 12, "call": true,
    "rpc_version": 73,
    ...
}
```

Standardized encoding: JSON, XML & etc.

Independent to a specific programming language

Textual formats:

- JSON
- XML
- CSV

Benefits:

Human-readable: easy to debug

Standardized encoding: JSON, XML & etc

Independent to a specific programming language

Textual formats:

- JSON
- XML
- CSV

Drawbacks:

- 1. Ambiguity around encoding of numbers
- 2. How to support binary strings?
 - Programmers have to re-encode the string as Base64, etc.
- 3. **Verbose**: use more bytes to store the data. E.g., tags like <xx> </xx> in XML

```
How may bytes should I use to store the number?

JSON rep (ation)

{ "xid" : 12, "call": true, "rpc_version": 73, ....
```

Binary formats

Encode the data using binary encoding

- Pros: more compact, more accurate, faster to parse and store
- Cons: less human-readablity
 - Not a problem when the data is only used internally
 - E.g., RPC message metadata

```
Typical example: Thrift & Protocol Buffers
```

Thrift, by Facebook, now in Apache

{ "userName": "Martin", "favoriteNumber": Protocol Buffers, by Google

1337, "interests": ["daydreaming",

"hacking"] }
JSON representation takes 81B in total (w/o spaces)

Binary formats can reduce it to 34B or even less

Binary formats: schema

Both Thrift and Protocol Buffers require a **schema** for any data that is encoded

Benefits: no need to encode things such as userName in the encoded data

```
Thrift interface definition language (IDL)

struct Person {

1: required string userName,

2: optional i64 favoriteNumber,

3: optional list<string> interests

}

Protocol Buffers IDL

message Person {

required string user_name = 1;

optional int64 favorite_number = 2;

repeated string interests = 3;

}
```

```
{ "userName": "Martin", "favoriteNumber":
1337, "interests": ["daydreaming",
JSON representation takes 81B in total (w/o spaces)
```

The BinaryProtocol of Thrift

Each field has:

- Type annotation (1B)
- Type field (1B)
- A length indication (optional, required for string, list, etc.)
- Data (like json)

For each type, the protocol assumes an internal encoding

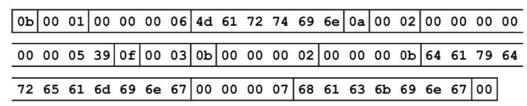
E.g., little endian

Total 59B

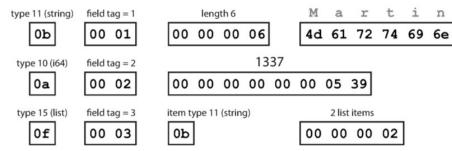
 no field names (userName, favoriteNumber, interests)

Thrift BinaryProtocol

Byte sequence (59 bytes):



Breakdown:



length 11

00 00 00 0ь

length 7

00 00 00 07



end of struct

00


64 61 79 64 72 65 61 6d 69 6e 67

68 61 63 6b 69 6e 67

Being more compact: Thrift CompactProtocol

Techniques:

- Packing field type & field tag in 1B
- 2 Variable-length integer: top-bit of each byte indicates whether there are more bytes



Byte sequence (59 bytes): 0b|00 01|00 00 00 06|4d 61 72 74 69 6e|0a|00 02|00 00 00 00 00 00 05 39 0f 00 03 0b 00 00 00 02 00 00 00 0b 64 61 79 64 72 65 61 6d 69 6e 67 00 00 00 07 68 61 63 6b 69 6e 67 00 (2)Breakdown: type 11 (string) field tag = 1length 6 artin 00 01 00 00 00 06 4d 61 72 74 69 6e 0b 1337 field tag = 2type 10 (i64) 0a 00 02 00 00 00 00 00 00 05 39 type 15 (list) field tag = 3item type 11 (string) 2 list items 0£ 00 03 0b 00 00 00 02 length 11 00 00 00 0ь 64 61 79 64 72 65 61 6d 69 6e 67 hacking length 7 end of struct 68 61 63 6b 69 6e 67 00 00 00 00 07

Being more compact: Thrift CompactProtocol

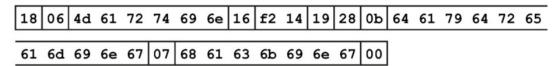
Techniques:

- Packing field type & field tag in 1B
- 2 Variable-length integer: *top-bit of each byte indicates whether there are more bytes*1337: from 8B to 2B

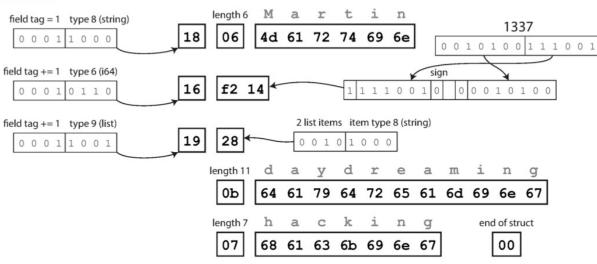
Now only consumes **34B**

Thrift CompactProtocol

Byte sequence (34 bytes):



Breakdown:

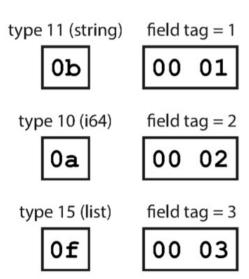


Schema simplifies supporting compatibility

Key observation: we only need to keep the <u>field tags</u> & <u>field type</u> compatible

Forward compatible: if new schema preserves the presentation of old schema, then new code can trivially decode the data encoded with old schema

Backward compatible: old code can simply **skip** fields with **unknown field tag**



Short summary of parameter passing in RPC

Passing data using RPC is non-trivial, should ensure:

- Correctness
- Compact
- Evolvable

Encoding/decoding

- Convert an object into an array of bytes with enough annotation so that the decode procedure can convert it back into an object
- Also known as marshal / unmarshal in RPC
- With other names, e.g., serialize / unserialize or deserialize

Automatic stub generation

Generate stubs from an interface specification

- Tool to look at argument and return types
- Generate the marshal and unmarshal code
- Generate stub procedures

Benefits:

- Saves programming (thus less error)
- Ensures agreement on argument types
 - E.g., consistent function ID

Typically a standard component in an RPC library

Transport protocol of RPC

Stubs also include implementations of sending/receiving messages

- TCP or UDP, which one to use?
- What about new networking features, e.g., RDMA?

Hide the transport protocol from the user

Benefits: e.g., transparent migrate to a more advanced network

Most support several

 Allow programmer (or end user) to choose at runtime based on their hardware setup or performance requirements

Short summary of encoding & decoding data for RPC

Transfer objects through the network

Correctly encode and decode a object to a byte stream

Compatibility

Uniformed format + Schema

Support multiple language, multiple versions of program

Efficiency

Reduce the traffic transferred from the network

Binary format

Network bandwidth is a scarce resource

Short summary of encoding & decoding data for RPC

Efficiency

- Reduce the traffic transferred from the network
- Network bandwidth is a scarce resource

Question

- Does binary format always improve the efficiency of RPC?
- E.g., What if the network is becoming faster?

Short summary: RPC so far

RPC simplifies programming w/ an interface similar to local function call

RPC uses stubs to avoid handling argument **encoding/decoding** and send/receiving messages, etc.

Ensure correctness & efficiency

Remain challenges of RPC:

- How to handle failure?
 - As we have mentioned, failure is common especially in large-scale distributed systems!

Local procedure calls do not fail (most of the time)

If that happens, the entire process dies

RPC may meet many different failures

- Parameter error -> just return
- Server crash -> how the client copes with it?

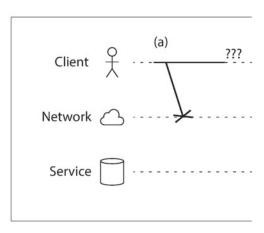
Crash can happen even the remote procedure runs correctly!

RPC could fail in a partial failure way

Like what we have presented in lecture 02


A client uses RPC to call a function at a service through the network

- The client does not get any reply
- What could possibly happen?



A user sends an RPC but the server does not reply, possible reasons:

- 1. Request may have been **lost** (e.g., someone unplugged a network cable).
- 2. Request may be waiting in a queue and will be delivered **later** (e.g., the network or the recipient is overloaded).
- 3. The remote node may have **failed** (e.g., it crashed or it was powered down).



A user sends an RPC but the server does not reply, possible reasons:

- 4. The remote node may **have temporarily stopped** responding (e.g., it is experiencing a long **garbage collection pause**)
- 5. The remote node may **have executed the function**, but the **response** has been **lost** on the network (e.g., a network switch has been misconfigured).
- 6. The remote node may **have executed the function**, but the response has been **delayed** and will be delivered later (e.g., the network or your own machine is overloaded).

RPC != PC

The **transparency** of RPC breaks here:

- Applications should be prepared to deal with RPC failure
- E.g., gRPC will return a status for application to check

Semantics of remote procedure calls

Local procedure call: Exactly once

A remote procedure call may be called:

- 0 time: server crashed or serve process died before executing server code
- 1 time: everything worked well, as expected
- 1 or more: excess latency or lost reply from server and client retransmission
- 0 or 1 time: the function can execute at most once

What is the most desirable RPC semantic for the developers?

What is the desirable RPC semantic for system developers?

RPC semantic

Most RPC systems will offer either:

- At-least-once semantics
- At-most-once semantics

Simple retransmission leads to "at-least-once"

Birrell's RPC semantics (1984):

- server says **OK**: executed once
- server says CRASH: zero or one time

much easier than exactly once, more useful than at-least-once

RPC semantic

Understand application:

- **Idempotent**: may be run any number of times without harm (e.g., i = 1)
- Non-idempotent: those with side-effects (e.g., i++)

When at-least-once is **OK**?

- if no side effects (e.g., read-only operation)
- if app has its own plan for detecting duplication


Cases of idempotence

"Idempotence means that multiple invocations of some work are identical to exactly one invocation."

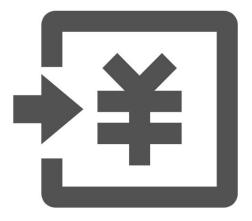
——Pat Helland



Sweeping the floor is idempotent



Pressing the elevator button is idempotent

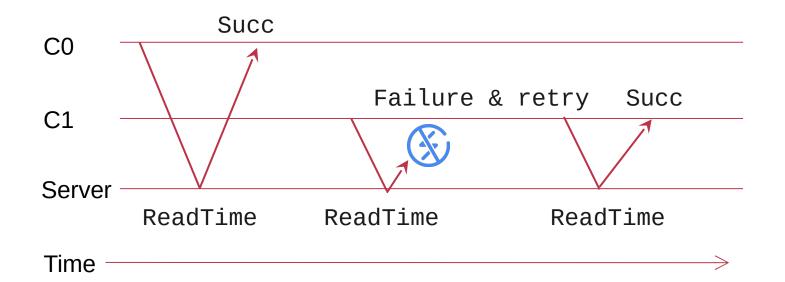


Saving money is nonidempotent

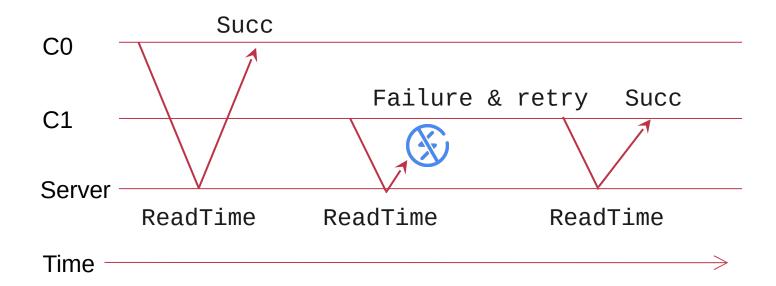
```
RPC client

Call(server, "ReadTime");

ReadTime() {
    return current_time;
}
```




Expect ReadTime is called **twice**Actually called **3 times**Is it ok? Why?



```
RPC client
                                          RPC server
Call(server, "Like");
                                    Like() {
                                       counter++;
                 Succ
    C0
                             Failure & retry
                                                Succ
    Server
            Like
                                             Like
                             Like
            counter=1
                                           counter=3
                            counter=2
    Time
```

Expect Like is called **twice**Actually called **3 times**Is it ok? Why?
What about something like **Vote**?





RPC semantic

Understand application:

- Idempotent: may be run any number of times without harm
- Non-idempotent: those with side-effects (e.g., call i++ at the server)

When at-least-once is OK?

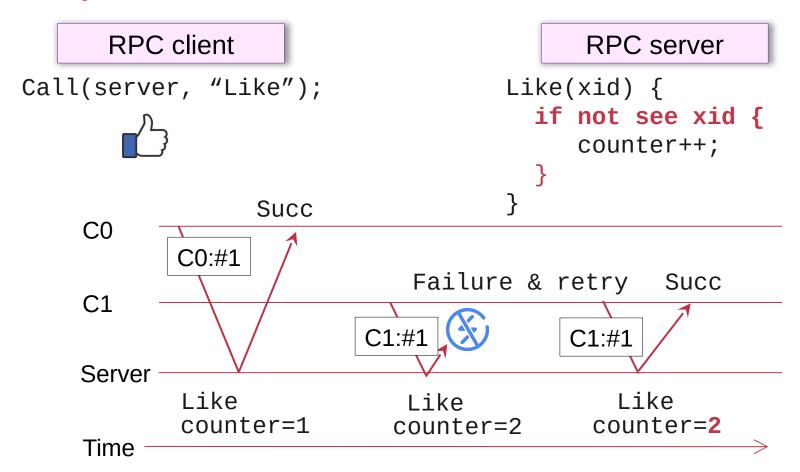
- If no side effect (e.g., read-only operation, ReadTime)
- If app has its own plan for detecting duplication
 - E.g., record how many times a counter has been added

Ideal RPC Semantics: exactly-once

Like single-machine function call

Implement exactly-once semantics:

- Server remembers the requests it has seen and replies to executed RPCs (need to across reboots)
- Detect duplicates, requests need unique IDs (XIDs)



Ideal RPC Semantics: exactly-once

Like single-machine function call

Implement exactly-once semantics:

- Server remember the requests it has seen and replies to executed RPCs (across reboots)
- Detect duplicates, reqs need unique IDs (XIDs)

Assumption: failures are **eventually** repaired, and client retries **forever**

How to correctly recover from failure? See later lectures

Put it all together: RPC system components

- 1. Standards for wire format of RPC message and data types
- 2. Library of routines to marshal / unmarshal data
- 3. Stub generator, or RPC compiler, to produce "stubs"
- For client: marshal arguments, call, wait, unmarshal reply
- For server: unmarshal arguments, call real function, marshal reply

Put it all together: RPC system components

- 4. Server framework:
- Dispatch each call message to correct server stub
- Recall each called functions ,if provide at-most-once semantic or exactly-once semantic
- 5. Client framework:
- Give each reply to correct waiting thread / callback
- Retry if timeout or server cache
- 6. Binding: how does client find the right server?