

# Network Layer

## All about routing

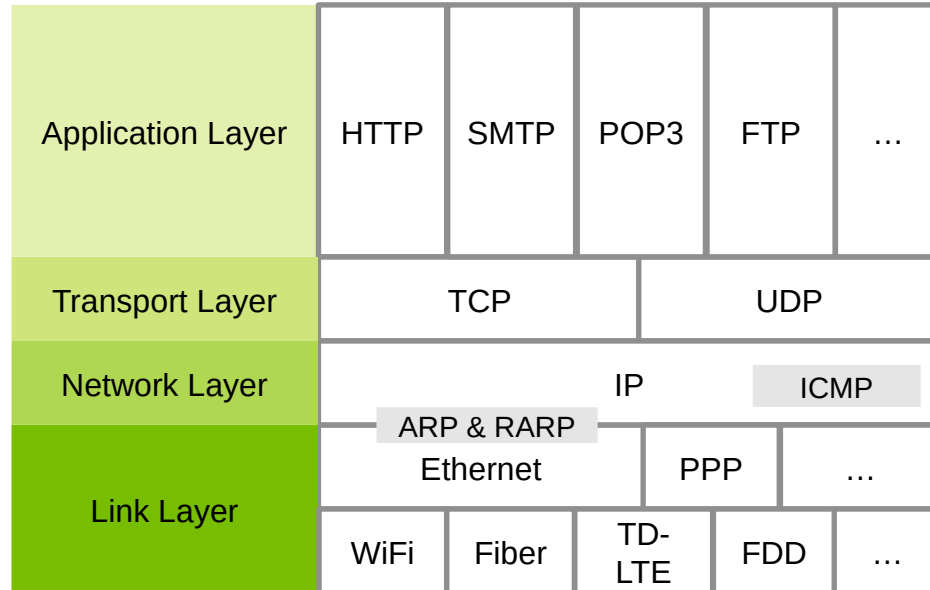
IPADS, Shanghai Jiao Tong University

<https://www.sjtu.edu.cn>

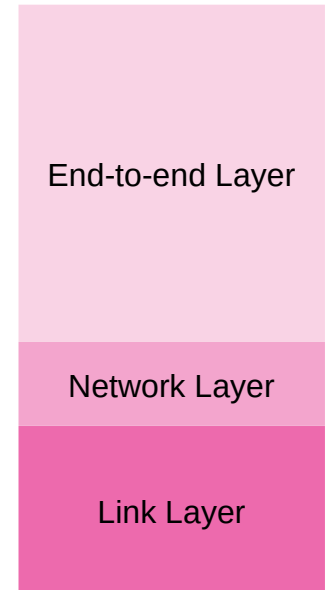
# Review: OSI, TCP/IP & Protocol Stack



OSI



TCP/IP



CSE

# Review: Simple Parity Check

## 2 bits -> 3 bits


- Detect 1-bit errors
- 8 patterns total

## Only 4 correct patterns

- 00 -> 00**0**
- 11 -> 11**0**
- 10 -> 10**1**
- 01 -> 01**1**

## Hamming distance of this code is 2

- 1-bit flipping will cause incorrect pattern





**NETWORK LAYER**

# IP: Best-effort Network

## 1. Best-effort network

- If it cannot dispatch, may discard a packet

## 2. Guaranteed-delivery network

- Also called [store-and-forward](#) network, no discarding data
- Work with complete messages rather than packets
- Use disk for buffering to handle peaks
- Tracks individual message to make sure none are lost


## In real world

- No absolute guarantee
- Guaranteed-delivery: higher layer; best-effort: lower layer

# The Network Layer

## Addressing interface

- Network attachment points
- Network address
- Source & destination



## NETWORK\_SEND

- (segment\_buffer, destination, network\_protocol, end\_layer\_protocol)

## NETWORK\_HANDLE

- (packet, network\_protocol)

# Managing the Forwarding Table: Routing

## Routing (or path-finding)

- Constructing the tables

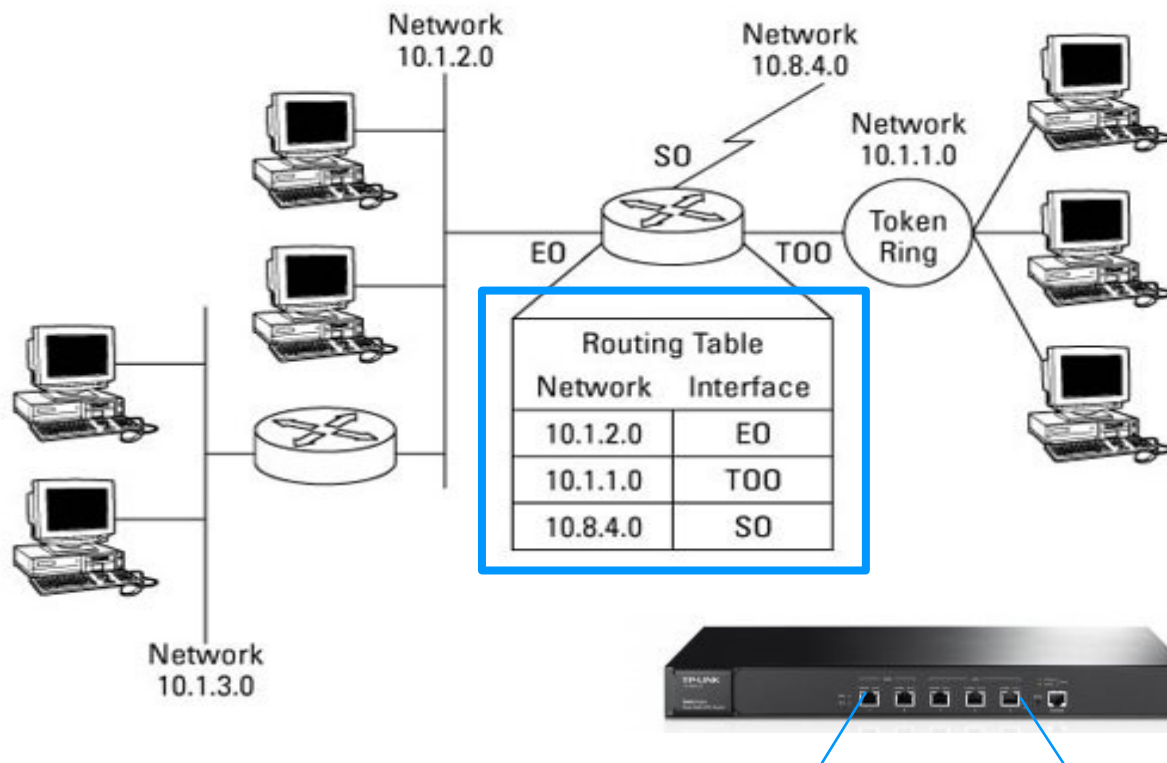
## Impractical by hand

- Determining the best paths requires calculation
- Recalculating the table when links change
- Recalculating the table when link fails
- Adapt according to traffic congestion

## Static routing vs. adaptive routing

- Adaptive routing requires exchange of info

# IP Route Table





# Control-plane VS. Data-plane

## Control-plane

- Control the data flow by defining rules
- E.g., the routing algorithm

## Data-plane

- Copies data according to the rules
- Performance critical
- E.g., the IP forwarding process

# Routing

How to generate the routing table?

# Goal of A Routing Protocol




Allow each switch to know, for every node **dst** in the network, a route to **dst**

Allow each switch to know, for every node **dst** in the network, a minimum-cost route to **dst**

Build a **routing table** at each switch, such that `routing_table[dst]` contains a minimum-cost route to **dst**

# Distributed Routing: 3 Steps in General

1. Nodes learn about their neighbors via the **HELLO protocol**
2. Nodes learn about other reachable nodes via **advertisements**
3. Nodes determine the **minimum-cost** routes (of the routes they know about)



# Two Types of Routing Protocol


## Protocol 1: Link-state

- A node's advertisements contain a list of its neighbors and its link costs to those nodes
- Nodes advertise to everyone their costs to their neighbors
  - via *flooding*
- Integrate using **Dijkstra's shortest path algorithm**

## Protocol 2: Distance-vector


- Nodes advertise to neighbors with their cost to all known nodes
- Update routes via **Bellman-Ford** integration

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



link state

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm




link state

what's in an advertisement

its link costs to each of  
its neighbors

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]


link state

what's in an advertisement

its link costs to each of  
its neighbors



**link-state routing:** disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]


**link state**

what's in an  
advertisement

its link costs to each of  
its neighbors

who gets a  
node's  
advertisement

**link-state routing:** disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

### link state


what's in an  
advertisement

its link costs to each of  
its neighbors

who gets a  
node's  
advertisement

effectively, every other  
node (via flooding)

**link-state routing:** disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

### link state


what's in an  
advertisement

its link costs to each of  
its neighbors

who gets a  
node's  
advertisement

effectively, every other  
node (via flooding)

**link-state routing:** disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

### link state


what's in an  
advertisement

its link costs to each of  
its neighbors

who gets a  
node's  
advertisement

effectively, every other  
node (via flooding)

**link-state routing:** disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

### link state


what's in an  
advertisement

its link costs to each of  
its neighbors

who gets a  
node's  
advertisement

effectively, every other  
node (via flooding)

**link-state routing:** disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

### link state


what's in an  
advertisement

its link costs to each of  
its neighbors

who gets a  
node's  
advertisement

effectively, every other  
node (via flooding)

**link-state routing:** disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

### link state


what's in an  
advertisement

its link costs to each of  
its neighbors

who gets a  
node's  
advertisement

effectively, every other  
node (via flooding)

**link-state routing:** disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

### link state

what's in an  
advertisement


its link costs to each of  
its neighbors

who gets a  
node's  
advertisement

effectively, every other  
node (via flooding)



**link-state routing:** disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

**link state**


what's in an  
advertisement

its **link costs** to each of  
its **neighbors**

who gets a  
node's  
advertisement

effectively, **every other  
node** (via flooding)

**link-state routing:** disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

**link state**


what's in an  
advertisement

its link costs to each of  
its neighbors

who gets a  
node's  
advertisement

effectively, every other  
node (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

nodes keep track of which advertisements they've forwarded so  
that they don't re-forward them

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other**  
**node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's advertisement: [(B,7),(D,2),(F,1)]

nodes keep track of which advertisements they've forwarded so  
that they don't re-forward them

they can also be a bit smarter about flooding, and not forward an  
advertisement back to the node that sent it

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other**  
**node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



nodes *integrate* advertisements by running  
Dijkstra's Algorithm

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other**  
**node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	7
C	?	$\infty$
D	A-D	2
E	?	$\infty$
F	A-F	1

### link state

what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

# Link-state Routing

**Keep track of W, the set of nodes haven't processed yet**

- Initially, W is all nodes in the network

**Keep track of the current costs and routes to all of the nodes. Initially:**

- `routing_table[self] = Self`; `routing_table[anyone else] = ?`
- `cost_table[self] = 0`; `cost_table[anyone else] = infinity`

**While W is not empty:**

- 1.  $u$  = the node in W we have the minimum cost to so far
- 2. Remove  $u$  from W
- 3. For every neighbor  $v$  of  $u$ :


$d = \text{cost\_table}[u] + \text{cost}(u, v)$

if  $d < \text{cost\_table}[v]$

$\text{cost\_table}[v] = d$

$\text{routing\_table}[v] = \text{routing\_table}[u]$

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	7
C	?	$\infty$
D	A-D	2
E	?	$\infty$
F	A-F	1

### link state

what's in an advertisement


its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)



link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	7
C	?	$\infty$
D	A-D	2
E	?	$\infty$
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	7
C	?	$\infty$
D	A-D	2
E	?	$\infty$
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	7
C	?	$\infty$
D	A-D	2
E	?	$\infty$

F | A-F | 1

F does not provide A  
with a better route to D

## link state


what's in an advertisement

its link costs to each of  
its neighbors

who gets a node's advertisement

effectively, **every other**  
**node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	7
C	?	$\infty$
D	A-D	2
E	?	$\infty$
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
-----	-------	------

B	A-B	7
---	-----	---

C	?	$\infty$
---	---	----------

D	A-D	2
---	-----	---

E	A-F	6
---	-----	---

F	A-F	1
---	-----	---

= the cost from A to F + the cost from F to E

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	7
C	?	$\infty$
D	A-D	2
E	A-F	6
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	7
C	?	$\infty$
D	A-D	2
E	A-F	6
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	7
C	?	$\infty$
D	A-D	2
E	A-F	6
F	A-F	1

### link state

what's in an advertisement


its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)



link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	?	$\infty$
D	A-D	2
E	A-F	6
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	?	$\infty$
D	A-D	2
E	A-F	6
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	A-D	7
D	A-D	2
E	A-F	6
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	A-D	7
D	A-D	2
E	A-F	6
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	A-D	7
D	A-D	2
E	A-D	5
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	A-D	7
D	A-D	2
E	A-D	5
F	A-F	1

we don't need to "visit" F; we  
already know the shortest path to it

### link state


what's in an  
advertisement

its link costs to each of  
its neighbors

who gets a  
node's  
advertisement

effectively, every other  
node (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	A-D	7
D	A-D	2
E	A-D	5
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	A-D	7
D	A-D	2
E	A-D	5
F	A-F	1

### link state

what's in an advertisement


its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)



link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
-----	-------	------

B	A-B	3
---	-----	---

C	A-D	6
---	-----	---

D	A-D	2
---	-----	---

E	A-D	5
---	-----	---

F	A-F	1
---	-----	---

notice that A's *route* doesn't change,  
but the *cost* needs to update  
(and the actual path of the packets from A to  
C has changed)

## link state


what's in an  
advertisement

its **link costs** to each of  
its **neighbors**

who gets a  
node's  
advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	A-D	6
D	A-D	2
E	A-D	5
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	A-D	6
D	A-D	2
E	A-D	5
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



A's routing table

dst	route	cost
B	A-B	3
C	A-D	6
D	A-D	2
E	A-D	5
F	A-F	1

### link state


what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

link-state routing: disseminate full topology information  
so that nodes can run a shortest-path algorithm



### link state

what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

what happens when things fail?

flooding makes link-  
state routing very  
resilient to failure

what limits scale?

the **overhead** of  
flooding

# Distance-vector Routing

In link-state, nodes calculate full shortest paths. But actually they only need the route (first-hop) to a destination

Advertisement format: Each node's advertisement is a list of all the nodes it knows about, and its current costs to those nodes

- Initially, this advertisement is just [(self, 0)].

Nodes who receive an advertisement: Node X's advertisement will be received only by its neighbors


Integrate step: When node X receives an advertisement from its neighbor Y, this advertisement will be a list of [(dst, cost)] pairs. Each cost represents Y's cost to dst

## Distance-vector Routing

**For each (dst, cost) in the advertisement, X needs to check for two things:**

- If X is already using Y to get to dst, update the cost information (remember, costs can change!)
- If X is not using Y to get to dst, see if Y could provide a better path; if so, update the routing and cost information

distance-vector routing: disseminate information about the current *costs* to each node, rather than the actual topology



link state

distance vector

what's in an advertisement

its **link costs** to each of  
its **neighbors**

who gets a node's advertisement

effectively, **every other  
node** (via flooding)

what happens when things fail?


flooding makes link-  
state routing very  
resilient to failure

what limits scale?

the **overhead** of  
flooding



distance-vector routing: disseminate information about the current *costs* to each node, rather than the actual topology



link state

distance vector

what's in an advertisement

its **link costs** to each of  
its **neighbors**

its **current costs** to  
**every node** it's aware of

who gets a node's advertisement

effectively, **every other**  
**node** (via flooding)


what happens when things fail?

flooding makes link-  
state routing very  
resilient to failure

what limits scale?

the **overhead** of  
flooding

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's first advertisement: [(B,7), (D,2),(F,1)]

A could also include (A,0) here


A's routing table

dst	route	cost
B	A-B	
7 D	A-D	
F	A-F	1

A's advertisement reflects its routing table, and right now, A only knows about its neighbors

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's first advertisement: [(B,7), (D,2),(F,1)]

A could also include (A,0) here

A's routing table		
dst	route	cost
B	A-B	
7 D	A-D	
F	A-F	1

A's advertisement reflects its routing table, and right now, A only knows about its neighbors

link state

distance vector

what's in an advertisement

its link costs to each of its neighbors

its current costs to every node it's aware of

who gets a node's advertisement

effectively, every other node (via flooding)

only its neighbors


what happens when things fail?

flooding makes link-state routing very resilient to failure

what limits scale?

the overhead of flooding

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's first advertisement: [(B,7), (D,2),(F,1)]

A could also include (A,0) here

A's routing table		
dst	route	cost
B	A-B	
7 D	A-D	
F	A-F	1

A's advertisement reflects its routing table, and right now, A only knows about its neighbors

link state

distance vector

what's in an advertisement

its link costs to each of its neighbors

its current costs to every node it's aware of

who gets a node's advertisement

effectively, every other node (via flooding)

only its neighbors


what happens when things fail?

flooding makes link-state routing very resilient to failure

what limits scale?

the overhead of flooding

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's first advertisement: [(B,7), (D,2),(F,1)]


A could also include (A,0) here

A's routing table		
dst	route	cost
B	A-B	
7 D	A-D	
F	A-F	1

A's advertisement reflects its routing table, and right now, A only knows about its neighbors

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	only its neighbors
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's first advertisement: [(B,7),(D,2),(F,1)]


A's routing table

dst	route	cost
B	A-B	
7 D	A-D	
F	A-F	1

A's neighbors **do not** forward A's advertisements; they *do* send advertisements of their own to A

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	only its neighbors
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's first advertisement: [(B,7),(D,2),(F,1)]

A's routing table

dst	route	cost
B	A-B	
7 D	A-D	
F	A-F	1

A's neighbors **do not** forward A's advertisements; they *do* send advertisements of their own to A

link state

distance vector

what's in an advertisement

its link costs to each of its neighbors

its current costs to every node it's aware of

who gets a node's advertisement

effectively, every other node (via flooding)

only its neighbors


what happens when things fail?

flooding makes link-state routing very resilient to failure

what limits scale?

the overhead of flooding

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's first advertisement: [(B,7),(D,2),(F,1)]

A's routing table


dst	route	cost
B	A-B	
7 D	A-D	
F	A-F	1

A's neighbors **do not** forward A's advertisements; they *do* send advertisements of their own to A

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	only its neighbors
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	



distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's first advertisement: [(B,7),(D,2),(F,1)]

A's routing table

dst	route	cost
B	A-B	
7 D	A-D	
F	A-F	1

A's neighbors **do not** forward A's advertisements; they *do* send advertisements of their own to A

link state

distance vector

what's in an advertisement

its link costs to each of its neighbors

its current costs to every node it's aware of

who gets a node's advertisement

effectively, every other node (via flooding)

only its neighbors


what happens when things fail?

flooding makes link-state routing very resilient to failure

what limits scale?

the overhead of flooding

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's routing table


dst	route	cost
B	A-B	
7 D	A-D	
F	A-F	1

B's first adv: [(A,7), (C,3), (D,1)]  
 D's first adv: [(B,1), (C,5), (E,3), (F,4)(A,2)]  
 F's first adv: [(A,1), (D,4), (E,5)]

A receives advertisements from B, D, and F

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	only its neighbors
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's routing table

dst	route	cost
B	A-B	
7 D	A-D	
F	A-F	1

B's first adv: [(A,7), (C,3), (D,1)]

link state

distance vector

what's in an advertisement

its link costs to each of its neighbors

its current costs to every node it's aware of

who gets a node's advertisement

effectively, every other node (via flooding)

only its neighbors


what happens when things fail?

flooding makes link-state routing very resilient to failure

what limits scale?

the overhead of flooding

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's routing table


dst	route	cost
B	A-B	7
C	A-B	10
D	A-D	2
F	A-F	1

B's first adv: [(A,7), (C,3), (D,1)]

A's cost to B + B's cost to C

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	only its neighbors
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology




A's routing table

dst	route	cost
B	A-B	7
C	A-B	10
D	A-D	2
F	A-F	1

B's first adv: [(A,7), (C,3), (D,1)]  
D's first adv: [(B,1), (C,5), (E,3), (F,4)(A,2)]  
F's first adv: [(A,1), (D,4), (E,5)]

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	only its neighbors
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's routing table

dst	route	cost
B	A-B	7
C	A-B	10
D	A-D	2
F	A-F	1

D's first adv: [(B,1), (C,5), (E,3), (F,4)(A,2)]

link state

distance vector

what's in an advertisement

its link costs to each of its neighbors

its current costs to every node it's aware of

who gets a node's advertisement

effectively, every other node (via flooding)

only its neighbors


what happens when things fail?

flooding makes link-state routing very resilient to failure

what limits scale?

the overhead of flooding

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's routing table

dst	route	cost
B	A-D	3
C	A-B	10
D	A-D	2
F	A-F	1

D's first adv: [(B,1), (C,5), (E,3), (F,4)(A,2)]

link state

distance vector

what's in an advertisement

its link costs to each of its neighbors

its current costs to every node it's aware of

who gets a node's advertisement

effectively, every other node (via flooding)

only its neighbors


what happens when things fail?

flooding makes link-state routing very resilient to failure

what limits scale?

the overhead of flooding

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's routing table

dst	route	cost
B	A-D	3
C	A-D	7
D	A-D	2
F	A-F	1

D's first adv: [(B,1), (C,5), (E,3), (F,4)(A,2)]

link state

distance vector

what's in an advertisement

its link costs to each of its neighbors

its current costs to every node it's aware of

who gets a node's advertisement

effectively, every other node (via flooding)

only its neighbors

what happens when things fail?


flooding makes link-state routing very resilient to failure

what limits scale?

the overhead of flooding



distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's routing table

dst	route	cost
B	A-D	3
C	A-D	7
D	A-D	2
E	A-D	5
F	A-F	1

D's first adv: [(B,1), (C,5), (E,3), (F,4)(A,2)]

link state

distance vector

what's in an advertisement

its link costs to each of its neighbors

its current costs to every node it's aware of

who gets a node's advertisement

effectively, every other node (via flooding)

only its neighbors


what happens when things fail?

flooding makes link-state routing very resilient to failure

what limits scale?

the overhead of flooding

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's routing table

dst	route	cost
B	A-D	3
C	A-D	7
D	A-D	2
E	A-D	5
F	A-F	1

B's first adv: [(A,7), (C,3), (D,1)]  
D's first adv: [(B,1), (C,5), (E,3), (F,4)(A,2)]  
F's first adv: [(A,1), (D,4), (E,5)]

link state

distance vector

what's in an advertisement

its link costs to each of its neighbors

its current costs to every node it's aware of

who gets a node's advertisement

effectively, every other node (via flooding)

only its neighbors


what happens when things fail?

flooding makes link-state routing very resilient to failure

what limits scale?

the overhead of flooding

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology




A's routing table

dst	route	cost
B	A-D	3
C	A-D	7
D	A-D	2
E	A-D	5
F	A-F	1

F's first adv: [(A,1), (D,4), (E,5)]

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	only its neighbors
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology




A's routing table

dst	route	cost
B	A-D	3
C	A-D	7
D	A-D	2
E	A-D	5
F	A-F	1

this is A's routing table after one round of advertisements; note that it does not have knowledge of the min-cost route to C yet

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	only its neighbors
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology




A's routing table

dst	route	cost
B	A-D	3
C	A-D	7
D	A-D	2
E	A-D	5
F	A-F	1

A's second adv:  
[(B,3), (C,7), (D,2), (E,5), (F,1)]

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	only its neighbors
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



A's routing table


dst	route	cost
B	A-D	3
C	A-D	7
D	A-D	2
E	A-D	5
F	A-F	1

A's second adv:  
[(B,3), (C,7), (D,2), (E,5), (F,1)]

A will learn about the correct min-cost path to C in the next round of advertisements; try that out for yourself!

link state	distance vector
what's in an advertisement	
its link costs to each of its neighbors	its current costs to every node it's aware of
who gets a node's advertisement	
effectively, every other node (via flooding)	only its neighbors
what happens when things fail?	
flooding makes link-state routing very resilient to failure	
what limits scale?	
the overhead of flooding	

distance-vector routing: disseminate information about the current costs to each node, rather than the actual topology



### link state

### distance vector

what's in an advertisement

its **link costs** to each of its **neighbors**

its **current costs** to every node it's aware of

who gets a node's advertisement

effectively, **every other node** (via flooding)

only its **neighbors**

what happens when things fail?

flooding makes link-state routing very resilient to failure

failures can be complicated because of timing

what limits scale?

the **overhead** of flooding

failure handling

## Problem of INFINITY

**When a node A has no route to destination B, it will advertise a cost of INFINITY to B**

- A cost of INFINITY B is interpreted as there being no route to B
- So INFINITY must be larger than the longest path in the network

**But because the order in which advertisements are sent matters, sometimes nodes can incorrectly think there's a route when there isn't one**

**This can last for up to INFINITY steps (usually 16 or 32)**



# INFINITY

A sends advertisements at  $t=0, 10, 20, \dots$ ; B sends advertisements at  $t=5, 15, 25, \dots$



A: Self, 0

A: B→A, 1

B: A→B, 1

B: Self, 0


C: A→B, 2

C: B→C, 1

$t=9$ : B↔C fails

# INFINITY

A sends advertisements at  $t=0, 10, 20, \dots$ ; B sends advertisements at  $t=5, 15, 25, \dots$



A: Self, 0  
B: A  $\rightarrow$  B, 1  
C: A  $\rightarrow$  B, 2

A: B  $\rightarrow$  A, 1  
B: Self, 0  
C: None, inf

A: Self, 0  
B: A  $\rightarrow$  B, 1  
C: A  $\rightarrow$  B, 2

A: B  $\rightarrow$  A, 1  
B: Self, 0  
C: B  $\rightarrow$  A, 3 (2+1)

A: Self, 0  
B: A  $\rightarrow$  B, 1  
C: A  $\rightarrow$  B, 4

A: B  $\rightarrow$  A, 1  
B: Self, 0  
C: B  $\rightarrow$  A, 3

A: Self, 0  
B: A  $\rightarrow$  B, 1  
C: A  $\rightarrow$  B, 4

A: B  $\rightarrow$  A, 1  
B: Self, 0  
C: B  $\rightarrow$  A, 5



$t=9$ : B  $\leftrightarrow$  C fails

$t=10$ : B receives the following advertisement from A:  
[(A, 0), (B, 1), (C, 2)]


$t=15$ : A receives the following advertisement from B:  
[(A, 1), (B, 0), (C, 3)]

$t=20$ : B receives the following advertisement from A:  
[(A, 0), (B, 1), (C, 4)]

**continues until both costs to C are INFINITY**

# Split Horizon

A sends advertisements at  $t=0, 10, 20, \dots$ ; B sends advertisements at  $t=5, 15, 25, \dots$



A: Self, 0	A: B- $\rightarrow$ A, 1
B: A- $\rightarrow$ B, 1	B: Self, 0
C: A- $\rightarrow$ B, 2	C: None, inf

A: Self, 0	A: B- $\rightarrow$ A, 1
B: A- $\rightarrow$ B, 1	B: Self, 0
C: A- $\rightarrow$ B, 2	C: None, inf

A: Self, 0	A: B- $\rightarrow$ A, 1
B: A- $\rightarrow$ B, 1	B: Self, 0
C: None, inf	C: None, inf



$t=9$ : B $\leftrightarrow$ C fails

$t=10$ : B receives the following advertisement from A:  
[(A,0)]

$t=15$ : A receives the following advertisement from B:  
[(B,0), (C,inf)]

**split horizon takes care of this particular case**

# Link-State Summary

## Pros:

- Fast convergence

## Cons:

- flooding is costly:  $2 \times \text{\#Node} \times \text{\#Line}$  advertisements

**Only good for small networks**

# Distance Vector

## Pros:

- Low overhead: 2x #Line advertisements

## Cons:

- Convergence time is proportional to longest path
- The infinity problem

**Only good for small networks**

# **How to Scale the Routing?**

# 3 Ways to Scale

## Path-vector Routing



- Advertisements include the path, to better detect routing loops

## Hierarchy of Routing

- Route between **regions**, and then within a region

## Topological Addressing

- Assign addresses in contiguous blocks to make advertisements smaller



18.0.0.0, ... , 18.0.0.255  
↓  
18.0.0.0/24

# Path Vector Exchange

## Each participant maintains a path vector

- A complete path to some destination
- E.g., zero-length path to itself
- Gradually learns about other paths
- Construct a new forwarding table from its new path vector


## Algorithm

- Advertising
- Path selection

to	path
G	< >



# Path Vector Exchange



Need coordination to ensure no loop

A	
destination	link
A	end-layer
all other	1

G	
destination	link
A	1
B	2
C	2
D	3
E	4
F	4
G	end-layer
H	2
J	3
K	4

From A,  
via link 1

to	path
A	< >

From H,  
via link 2:

to	path
H	< >

From J,  
via link 3:

to	path
J	< >

From K,  
via link 4:

to	path
K	< >

path vector

to	path
A	<A>
G	< >
H	<H>
J	<J>
K	<K>

forwarding table

to	link
A	1
G	end-layer
H	2
J	3
K	4

path vector

to	path
A	<A>
B	<H, B>
C	<H, C>
D	<J, D>
E	<J, E>
F	<K, F>
G	< >
H	<H>
J	<J>
K	<K>

forwarding table

to	link
A	1
B	2
C	2
D	3
E	3
F	4
G	end-layer
H	2
J	3
K	4

From A,  
via link 1

to	path
A	< >
G	<G>

From H,  
via link 2:

to	path
B	<B>
C	<C>
G	<G>
J	< >
K	<J>
K	<K>

From J,  
via link 3:

to	path
D	<D>
E	<E>
G	<G>
H	<H>
J	< >
K	<K>

From K,  
via link 4:

to	path
E	<E>
F	<F>
G	<G>
H	<H>
J	<J>
K	< >

# Path Vector Exchange

## Better for scaling

- Like Distance-Vector, but include the **full path** in the routing advertisements
- Overhead increases (advertisements are larger), but convergence time decreases (avoid counting to infinity)
- Overhead is still lower than Link-State

## Questions on Path Vector

**How do we avoid permanent loops?**

- When a node updates its paths, it never accepts a path that has itself

**What happens when a node hears multiple paths to the same destination?**

- It picks the better path

**What happens if the graph changes?**

- Algorithm deals well with new links
- To deal with links that go down, each router should discard any path that a neighbor stops advertising

# Hierarchical Address Assignment & Routing


## Two problems of the path vector implementation

- Every attachment point must have a unique address
- The path vector grows in size with the number of attachment points

## Hierarchy for better scalability

- Two parts of network address: region & station, e.g., "11,75"
- Regions correspond to the set of closely-connected entities
- E.g., region-11 has **only 1 entry** in other region routers' table

# Hierarchical Address Assignment & Routing



forwarding table in R1.B

region forwarding section		local forwarding section	
to	link	to	link
R1	local	R1.A	1
R2	1	R1.B	end-layer
R3	1	R1.C	2
R4	3	R1.D	3

Region is also as known as **AS**: Autonomous System

# Hierarchical Address Assignment & Routing

## Problems introduced by hierarchy: more complex

- Binding address with location
  - Has to change address after changing location
- Paths may no longer be the shortest possible
  - Algorithm has less detailed information


## More about hierarchy

- Can extend to more levels
- Different places can have different levels

# Routing Hierarchy

## Across Region

- Use one routing protocol to route **across regions**, and a different protocol to route **within regions**
- Implies that there are devices on the edge of each region that can "translate" between or "speak" both protocols



**BGP** is the path-vector protocol used across regions

- Border Gateway Protocol



# Topological Addressing

## Further reduce the routing table

- Despite being between regions, BGP still routes to IP addresses (e.g., to *18.0.0.1*, not to *region-3* )
- Addresses are given to regions in **contiguous blocks**, so that they can be specified succinctly via a particular notation ("CIDR" notation)
  - CIDR: Classless Inter Domain Routing
- Keeps advertisements small

**18.0.0.0, ... ,18.0.0.255**



**18.0.0.0/24**

CIDR Notation

# **Data Plane: Packet Forwarding**

# Network Layer Interface

```
structure packet
```

```
    bit_string source
```

```
    bit_string destination
```

```
    bit_string end_protocol
```

```
    bit_string payload
```

```
1 procedure NETWORK_SEND (segment_buffer, destination,
2                           net_protocol, end_protocol)
3   packet instance outgoing_packet
4   outgoing_packet.payload ← segment_buffer
5   outgoing_packet.end_protocol ← end_protocol
6   outgoing_packet.source ← MY_NETWORK_ADDRESS
7   outgoing_packet.destination ← destination
8   NETWORK_HANDLE (outgoing_packet, net_protocol)

9 procedure NETWORK_HANDLE (net_packet, net_protocol)
10  packet instance net_packet
11  if net_packet.destination != MY_NETWORK_ADDRESS then
12    next_hop ← LOOKUP (net_packet.destination,
13                      forwarding_table)
14    LINK_SEND (net_packet, next_hop, link_protocol,
15              net_protocol)
16  else
17    GIVE_TO_END_LAYER (net_packet.payload,
```



# Forwarding an IP Packet

## Lookup packet's destination in forwarding table

- If known, find the corresponding outgoing link
- If unknown, drop packet

## Decrement TTL (Time To Live)

- Drop packet if TTL is zero

## Update header checksum

## Forward packet to outgoing port

## Transmit packet onto link

# Data-plane Case Study: Intel's DPDK

## DPDK: Data Plane Development

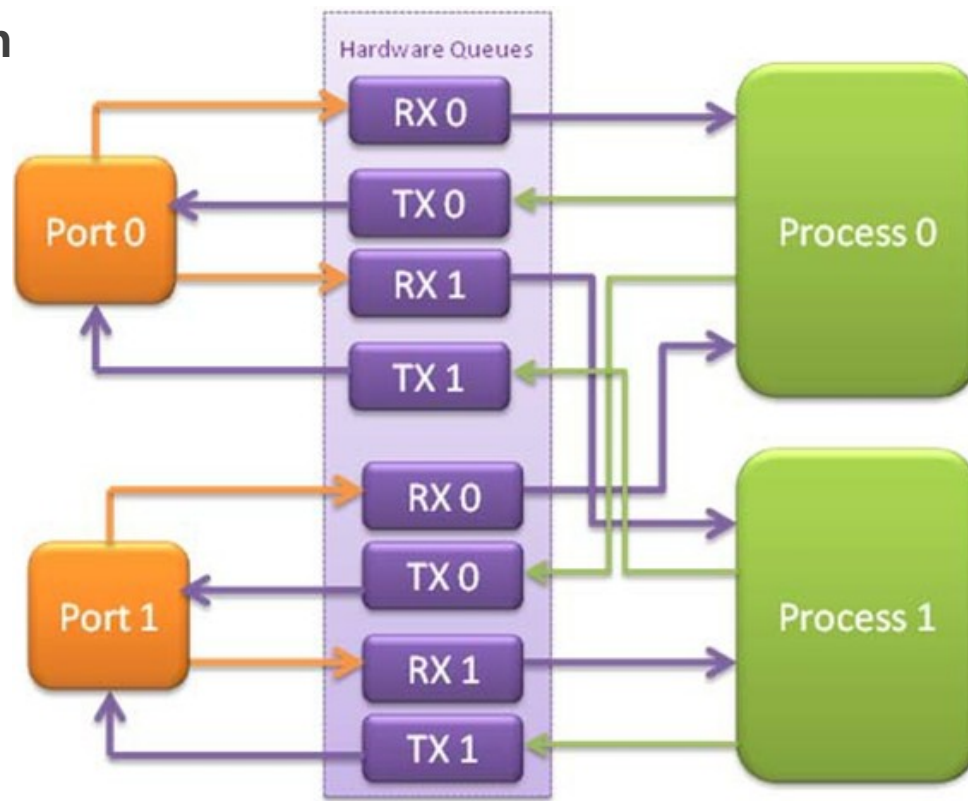
- Bypass kernel

## Network card

- Has several ports
- A port has RX/TX

## Processor

- Read packets from RX
  - Polling
- Find output port
- Write packets to TX



RX for receiving  
TX for sending

**NAT**



# NAT (Network Address Translation)

## Private network

- Public routers do not accept routes to network 10 (e.g., 10.8.8.8)


## NAT router: bridge the private networks

- Router between private & public network
- Send: modify source address to temp public address
- Receive: modify back by looking mapping table

## Limitations

- Some end-to-end protocols place address in payloads
- The translator may become the bottleneck
- What if two private network merge?

# NAT



PAT  
Table

Src PC	Src IP Address	Src. Port # (Many Devices also Translate the SRC Port Number to a New Value)	Translated Src. IP Address	Destination IP Address	Destination Port Number	Translated Destination IP Address
A	10.1.1.1	3759	200.1.1.1	100.1.1.1	80	100.1.1.1
B	10.1.1.2	2119	200.1.1.1	100.1.1.1	80	100.1.1.1



# **CASE: Ethernet Mapping**

Mapping Internet to Ethernet

# Case Study: Mapping Internet to Ethernet

**Listen-before-sending rule, collision**

**Ethernet: CSMA/CD**

- Carrier Sense Multiple Access with Collision Detection

**Ethernet type**

- Experimental Ethernet, 3 mpbs
- Standard Ethernet, 10 mbps
- Fast Ethernet, 100 mbps
- Gigabit Ethernet, 1000 mbps

# Overview of Ethernet

## A half duplex Ethernet

- The max propagation time is less than the 576 bit times, the shortest allowable packet
- So that two parties can detect a collision together
- If collision: wait random first time, exponential backoff if repeat

## A full duplex & point-to-point Ethernet

- No collisions & the max length of the link is determined by the physical medium

leader	destination	source	type	data	checksum
64 bits	48 bits	48 bits	16 bits	368 to 12,000 bits	32 bits

# Difference between Hub and Switch


## Hub

- A frame is "broadcast" to every one of its ports
- A 10/100Mbps hub must share its bandwidth with each port

## Switch

- Keeps a record of the MAC addresses of all the devices
- A 10/100Mbps switch will allocate a full 10/100Mbps to each of its ports

# Broadcast Aspects of Ethernet



## Broadcast network

- Every frame is delivered to every station
- (Compare with forwarding network)

## ETHERNET\_SEND

- Pass the call along to the link layer

## ETHERNET\_HANDLE

- Simple, can even be implemented in hardware



## Broadcast Aspects of Ethernet

```
procedure ETHERNET_HANDLE (net_packet, length)
  destination ← net_packet.target_id
  if destination = my_station_id

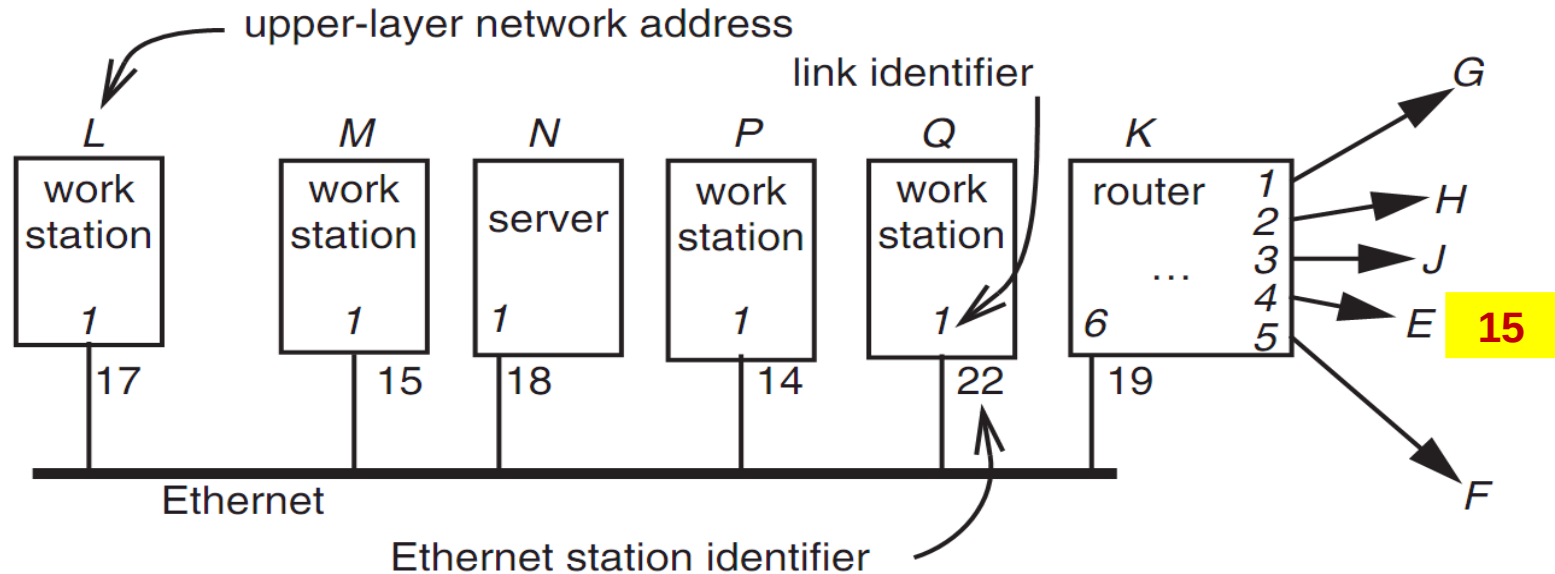
    then
      GIVE_TO_END_LAYER (net_packet.data,
                        net_packet.end_protocol,
                        net_packet.source_id)
    else
      ignore packet
```

**no need to do any  
forwarding**

## Broadcast Aspects of Ethernet

```
procedure ETHERNET_HANDLE (net_packet, length)
  destination ← net_packet.target_id
  if destination = my_station_id
    or destination = BROADCAST_ID
  then
    GIVE_TO_END_LAYER (net_packet.data,
                      net_packet.end_protocol,
                      net_packet.source_id)
  else
    ignore packet
```

# Layer Mapping: Attach Ethernet to Forwarding Network



L sends a RPC to N by sending to station 18 of link 1

L sends a RPC to E by sending to K, E may have 15 as address, as well as M

# Layer Mapping

## The Internet network layer

- **NETWORK\_SEND** (data, length, RPC, INTERNET, N)
- **NETWORK\_SEND** (data, length, RPC, ENET, 18)

L must maintain a table

internet address	Ethernet/ station
<i>M</i>	enet/15
<i>N</i>	enet/18
<i>P</i>	enet/14
<i>Q</i>	enet/22
<i>K</i>	enet/19
<i>E</i>	enet/19

## ARP (Address Resolution Protocol)

NETWORK\_SEND ("where is M?", 11, ARP, ENET, BROADCAST)

NETWORK\_SEND ("M is at station 15", 18, ARP, ENET, BROADCAST)

L asks E's Ethernet address, E does not hear the Ethernet broadcast, but the router at station 19 does, and it sends a suitable ARP response instead

Manage forwarding table as a cache

internet address	Ethernet/ station
<i>M</i>	enet/15

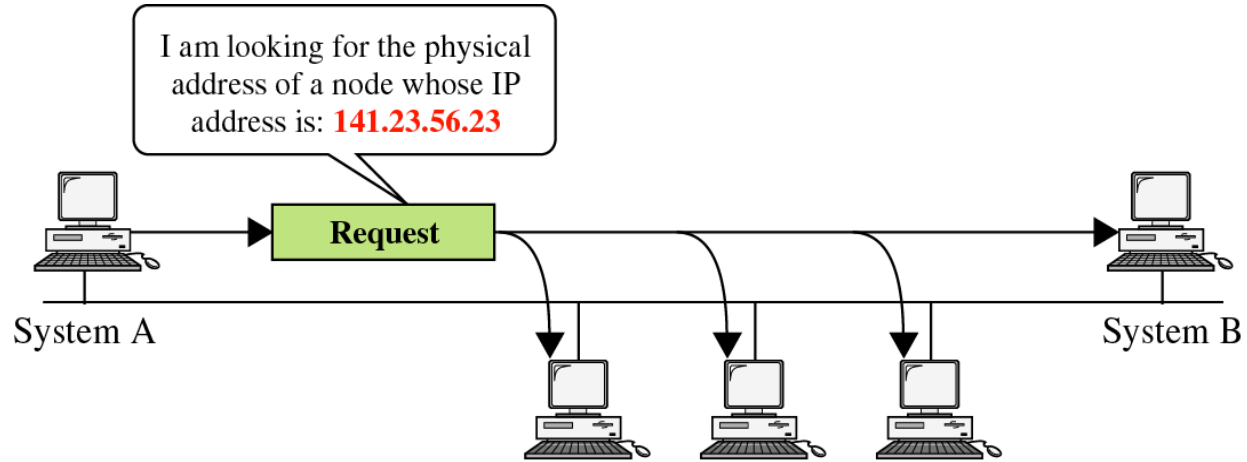
internet address	Ethernet/ station
<i>M</i>	enet/15
<i>E</i>	enet/19

## ARP & RARP Protocol

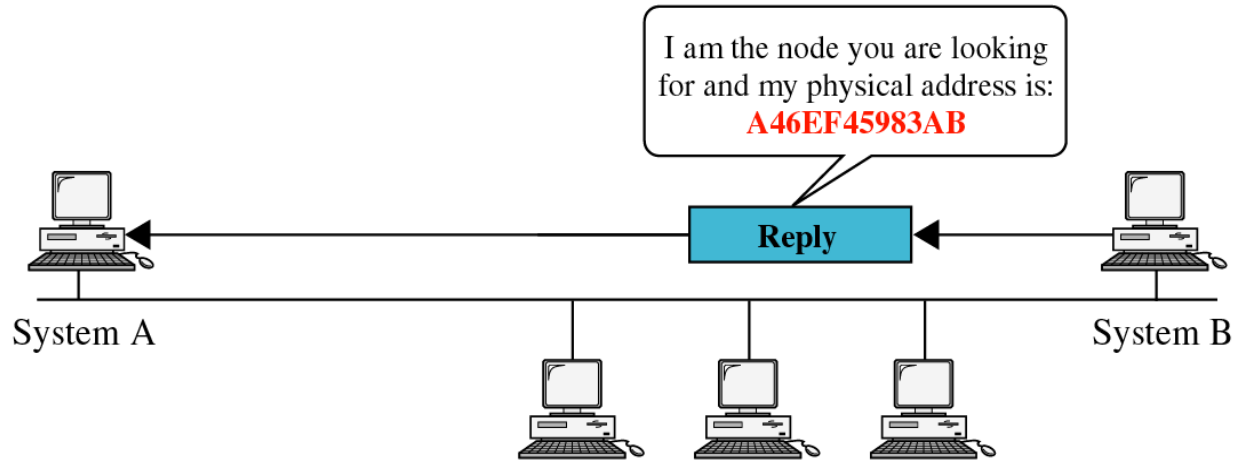
Hardware Type (16 bits)		Protocol Type (16 bits)
HA Length (8 bits)	PA Length (8 bits)	Operation (16 bits)
Sender Hardware Address (Octets 0-3)		
Sender Hardware Address (Octets 4-5)		Sender Protocol Address (Octets 0-1)
Sender Protocol Address (Octets 2-3)		Target Hardware Address (Octets 0-1)
Target Hardware Address (Octets 2-5)		
Target Protocol Address (Octets 0-3)		

Name mapping: IP address <-> MAC address

# ARP & RARP



a. ARP request is broadcast



b. ARP reply is unicast

# Network Topology

## Take SJTU network for example

- Subnet: usually like 192.168.0.2 or 10.0.0.2
- Gateway: usually like 192.168.0.1
  - Get the global IP address: 202.120.40.82
  - A gateway usually has two (or more) IP address
- Proxy: get proxy's address
  - E.g., 106.185.46.164 (Japan)





# Network Topology

## How to Use socket to Access [www.baidu.com](http://www.baidu.com) ?

**You code as if your PC connect directly with Baidu**

- Call `connect()` with Baidu's IP address

**But how does the system find next hop?**

## Putting All Together

**App: I want to send a packet to Baidu, here is the packet with Baidu's IP in its header as target IP, and client's IP as source IP (Node-C)**

**OS: I don't know how to get to Baidu, I'll just send it to the router (gateway). But I cannot change the source IP of the packet, so I'll just change the MAC target address of the packet to the router's MAC address**

## Putting All Together

**The router-1 (gateway):** I get a packet with my MAC as target address. Is it my IP? No... So I'll just forward it to next hop, by changing the target MAC address to next hop's MAC address (NAT: change source IP and source port as well)

**Router-2:** I connect directly to Baidu, I'll just change the target MAC address to Baidu

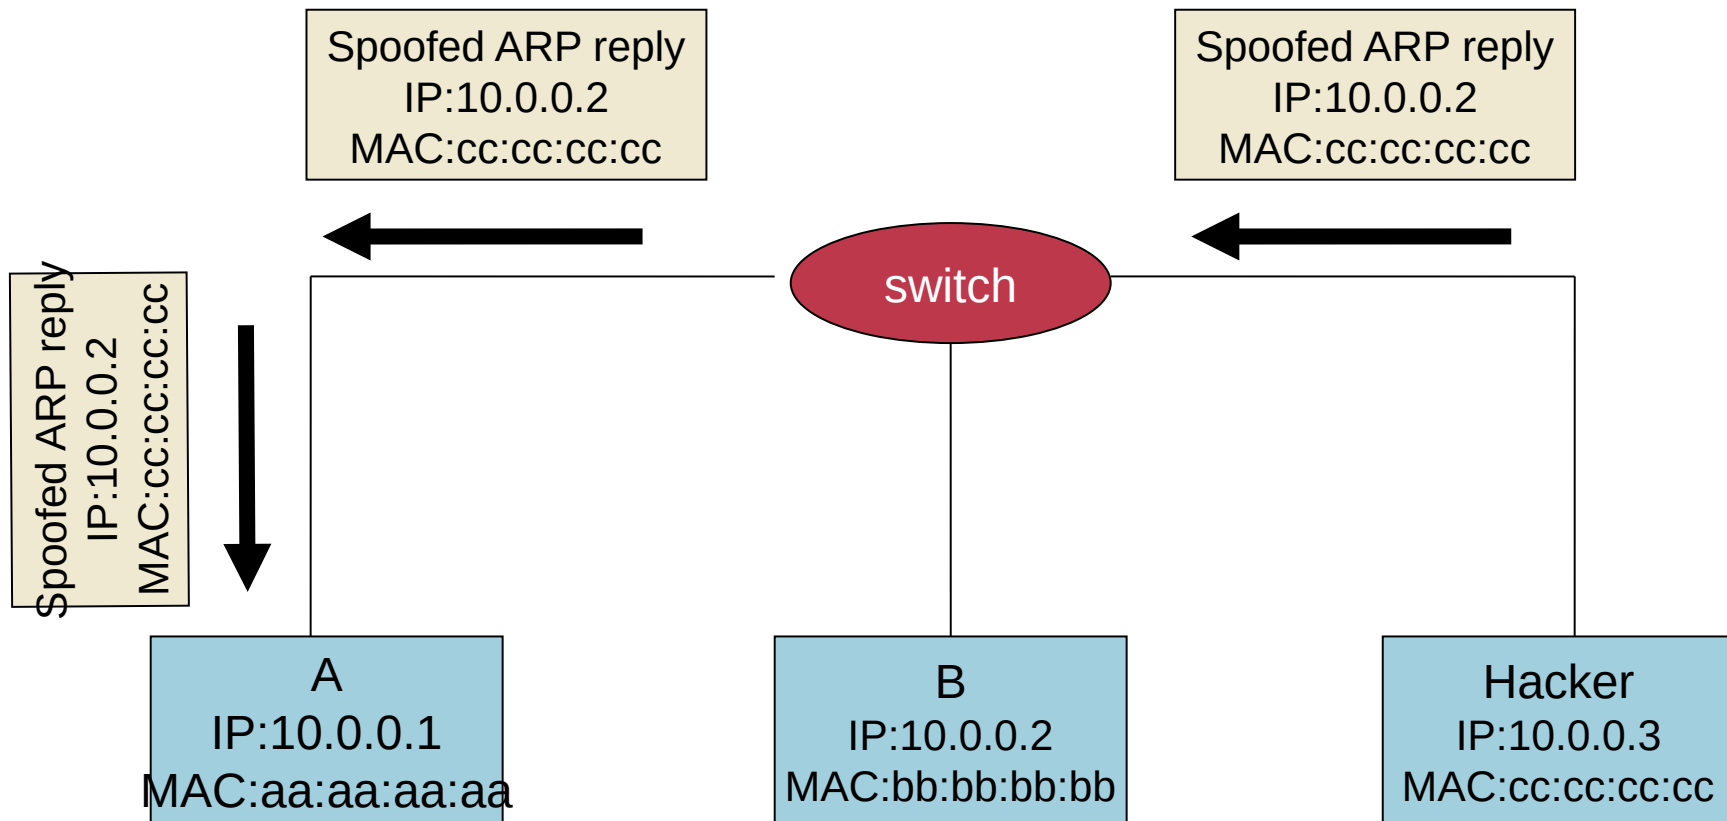
# ARP Spoofing

Construct spoofed ARP replies

A target computer could be convinced to send frames destined for computer A to instead go to computer B

Computer A will have no idea that this redirection took place

This process of updating a target computer's ARP cache is referred to as "**ARP poisoning**"




ARP cache

IP	MAC
10.0.0.2	bb:bb:bb:bb

ARP cache

IP	MAC
10.0.0.1	aa:aa:aa:aa



**A's cache is poisoned**

ARP cache

IP	MAC
10.0.0.2	cc:cc:cc:cc

ARP cache

IP	MAC
10.0.0.1	aa:aa:aa:aa

## ARP Spoofing

Now all the packets that A intends to send to B will go to the hacker's machine

Cache entry would expire, so it needs to be updated by sending the ARP reply again

- How often?
- depends on the particular system
- Usually every 40s should be sufficient

In addition the hacker may not want his Ethernet driver talk too much

- Accomplish with `ifconfig -arp`



## **Man-in-the-Middle Attack**

**A hacker inserts his computer between the communications path of two target computers**

**The hacker will forward frames between the two target computers so communications are not interrupted**

**E.g., Hunt, Ettercap, etc.**

- Can be obtained easily in many web archives

# Man-in-the-Middle Attack

**The attack is performed as follows:**

- Suppose X is the hacker's computer
- T1 and T2 are the targets
- 1. X poisons the ARP cache of T1 and T2
- 2. T1 associates T2's IP with X's MAC
- 3. T2 associates T1's IP with X's MAC
- 4. All of T1 and T2's traffic will then go to X first, instead of directly to each other

Spoofed ARP reply  
IP:10.0.0.2  
MAC:cc:cc:cc:cc

Spoofed ARP reply  
IP:10.0.0.2  
MAC:cc:cc:cc:cc

Spoofed ARP reply  
IP:10.0.0.2  
MAC:cc:cc:cc:cc

switch

T1  
IP:10.0.0.1  
MAC:aa:aa:aa:aa

T2  
IP:10.0.0.2  
MAC:bb:bb:bb:bb


Hacker  
IP:10.0.0.3  
MAC:cc:cc:cc:cc

ARP cache

IP	MAC
10.0.0.2	bb:bb:bb:bb

ARP cache

IP	MAC
10.0.0.1	aa:aa:aa:aa



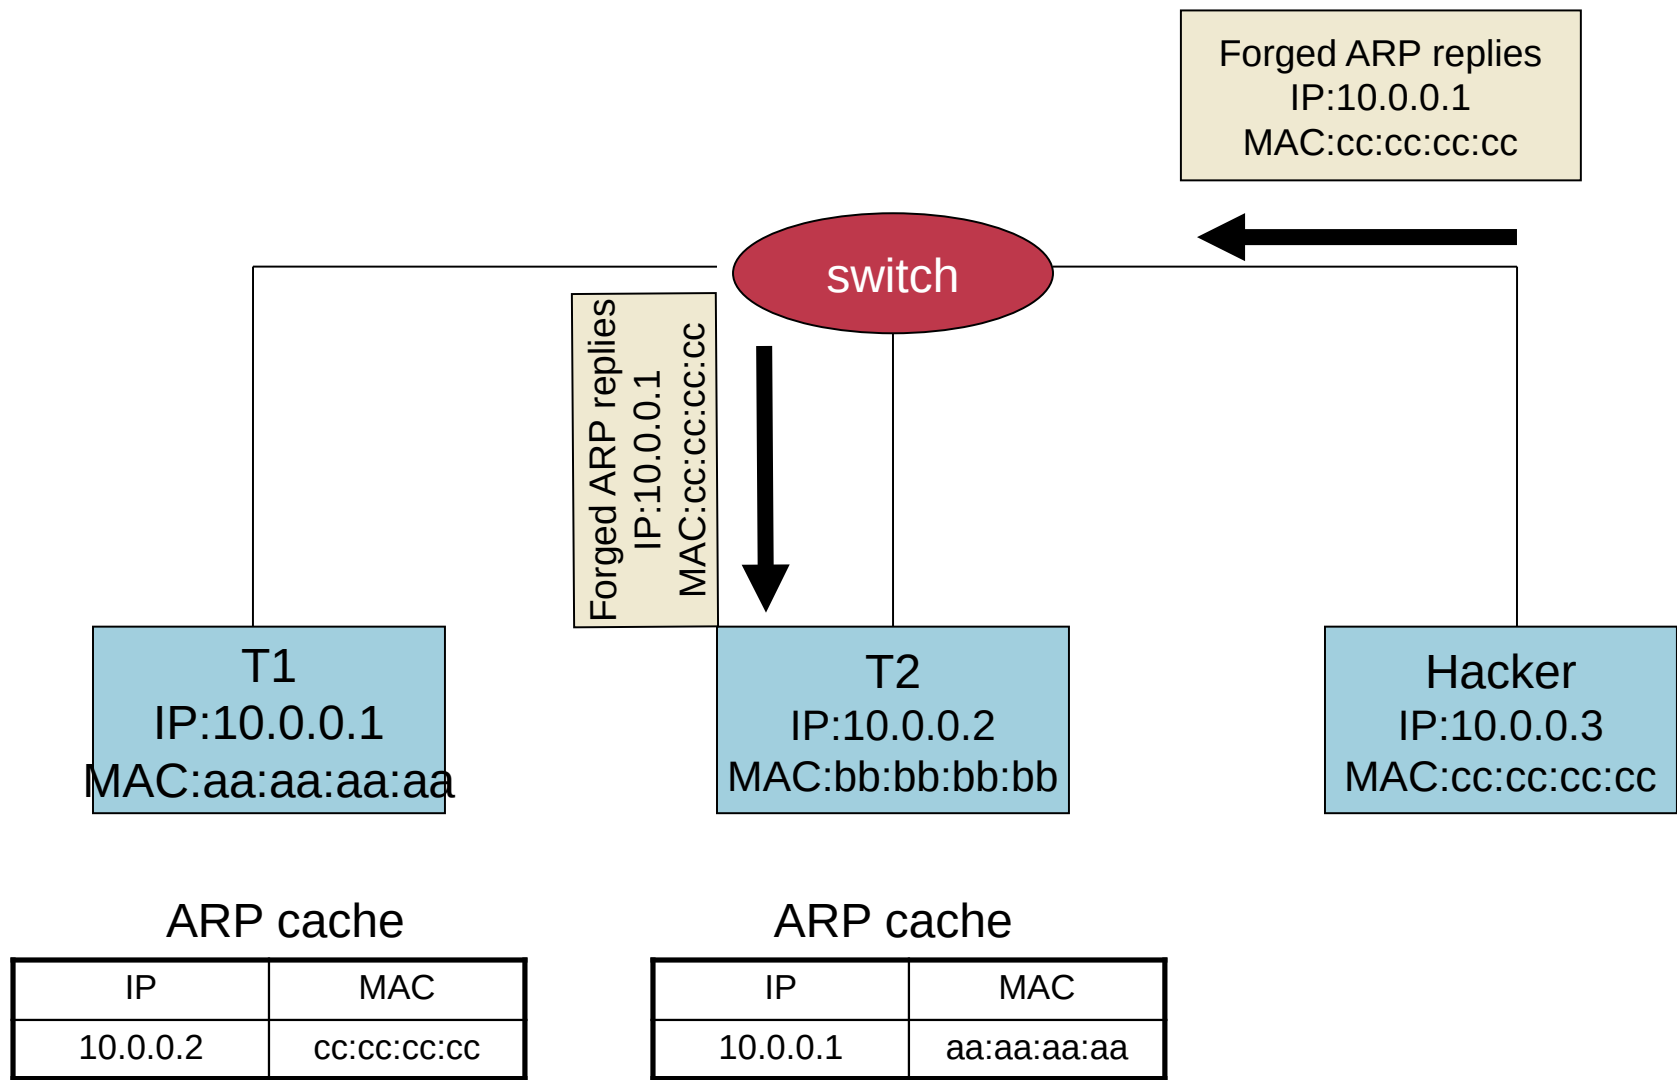
ARP cache


IP	MAC
10.0.0.2	cc:cc:cc:cc

T1's cache is poisoned

ARP cache

IP	MAC
10.0.0.1	aa:aa:aa:aa





ARP cache


IP	MAC
10.0.0.2	cc:cc:cc:cc

ARP cache

IP	MAC
10.0.0.1	cc:cc:cc:cc

T2's cache is poisoned

Message intended to send to T2



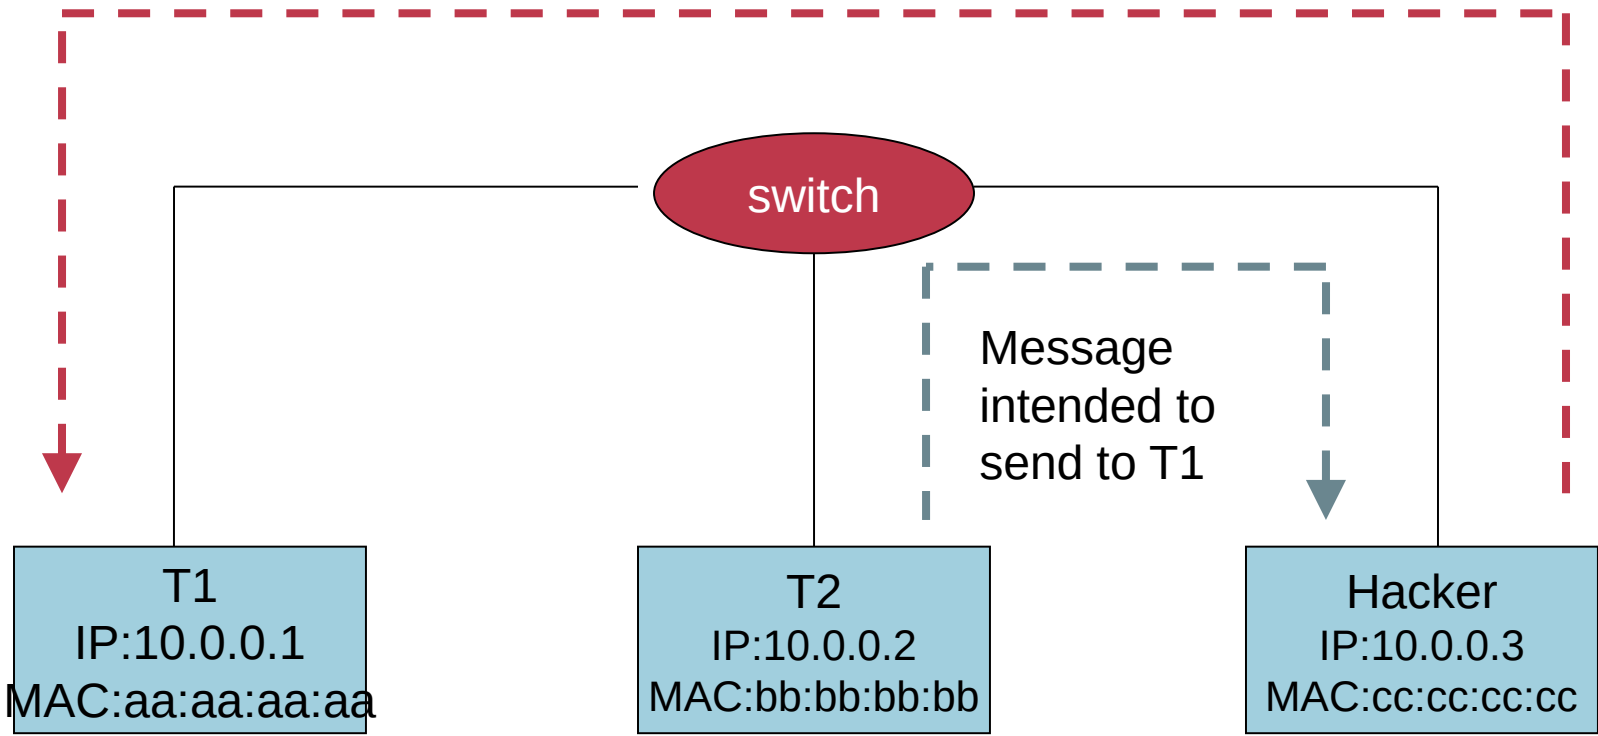
ARP cache

IP	MAC
10.0.0.2	cc:cc:cc:cc

ARP cache

IP	MAC
10.0.0.1	cc:cc:cc:cc

Hacker will relay the message



ARP cache

IP	MAC
10.0.0.2	cc:cc:cc:cc

ARP cache

IP	MAC
10.0.0.1	cc:cc:cc:cc



# Defenses against ARP Spoofing

## No Universal defense (!)

### Use static ARP entries

- Cannot be updated; Spoofed ARP replies are ignored
- ARP table needs a static entry for each machine on the network
- Large overhead
  - Deploying these tables; Keep the table up-to-date

### Arpwatch

- A free UNIX program which listens for ARP replies on a network
- Build a table of IP/MAC associations and store it in a file
- When a MAC/IP pair changes (flip-flop), an email is sent to an administrator